

Msc-generator

A tool to draw various charts from text description
(version 8.6.3, 4 September 2025)

Zoltan R. Turanyi

Table of Contents

1	Introduction	1
2	What's new in Msc-generator 8.6	2
3	Getting started	4
3.1	The MFC GUI on Windows	4
3.1.1	Working with Charts	5
3.1.2	Zooming	6
3.1.3	Quick Access Toolbar	7
3.2	The CLI GUI	7
3.3	Tracking Mode	10
3.4	Auto Split	11
3.5	Collapsing and Expanding	11
3.6	The command-line tool	12
3.7	Office integration	12
3.7.1	OLE embedding	12
3.7.2	Alt-text embedding	13
4	Signalling Chart Language Tutorial	15
4.1	Defining Arrows	15
4.2	Defining Entities	19
4.3	Dividers	25
4.4	Drawing Boxes	27
4.5	Drawing Things in Parallel	32
4.6	Annotating the Chart	33
4.7	Other Features	34
5	Graph Language Tutorial	42
6	Block Diagram Language Tutorial	45
6.1	Defining and Arranging Blocks	45
6.2	Arrows and Lines	51
7	Usage Reference	56
7.1	Multiple Chart Types	56
7.2	Design Library	56
7.3	External Editor	57
7.4	Internal Editor	58
7.4.1	Smart Indent	58
7.4.2	Color Syntax Highlighting	59

7.4.3	Typing Hints and Autocompletion	60
7.5	Example Library	61
7.6	Renaming elements	61
7.7	Options	62
7.8	Working with Multi-page Charts	63
7.9	Scaling Options	64
7.10	Advanced OLE Considerations	65
7.10.1	Graphics of Embedded Charts	65
7.10.2	Linking	66
7.11	Autosave and Recovery	66
7.12	International Text Support	67
7.13	Command-Line Reference	68
7.13.1	Label Maps	68
7.13.2	Coloring Input Files	69
7.13.3	Embedding Charts	69
7.13.4	Text Editor Integration	70
7.14	Fonts	74

8 Reference for Common Language Elements 75

8.1	Labels	75
8.2	Text Formatting	76
8.3	Links	79
8.4	Numbering	80
8.5	Specifying Colors	83
8.6	Common Attributes	85
8.6.1	Line and Fill Attributes	85
8.6.2	Shadow Attributes	87
8.6.3	Text Formatting Attributes	87
8.6.4	Styles	88
8.7	Scoping	89
8.8	Defining Styles	90
8.9	Chart Designs	91
8.10	Defining Shapes	92
8.11	Procedures	94
8.12	Variables	97
8.13	File Inclusion	98

9 Signalling Chart Language Reference 99

9.1	Titles	99
9.2	Specifying Entities	99
9.2.1	Entity Positioning	99
9.2.2	Group Entities	100
9.2.3	Entity Attributes	101
9.2.4	Implicit Entity Definition	103
9.2.5	Entity Headings	103
9.2.6	Entity Shapes	104
9.3	Specifying Arrows	104

9.3.1	Lost Messages	106
9.3.2	Arrow Attributes	107
9.3.3	Arrow Appearance	112
9.3.4	Block Arrows.....	112
9.4	Boxes	114
9.4.1	Box Series.....	115
9.4.2	Box Tags.....	116
9.5	Pipes	117
9.6	Verticals	119
9.7	Dividers.....	124
9.8	Notes and Comments.....	125
9.8.1	Notes	126
9.8.2	Comments and Endnotes	127
9.9	Parallel Blocks	127
9.9.1	Parallel Keyword	131
9.9.2	Overlap Keyword.....	131
9.9.3	Joining Arrows and Boxes	132
9.10	Signalling Chart Attributes and Styles	133
9.10.1	Appearance	133
9.10.2	Word Wrapping and Long Labels	135
9.10.2.1	Word Wrapping for Signalling Charts	135
9.10.2.2	Word Wrapping for Block Diagrams	136
9.10.3	Compression and Vertical Spacing	136
9.10.4	Default and Refinement Styles for Signalling Charts.....	137
9.11	Chart Options	139
9.12	Multiple Pages	142
9.13	Free Drawing.....	143
9.13.1	Spacing	143
9.13.2	Symbols.....	144
9.13.3	Inlining other charts.....	149
9.13.4	Inline text.....	150
9.14	Commands.....	150
9.15	Mscgen Backwards Compatibility.....	152
10	Graph Language Reference	154
10.1	Graph Attributes.....	154
10.2	Default and Refinement Styles for Graphviz Graphs.....	155
10.3	Clusters.....	156
11	Block Diagram Language Reference	157
11.1	Block Name Resolution.....	157
11.2	Block Types and Definition.....	158
11.3	Block Attributes	162
11.4	Block Layout.....	166
11.4.1	Default alignment	169
11.4.2	Alignment modifiers	171
11.4.3	Content placement inside another block.....	173

11.4.4	Block Layout Conflicts	174
11.5	Arrows in Block Diagrams.....	175
11.5.1	Defining Arrows and Lines.....	175
11.5.2	Ports and Directions.....	176
11.5.3	Fine-tuning Arrow Ends	178
11.5.4	Automatic De-Overlapping at Arrow Ends	178
11.5.5	Defining Coordinates	180
11.5.6	Arrow Labels and Markers.....	181
11.5.7	Arrow and Line Attributes	183
11.6	Replicating parts of the Diagram	186
11.6.1	Copying a Block.....	186
11.6.2	Block Templates.....	187
11.6.3	Repeating a Block Many Times.....	188
11.7	Chart Options and Commands	188

1 Introduction

This manual is for Msc-generator (version 8.6.3, 4 September 2025), a tool to draw various charts from a textual description.

Please visit <https://gitlab.com/msc-generator/msc-generator/-/packages> to download the latest version.

Msc-generator is a program that parses textual chart descriptions and produces graphical output in a variety of file formats, or can embed charts in documents, such as Word or PowerPoint. It currently supports three kinds of charts: Message Sequence Charts (MSCs, this is where the name of the tool comes from); general graphs in the DOT language of graphviz; and experimental Block Diagrams.

Message Sequence Charts are a way of representing entities and message interactions between those entities over some time period. MSCs are often used in combination with SDL. MSCs are popular in telecom and data networks and standards to specify how protocols operate. MSCs need not be complicated to create or use. Msc-generator aims to provide a simple text language that is clear to create, edit and understand, and which can be transformed into images. Msc-generator is a potential alternative to mouse-based editing tools, such as Microsoft Visio.

The signalling chart part of msc-generator is heavily extended and completely rewritten version of the 0.08 version of Michael C McTernan's mscgen. The original tool was more geared towards describing interprocess communication, this version is more geared towards networking. Msc-generator has a number of enhancements compared to mscgen. The command-line syntax of Msc-generator is compatible to that of mscgen, so any tool integrated with mscgen (such as Doxygen) can also be used with Msc-generator. Since version 4.5 Msc-generator also contains an *mscgen compatibility mode*, which aims to interpret mscgen chart descriptions in a fully backwards compatible manner. See Section 9.15 [Mscgen Backwards Compatibility], page 152.

The graph part of msc-generator uses the graphviz library to lay out graphs. It uses the DOT language to describe charts, with a few extensions.

Block Diagrams use a language similar in logic to the above two. The language aims to capture the relation among blocks to lay them out in relation to one another.

Msc-generator builds on lex, yacc, graphviz, glpk, cairo and ImGui. A Linux/Mac and Windows port is maintained. The Windows version is written using MFC. There is a GUI for Linux/Mac based on ImGui (called *CLI GUI* also available on Windows, if you prefer that better than the MFC GUI). The reference documentation of the source code of Msc-generator itself can be found at <https://msc-generator.gitlab.io/msc-generator/doxygen>.

2 What's new in Msc-generator 8.6

The improvements added since version 8.5 are listed below. If you are new to Msc-generator, you should probably jump to Chapter 3 [Getting Started], page 4.

- Added chart inlining which allows to draw a chart of any type in a signalling chart or block diagram. This means that you can draw a signalling chart in a block or add a graph to a signalling chart. Graphviz graphs cannot inline other charts, as this would require large changes to its grammar, not only simple extensions.

In signalling charts use the `symbol block` or `symbol gv` command followed by a location and the inlined chart's text enclosed between <<> and >>> symbols.

For Block Diagrams, the new `chart` block type can be used followed by chart type and the chart text similarly enclosed. See the Examples.

- [signalling] Now you can position symbols (including inlined charts as above) in relation to other elements of the chart. So far you had to specify horizontal and vertical position in relation to entity lines and markers. Now you can position symbols left of, right of, above, below or even around other chart elements. See the example when searching for '`ref`'.
- [block] Added `aspect_ratio` attribute to blocks.
- [block] Added `grid` routing type to route either horizontally or vertically.
- [block] Added new ports to boxes, like `A@s30`, which means places the end of an arrow at 30% of the south (downward) edge of the box.
- [block] Improvements to arrow layout and de-overlapping
- [block] Added `tangent` and `tangent_upside_down` for `label.orient` for arrows. These rotate the label to align with the direction of the arrow.

And many fixes and improvements in the CLI GUI, hinting and elsewhere:

- Improved contour performance by another 24%, bringing the total average reduction in execution time of Expand and Union operations since v8.4 to 42% (almost x2 speed with operation on large contours being 3-6x faster).
- [gui] fix: Tab key moves the cursor when no indenting
- [gui] fix: Ctrl+bksp shall only delete backwards
- [gui] fix: Ctrl+X/C on a full line shall make Ctrl+V insert a full line (even on last line)
- [gui] fix: pressing End shall remove selection even if at the end of the selection (even on last line)
- [gui] fix: last character not colored
- [gui] fix: Shift+Left shall select only to the end of the word
- [gui] fix: Shortcuts to dialog buttons.
- [gui] fix: Starting help does not block the app
- [gui] fix: Error squiggles now show up at the line ends (8.6.1)
- [signalling] fix: note keyword with nothing after shall parse gracefully
- [signalling] fix: byref vertical to {}
- [block] fix: color token should allow negative numbers, to be used in setting `margin`

- [signalling] fix: `adjust` shall not decrease the size of the symbol negative. Emit a warning if the symbol disappears.
- [signalling] fix: `size` for symbols should also accept a number.
- [signalling] fix: increase chart size when closing an activated entity
- [block] fix: names for cloned unnamed blocks
- fix: pedantic the same for inline chart CSH
- [graph] fix: Rename Entities did not work with multiple graphs.
- [gui] more CLI options available with GUI
- [gui] fix: hinting (also in languages)
- [block] fix: retail labels of joined blocks even with fill color
- [graph] fix: a crash on Windows with A->A edges (#94, thanks Jiří)
- [cli] fix: starting with `-tei-mode` without `-S` no longer crashes (8.6.1)
- [signalling] fix: 2 kinds of crashes (8.6.3, thanks Gořan!)


3 Getting started

Msc-generator is a program that parses textual chart descriptions and produces graphical output in a variety of file formats, or can embed in documents, such as Word or PowerPoint. It currently supports three kind of charts: Message Sequence Charts (MSCs, this is where the name of the tool comes from), Graphs using the graphviz language and Block Diagrams.

On Windows Msc-generator is installing as two applications: `Msc-generator.exe` and `msc-gen.exe`. The first is called the *MFC GUI* that can use OLE to embed charts into PowerPoint and Word documents and is supported from version 2.2. The latter implements the command line, but when invoked as `msc-gen.exe --gui` the *CLI GUI* starts, which supports *alt-text embedding*, see more on this in Section 3.7 [Office integration], page 12. You can start both GUIs directly. Clicking on a file with `.signalling`, `.graph` or `.block` extension starts the MFC GUI. You can also double-click a chart OLE embedded in a document to edit it in the MFC GUI.

Msc-generator is also supported on Debian Linux (package `msc-generator`) and on MacOS (`brew install msc-generator`). On these platforms only the CLI GUI is available.

3.1 The MFC GUI on Windows



The MFC GUI window has the usual elements of a Windows application: menu bar, a ribbon and a status bar. We will briefly discuss these here and give a more detailed description in Chapter 7 [Usage Reference], page 56.

You can use the scrollbars to navigate around in the chart. You can also grab the chart by the mouse and drag it (if not all of it fits into the window).

You can also reposition the pane of the internal editor and the error list by clicking on their title bar and dragging them to a new location. On the example above, the internal editor has been moved to the right side from the left (which is the default). You can even create floating windows out of these panes.


If you accidentally close the internal editor, use the ‘Text Editors...’ button on the ribbon and re-select ‘Internal Editor’.

3.1.1 Working with Charts


Msc-generator has a built-in text editor with color syntax highlighting. You can edit the chart description there. When you are ready, press the ‘Update Chart’ button on the very left of the ribbon (or F2 on the keyboard) and the visual view of the chart will get updated. There is also an option to automatically update the chart as you type, see Section 7.7 [Options], page 62. Any error or warning messages will show up in the error panel at the bottom. You can click on them to jump to the location of the error in the input file or use F8 and Ctrl+F8 to access them. If the error is not in the main file (but in an included file or design library), Msc-generator does not load it, but merely gives a small sound.

You can use the Main button on the ribbon or the quick access items in the window title to load/save the file. The file format is simply text, the very same that you edit inside Msc-generator’s text editor. You can also save the graphical rendering of the file in various graphics formats using the Main|Export... item. A third alternative is to export the file as a PNG image and save the chart text with it. Then opening the PNG file, will also bring back the chart text. To do this, you simply need to select PNG as a file format in the Open and Save dialog boxes. Pressing the Main button you also find the usual Print and Print Preview commands.

The first pane of the Ribbon always shows the type of chart currently opened (such as Signalling Chart, Graphviz Graph or Block Diagram). The buttons in the second column of this pane enables you to enter tracking mode (see Section 3.3 [Tracking Mode], page 10); to turn automatic splitting (Section 3.4 [Auto Split], page 11) on or off; or to enable the showing of collapse/expand controls for entity groups and boxes (Section 3.5 [Collapsing and Expanding], page 11)¹.



¹ Some of these buttons may not be available for all chart types. E.g., graphs have no heading and thus does not support AutoSplit.



The third column has additional controls to govern chart appearance. The exact controls depend on the chart type. For graphs, the layout algorithm used can be selected here. For graphs and signalling charts there is a design and page selector. By selecting a chart design here you can override the selection in the source file. This is an easy way of reviewing how your chart would look like in a particular design. See Section 8.9 [Chart Designs], page 91, for more info on chart designs. With the page selectors, you can navigate around in multi-page charts. If ‘all’ is selected then pagination is ignored and the whole chart is shown. (See Section 9.12 [Multiple Pages], page 142, for more info on pagination commands for signalling charts. For graphs each graph defined in sequence ends up in a separate page.) Finally, for graphs using the ‘Collapse all’ and ‘Expand all’ buttons you can collapse or expand all cluster subgraphs with a single click.

The Clipboard pane on the ribbon has two set of Copy/Paste operations: one for text in the text editor and a separate set for the entire chart. If you use paste for the entire chart, then its whole content is replaced, whereas if you paste into the editor, the content of the clipboard will be inserted.

You can also perform undo or redo from the Edit pane of the ribbon or by pressing Ctrl+Z or Ctrl+Y. Similar search and replace operations for the text editor can also be accessed from the Edit pane.

Finally, there is a separate button in the Edit pane to start and stop the internal or an external text editor (see Section 7.3 [External Editor], page 57). The latter is useful in case you prefer to use your own editor.

3.1.2 Zooming

You can zoom the chart in and out using the commands on the Zoom pane. The zoom drop-down allows setting a specific zoom value. However, the easiest way to zoom is to use the mouse wheel with the Ctrl key pressed.

You can easily select an appropriate zoom factor by clicking certain Zoom pane buttons. Overview adjusts zoom to fit the entire chart into the window. This is useful to get an overview of a chart. *Fit to width* changes the zoom factor to fit the width of the chart to the current window. Finally, *Original Size* sets the magnification back to 100%.²


You can also make Msc-generator apply one of the above three zoom adjustments after every update, page change or window re-size by selecting checkboxes besides the above command buttons.

You can also view the chart in full screen mode, by pressing F11. Mouse zooming and panning works in this mode, as well. A small toolbar enables you to flip pages, return to the all pages view or to toggle Auto Split (if applicable to the chart, see below). You can exit full screen mode by pressing Escape.

² Msc-generator applies a maximum zoom factor for the *Overview* and the *Fit to width* modes, in order to avoid a small chart being enormously magnified. This defaults to 150%. On very large screens this may prove to be a tool small factor - thus it is possible to increase this value on the *Preferences* category of the ribbon using the *Max Auto Zoom* edit box.


3.1.3 Quick Access Toolbar

Msc-generator allows you to select some of your most frequent commands or settings to be displayed all the times at the *Quick Access Toolbar*. (That is not just when their category is selected.) The Quick Access Toolbar can be displayed above or below the ribbon. Right of the buttons there is a small arrow opening for more actions and a menu to customize them. This is also where you can select whether to display the Quick Access Toolbar above or below the Ribbon.




3.2 The CLI GUI

On all supported platforms, you can start Msc-generator in the command line with the `--gui` switch. This opens a GUI, where you can work with charts. Most of the GUI is quite straightforward, here we just list a few things you may not discover by yourself at first.



- You can use the mouse wheel to scroll and scale the editor and the chart.
- You can drag the chart to pan.
- You can press the ‘Fit to window’ and ‘Fit to width’ buttons in the menu bar to scale the chart accordingly. Checking the boxes beside these buttons will automatically re-scale the chart after each compilation.
- Right-clicking an element in the editor brings up a context menu. You can then rename elements (entities, graph nodes or blocks) or briefly flash the element (arrows, entities, boxes, graph nodes, etc.) in the chart graphics. You can also do these by pressing Ctrl-R and F4, respectively while the cursor is positioned in the desired element.

- Double-clicking the chart will enter *Tracking mode*. In this mode hoovering over an element of the chart will highlight the element and select its definition in the text editor. While in tracking mode, moving around in the text editor will highlight the element the cursor is currently in. You can also click to select and unselect chart elements - this is useful when explaining the chart in a screen sharing or presentation session. See *Presentation mode* below.
- If the chart contains multiple pages, you can use the page selector to cycle among them. (The page selector does not appear, if the chart has only one page.) Selecting page zero, will show all the pages.
- If the clipboard contains an alt-text embedded chart or an alt-text encoded chart source, the **Paste Clipboard** button on the menu bar becomes active. Clicking it will paste the chart for editing. For embedded objects you can then press Ctrl+S to save the object back to the clipboard from where you can re-insert it into PowerPoint or Word. See Section 3.7 [Office integration], page 12, for more details.



There are also a few settings to tweak.

- You can set certain chart options, which will impact the chart appearance, the errors/warnings produced, etc. By default, Msc-generator will recompile the chart every time you change it. Unchecking 'Instant compilation' will prevent that. Then you can use F2 or the 'Compile' button to manually compile.
- In the middle column, you can adjust the zoom and its settings. You can also turn *Auto heading* mode off if you do not want the signalling chart headers to be shown always. (*Auto heading* has an effect only for signalling charts.) You can also enter *Presentation mode* here. In this mode only the chart is shown, but in full screen. You can press the Escape key to exit Presentation mode. Double-clicking the chart in presentation mode will enter tracking mode allowing you to explain the chart easier.
- In the third column you can experiment with editor settings. First, Msc-generator saves the chart text every second if the chart is associated with a file. You can turn this off by setting 'Auto Save' off. Turning 'Auto Paste Clipboard' on will make Msc-generator automatically open any alt-text embedded chart copied to the clipboard and bring Msc-generator to the front. See more on this in Section 3.7.2 [Alt-text embedding], page 13. The error squiggles provided in the editor are a subset of the error messages

received during a compilation. They are generated with color syntax highlighting and are useful as instant error feedback for charts that take longer to compile. If you turn *Smart indent* off, the Tab key will indent (Shift+Tab will de-indent) the line even if you are at the end of it.

You can use the following command-line switches alongside `--gui`.

-S <chart type>

Adding this will skip the welcome screen and start the GUI with a fresh copy of the specified chart. If you also specify a filename, then here you can force its type.

<input filename>

The file specified will be opened. Its chart type will be deduced from its extension or its content if it is a PNG file containing a chart (saved by Msc-generator).


--nodedesigns

This will skip loading the design libraries as usual.

Some other switches will abort (such as `--utf16` to force a UTF-16 encoded input), because the GUI does not support them, but some switches will just be silently ignored (such as `-q` for quiet output).

The GUI has a settings file `.msc-generator-ini`, which stores window positions, settings, the recently opened file list and the autosave. On Linux and the Mac it is placed into `$HOME/.msc-genrc` by default or to `$MSC_GEN_RC/.msc-genrc`, if `MSC_GEN_RC` is set. On Windows it is in the `AppData\Roaming\Msc-generator` folder. You can freely delete this file to erase settings and history.

Finally, the command-line GUI has an example store similar to the other GUI (See Section 7.5 [Example Library], page 61). Selecting **Examples** from the menu bar opens a dialog box with example snippets demonstrating the features of each language.



The examples are keyword searchable, just start typing the topic you have in mind to narrow the list of examples. You can then copy and paste chart text from the examples to your chart. In addition, the examples are editable allowing you to experiment with the language there and then. (Hinting with **Ctrl+Space** works, too.) Any changes made are immediately visible in the chart on the right. The examples are reset to their original text each time you exit Msc-generator.

3.3 Tracking Mode

In both GUIs if you click an arrow, entity, graph node, block or any other visual element on the chart, it is briefly highlighted and the corresponding text is selected in the internal editor. This is useful to quickly jump to a certain element in the chart text. The same way, when pressing F4, if the cursor is inside a chart element, the element is briefly flashed.

If you double-click the chart (try the background) you enter Tracking Mode, where you can select and highlight multiple chart elements at the same time. Visual elements are faintly outlined just by hovering above them. When actually clicking on them (or pressing F4 when the internal editor cursor is inside them) will make them actually highlighted. Then you can move on to highlight more. You can also remove highlighting from an element by clicking it again or by pressing F4 again.

If you click the chart background or press Escape, all highlighting is removed. Pressing Escape again will make you exit tracking mode. You can also enter tracking mode by pressing **Ctrl+T**. In the MFC GUI pressing the ‘Tracking Mode’ button on the Chart pane will also enter tracking mode.

Tracking mode is useful when you are discussing a chart (e.g., in a shared screen session) and you want to talk about various elements.

3.4 Auto Split

When working with a large signalling chart, it is sometimes needed to zoom in to an area of it. In case the viewing area is towards the bottom of the chart, it is often difficult to know which entity line belongs to which entity. In such cases turning Auto Split on will result in the splitting of the view into two parts, the upper one showing the entity headings. If zooming is applied Msc-generator always attempts to resize the upper view part to show the entities only.


Note that it is possible to define charts where there is no meaningful row of entity headings at the top. In such cases, Msc-generator will get confused and Auto Split is of no use. In case of multi-page charts, Auto Split will show automatic headers only (Instead of ‘newpage; heading;’ use rather ‘newpage [auto_heading=yes];’ command.)

Note that you can use the slider to change, where the chart is split, both with and without Auto Split. The difference is that when Auto Split is on, the split is reset to headings after a compilation or a page change.

Auto Split also works in Full Screen mode, but is available only for signalling charts.

3.5 Collapsing and Expanding

Msc-generator allows you to collapse boxes, entity groups (signalling charts); blocks (block diagrams); and cluster subgraphs (graphviz graphs). This way you can show only a simplified view of the procedure, diagram or graph described by the chart text. E.g., instead of many arrows comprising a part of the procedure, a simple box is shown as a summary.



If you move the mouse over a chart element that can be collapsed (or is already collapsed), control icon(s) appear at its top right corner. The control with the minus sign collapses the element, the control with the plus sign expands a collapsed element, while the green arrow collapses the element into a block arrow. The last icon will only appear for boxes of signalling charts, which are not part of a box series (Section 4.4 [Drawing Boxes], page 27).

In the MFC GUI you can disable the showing of such controls via the red plus button on the Chart pane. This improves compilation time a bit.

For signalling charts and block diagrams expanding and collapsing can also be set via the ‘collapsed’ attribute (for group entities/boxes or blocks, respectively) and hence is available for the command-line version, as well. It is most useful, however, for interactive work. Any collapse/expand setting via the GUI overrides the one specified by attributes. Such overrides are saved with embedded charts, but naturally not when the chart is saved to disk as text.

In both GUIs, double clicking any element that has controls (can be collapsed/expanded) activates the first control (even if controls are not shown). This essentially toggles collapse/expand status.

3.6 The command-line tool

The command-line version of Msc-generator runs on all supported platforms. On Windows it is installed to the same directory as the windowed application. That directory is included in the PATH, so you can call it from anywhere.

The command line version of Msc-generator supports PNG, PDF, EPS, SVG and PPT graphics output formats, plus EMF on Windows. It can compile and extract chart texts embedded in PNG and PPT files and it can also embed the result back to a PNG or PPT file (the latter using Section 3.7.2 [Alt-text embedding], page 13).

To use Msc-generator to generate a signalling chart from a text file (containing a signalling chart description in the appropriate language) simply type

```
msc-gen -T pdf inputfile.signalling
```

To use Msc-generator to generate a graph or block diagram from a text file (containing a graph description in the DOT or block languages, respectively) simply type

```
msc-gen -T pdf inputfile.graph
```

or

```
msc-gen -T pdf inputfile.block
```

All of these will give you `inputfile.pdf`. You can change ‘pdf’ to get the other file formats. If you omit the ‘-T’ switch altogether, a PNG will be generated.

If Msc-generator has successfully generated an output, it prints ‘Success.’. Instead, or in addition, it may print warnings or errors, when it does not understand something.

3.7 Office integration

3.7.1 OLE embedding

Since one of the main uses of Msc-generator is to paste charts into presentations or documents, Msc-generator supports OLE embedding of charts in the MFC GUI from version 2.2.

You can take a chart and embed it as a component in a compound document such as a Word, Excel or Powerpoint document. To do this, copy the chart to the clipboard by clicking on the *Copy Chart (OLE obj)* button and paste it into the compound document³. Later you can edit the chart by double clicking the chart in the document⁴.

Right clicking an embedded chart in a document will bring up a menu of options, where you can select **Edit** or **Open** for editing in a separate window; or **View Full Screen** to view (but not edit) the chart in full screen.

Page, chart design and layout settings you select on the ribbon are saved with embedded documents along with entity/box collapse/expand state. Thus if you copy a multi-page

³ Make sure you paste the chart using ‘Paste Special...’ as an ‘Msc-generator Chart Object’.

⁴ In place editing is no longer supported from version 3.4.1.

chart to a document with a single page selected, only the selected page will be visible in the document. (But the source text of the whole chart is stored there.) Likewise, if you collapse/expand certain boxes or entities or subgraphs, they will show in that state in the document. Later, if you double-click the chart and open Msc-generator, you can change the page, design or collapse/expand status and these will update in the document. Note that page, chart design, layout and collapse/expand settings are not saved with the chart when saving to a text file. Only when embedded (either as OLE object, as alt-text (see below) or into a PNG file).

See more details in Section 7.10 [Advanced OLE Considerations], page 65.

Best common practice workflow You start editing the chart in Msc-generator. At some point select *Copy Chart (OLE obj)* and then use *Paste Special...* to paste into a PowerPoint or Word document (select **Msc-generator Signalling Chart Object**). Then any time later, simply double-click the chart to open and edit it. Updating the chart (with the F2 key, for example) in such a session will update the embedded graphics.

3.7.2 Alt-text embedding

Over the years OLE technology became increasingly outdated⁵ and is not perfectly supported by Sharepoint and other modern Office tools. In addition, of course, it was never supported on Linux or on the Mac, whereas Msc-generator now has a GUI for both.

Therefore a new way to add charts to Office documents has been developed in v8.0, called *alt-text embedding*. In this method the chart is inserted as a PNG image, where the alt-text (a.k.a. *alternative text*, usually used for accessibility or when the graphics cannot be shown) stores the chart text, language and other info. This method of embedding requires no special OS support, like OLE and also allows to apply various properties to the chart, such as cropping, outline, shadow, rotation, etc. Alt-text embedding is not compatible with OLE embedding in any which way.

Both the CLI⁶ and MFC⁷ GUIs of Msc-generator can now work with alt-text embedding the following ways.

1. Open a PPT file (with `pptx` extension) and simultaneously edit all the charts embedded in it. It is possible to copy snippets between the charts, but also to add a new one. Note that you cannot save the PPT file, if it is opened in PowerPoint, as PowerPoint locks the file for writing. Also, PPT files on Sharepoint drives may not merge well, when edited this way. (CLI GUI only.)
2. Copy the currently edited chart to the clipboard using the ‘File | Copy to clipboard’ menu item (CLI GUI) or the ‘Copy Chart (alt-text)’ ribbon button (MFC GUI). This will package the chart text, language, collapse/expand state, design and page with a PNG rendering of the chart into an Office document snippet. You can then paste it into PowerPoint or Word documents. (Both CLI and MFC GUIs.)
3. If you pasted an Msc-generator chart this way into a into PowerPoint or Word document and later copy it to the clipboard from PowerPoint or Word, you can paste it into Msc-generator for editing. This will re-use the chart text and other metadata stored

⁵ For example, it still uses 16-bit Windows metafiles

⁶ The CLI GUI can be invoked with `msc-gen --gui`

⁷ The MFC GUI can be invoked wit `Msc-generator.exe` and supports alt-text embedding since version 8.5.

with the image. Msc-generator stores the sizing and formatting of the chart (along with other graphical elements or text copied with the chart) so when you save, those will be preserved. Saving will happen to the clipboard from where you can re-insert the chart back to the PowerPoint or Word document (or to another one). If a slide contains multiple charts, you can copy all of them to the clipboard and edit all of them simultaneously in Msc-generator (potentially copying snippets between them). (Both CLI and MFC GUIs. The MFC GUI can only edit one chart at a time.)

4. Finally, Msc-generator can copy the alt-text for the current chart to the clipboard. Adding this to any graphics in PowerPoint will *upgrade* the graphics to an Msc-generator chart and Msc-generator will be able to edit it (either by opening the PPT file as #1 above or via the clipboard). Similarly, if you copy an alt-text to the clipboard, Msc-generator can paste it as if the whole chart has been copied. (The MFC GUI cannot copy the alt-text, but can paste it.)

Turning the setting **Auto Paste Clipboard** on, will make Msc-generator automatically open any chart copied to the clipboard (and come to the foreground) if the currently open chart has no unsaved changes.

See Section 7.13.3 [Embedding Charts], page 69, on how to use the command-line tool to embed and extract charts into and from PNG and PPT files.

Best common practice workflow You start editing the chart in Msc-generator. At some point select *File|Copy to Clipboard* (CLI GUI) or *Home|Copy Chart (alt-text)* from the ribbon (MFC GUI) and then simply paste into a PowerPoint or Word document. (Turn **Auto Paste Clipboard** on.) Then any time later (with Msc-generator running), 1) cut the chart (and maybe other graphics elements or multiple charts) to the clipboard; 2) Msc-generator pops up automatically; 3) edit the chart; 4) Save; and 5) Paste the chart(s) (and other elements) back to the PowerPoint or Word document.


4 Signalling Chart Language Tutorial

In this chapter we give a step-by-step introduction into the language of Msc-generator for signalling charts. At the end you will master most of the language to create charts. Further details (mostly on controlling appearance) are provided in Chapter 9 [Signalling Chart Language Reference], page 99.

4.1 Defining Arrows

Message sequence charts consists of *entities* and *messages*. The simplest file consists of a single message between two entities: a ‘*Sender*’ and a ‘*Receiver*’.

Sender->Receiver;



The message may have a label, as well.


Sender->Receiver: Message;



A more complicated procedure would be to request some information from a server, which, in turn, queries a backend. Note that everything in a line after a '#' is treated as a comment and is ignored by Msc-generator.

```


#A more complex procedure
Client->Server: Request;
Server->Backend: Query;
Server<-Backend: Answer;
Client<-Server: Response; #final
  
```



Arrows can take various forms, for example they can be bi-directional or can span multiple entities. They can also start and end at the same entity and can come from or go to “outside”

```

->Client: Hit;
Client->Server: Request;
Server->Backend: Query;
Backend->Backend: Lookup;
Client<-Server<-Backend: Reply;
Client<->Backend: Direct Query;
  
```




Sometimes one wants to indicate that a message came from an entity not shown, but not at the far left or right. In this case use the pipe symbol ‘|’, like a->|;.

```

Client;
|->Server: Req;
Server->Backend: Query;
Backend->Backend: Lookup;
|<-Server<-Backend: Reply;

```




It is also possible to make use of various arrow types, such dotted, dashed and double line. To achieve this the ‘->’ symbol need to be replaced with ‘>’, ‘>>’ and ‘=>’, respectively.

```

->Client: Hit;
Client=>Server: Request;
Server>>Backend: Query;
Server<<Backend: Response;
Client<=Server: Reply;
Client<>Backend: Optional Query;

```




It is also possible to use different line styles for different segments of an arrow - but all must be of the same direction. (That is, it is not possible to write ‘a->b<-c’, for example.) In addition, for multi-segment arrows the dash ‘-’ symbol can be used in the second and following segments, as a shorthand. In this case the added segment will have the same line style as the first one.

```


Client<=>Server<>Backend:
    Request with optional part;
Client=>Server-Backend: Request;
Client<=Server-Backend: Response;
Client<>>Server-Backend: Query;

```




It is possible to change the type of the arrowhead. The arrowhead type is an *attribute* of the arrow. Attributes can be specified between square brackets before or after the label, as shown below. A variety of arrow-head types are available, for a full list of arrow attributes and arrowhead types See Section 9.3 [Specifying Arrows], page 104.

```
Client->Server-Backend: Request
  [arrow.midtype=dot];
Client<-Server-Backend: Response
  [arrow.type=line];
Client<>>Server-Backend: Query
  [arrow.type=empty];
```




Msc-generator generates an error if the order of entities in a multi-segment arrow does not follow the order of the entities (either left-to-right or back). However, sometimes it is important to show such a message zig-zagging among the entities as one message. This is possible by *joining* arrows. Note that the ‘join’ keyword can do more, see Section 9.9.3 [Joining Arrows and Boxes], page 132.

```
A, B, C, D;
A->C: First segment;
join C->B;
join B->D: Last segment;
```




Often the message has not only a name, but additional parameters, that need to be displayed. The label of the arrows can be made multi-line and one can apply font sizes and formatting, as well. This is achieved by inserting formatting characters into the label text. Each formating character begins with a backslash ‘\’. ‘\b’, ‘\i’ and ‘\u’ toggles bold, italics and underline, respectively. ‘\-' switches to small font, ‘\+' switches back to normal size, while ‘\^’ and ‘_’ switches to superscript and subscript, respectively. ‘\n’ inserts a line break. You can also add a line brake by simply typing the label into multiple lines. Leading and tailing whitespace will be removed from such lines so you can indent the lines in the source file to look nice.

```
->Client: Hit\_\_1;
Client<->Server-Backend: Request\_\_1\+\ with reply;
->Client: Hit\_\_2;
Client->Backend: Request\_\_2\-\Parameter=\b2;
Client<-Backend: Reponse
  \ -May contain an error;
```




Arrows can further be differentiated by applying styles to them. Styles are packages of attributes with a name. They can be specified in square brackets like an attribute that takes no value. Msc-generator has two pre-defined styles ‘`weak`’ and ‘`strong`’, that exists in all chart designs¹. They will make the arrow look less or more emphasized, respectively. The actual appearance depends on the chart design, in this basic case they represent gray color and thicker lines with bold text, respectively².

```
->Client: Hit [strong];
Client->Server: Request;
Server->Backend: Query;
Server<-Backend: Response;
Client<-Server: Reply;
Client<->Backend: Optional Query
    [weak] ;
```



Msc-generator places arrows one-by-one below each other. In case of many arrows, this may result in a lot of vertical space wasted. To reduce the size of the resulting diagram, a *chart option* can be specified, which compresses the diagram, where possible. You can read more on chart options, see Section 9.10.3 [Compression and Vertical Spacing], page 136.

```
compress=yes;
->Client: Hit [strong];
Client->Server: Request;
Server->Backend: Query;
Server<-Backend: Response;
Client<-Server: Reply;
Client<->Backend: Optional Query
    [weak] ;
```



You can use the ‘`angle`’ chart option (or attribute) to make the arrows slanted. Simply specify a value in degrees. Note that bi-directional arrows will not be slanted.


¹ You can define your own styles, as well, see Section 8.8 [Defining Styles], page 90.

² For more on chart designs Section 8.9 [Chart Designs], page 91.

```

compress=yes;
angle=3;
->Client: Hit [strong];
Client->Server: Request;
Server->Backend: Query;
Server<-Backend: Response;
Client<-Server: Reply;
Client<->Backend: Optional Query
[weak];

```




Normally, Msc-generator attempts to avoid overlaps between elements by placing them one below (or sometimes besides) each other. If you want to show messages crossing each other, you need overlapping arrows. For this, you can use the `overlap` keyword. Arrows marked such are allowed to be overlapped by subsequent arrows.

```

hscale=1.5;
angle=10;
H1: Host 1;
H2: Host 2;
overlap H1->H2: \p1Message1;
H2->H1: \prMessage2;

```




Finally, you can also indicate a lost message by marking the segment of the loss with an asterisk '*'.

```

defstyle arrow [text.ident = left];
->Client: Hit;
Client->Server: Request;
Server->*Backend: Query;
...: Timeout...
Server->Backend: Retransmission;
Server<-Backend: Response;
Client<-Server: Reply;

```




4.2 Defining Entities

Msc-generator, by default draws the entities from left to right in the order they appear in the chart description. In the examples above, the first entity to appear was always the 'Client', the second 'Server' and the third 'Backend'.

Often one wants to control, in which order entities appear on the chart. This is possible, by listing the entities before actual use. On the example below, the order of the entities are reversed. Note that we have reversed the first arrow to arrive to the ‘Client’ from the right.


```
Backend, Server, Client;
Client<- : Hit;
Client==>Server: Request;
Server>>Backend: Query;
Server<<Backend: Response;
Client<=Server: Reply;
```



Often the name of the entity need to be multi-line or need to contain formatting characters, or is just too long to type many times. You can overcome this problem by specifying a label for entities. The name of the entity then will be used in the chart description, but on the chart the label of the entity will be displayed. The ‘label’ is an attribute of the entity and can be specified between square brackets after the entity name, before the comma, as shown below. (You can specify entity attributes only when explicitly defining an entity and not if you just start using them without listing them first.)

```
B [label="Backend\\n\\- (Datastore)"] ,
S [label="Server\\n\\- (Frontend)"] ,
C [label="Client"] ;

C<- : Hit;
C=>S: Request;
S>>B: Query;
S<<B: Response;
C<=S: Reply;
```




You can also use the colon-notation to specify entity labels, similar to arrows. The above example can thus be written as below. Note that the entity definitions are now terminated by a semicolon – commas would be treated as part of the label.

```

B: Backend<br> - (Datastore) ;
S: Server<br> - (Frontend) ;
C: Client;

C<-: Hit;
C=>S: Request;
S>>B: Query;
S<<B: Response;
C<=S: Reply;

```




Entities can also be specified as ‘weak’ or ‘strong’, by applying these styles the same way as for arrows. You can also assign various shapes to the entity headings via the `shape` attribute.

```

Backend [shape=def.oval];
S: Server<br> - (Frontend) [strong] ;
C: Client [weak] ;

C<-: Hit;
C=>S: Request;
S>>B: Query;
S<<B: Response;
C<=S: Reply;

```




Entities can be turned on and off at certain points in the chart. An entity that is turned off, will not have its vertical line displayed. This is useful if the chart has many entities, but one is involved only in a small part of the process. An entity can be turned off by typing `hide` followed by the name of the entity. You can turn it later back on with the `show` keyword followed by the entities to turn on. When `hide` is used for an entity right at its definition, it will start hidden and its heading is not drawn at the place of definition. However, when it is later turned on, a heading will be shown.

```

C: Client;
S: Server\n\t- (Frontend) ;
hide B: Backend;
->C: Hit;
C=>S: Request;
show B;
S>>B: Query;
S<<B: Response;
hide B;
C<=S: Reply;

```




Not showing an entity from the beginning of the chart can also be achieved by simply defining the entity later. Note that this is different from simply starting to use an entity later. When you start using an entity without explicitly defining it first, it will appear at the top of the chart, not only where started using it first. (See earlier examples.)

```

C: Client;
S: Server\n\t- (Frontend) ;
->C: Hit;
C=>S: Request;
B: Backend;
S>>B: Query;
S<<B: Response;
hide B;
C<=S: Reply;

```




Sometimes the vertical space between entities is just not enough to display a longer label for an arrow. In this case use the 'hscale' chart option to increase the horizontal spacing. It can be set to a numerical value, 1 being the default.

```

hscale=1.3;
C: Client;
S: Server;
->C: Hit;
C=>S: Very Long Request;
B: Backend;
S>>B: Query;
S<<B: Response;
B [show=no];
C<=S: Very Long Reply;

```




Or you can simply set it to ‘auto’, which creates variable spacing, just as much as is needed.

```

hscale=auto;
C: Client;
S: Server;
->C: Hit;
C=>S: Very Long Request;
B: Backend;
S>>B: Query;
S<<B: Response;
B [show=no];
C<=S: Very Long Reply;

```




Alternatively, you can instruct Msc-generator to apply word wrapping to the labels of arrows, to fit into the available space, by setting the ‘text.wrap’ chart option to ‘yes’.

```


text.wrap=yes;
C: Client;
S: Server;
->C: Hi;
C=>S: Very Long Request;
B: Backend;
S>>B: Query;
S<<B: Response;
B [show=no];
C<=S: Very Long Reply;

```




It is possible to define entity groups, to indicate logical relations between various entities. Use curly braces ('{' and '}') after an entity definition (after any potential label and attributes).

```
hscale=auto;
C: Client;
SI: Server Infrastructure {
    S: Server;
    B: Backend;
};
->C: Hit;
C=>S: Request to Server;
S>>B: Internal Query;
S<<B: Internal Response;
C<=S: Reply from Server;
```



Instead of a group heading, you can also shade the background behind entities in an entity group.

```
hscale=auto;
C: Client;
SI: Server Infrastructure
    [large=yes] {
        S: Server;
        B: Backend;
    };
->C: Hit;
C=>S: Request to Server;
S>>B: Internal Query;
S<<B: Internal Response;
C<=S: Reply from Server;
```




It is also possible to collapse a group entity hiding details of the process. This can be done either via the 'collapsed' attribute or, on Windows, using the GUI. Elements that disappear leave a small indicator (box with 3 dots). The collapsed entity group also includes an indicator to show that further entities are hidden within. (Indicators can be turned off by the 'indicator' chart option).

```

hscale=auto;
C: Client;
SI: Server Infrastructure
[collapsed=yes] {
    S: Server;
    B: Backend;
};
->C: Hit;
C=>S: Request to Server;
S>>B: Internal Query;
S<<B: Internal Response;
C<=S: Reply from Server;

```




Entities can be *activated*. This results in the entity line becoming a thin rectangle instead. If you do this immediately after an arrow the activation will happen at the tip of the arrow indicating that the cause of the activation is the arrow.

```

hscale=auto;
C: Client;
S: Server;
B: Backend;
->C: Hit;
C=>S: Request to Server;
activate S;
S>>B: Internal Query;
S<<B: Internal Response;
C<=S: Reply from Server;
deactivate S;


```



4.3 Dividers


In a message sequence chart it is often important to segment the process into multiple logical parts. You can use the '---' element to draw a horizontal line across the chart with some text, e.g., to summarize what have been achieved so far.

C: Client;
S: Server;
B: Backend;
->C: Hit;
C=>S: Request;
S>>B: Query;
S<<B: Response;
C<=S: Reply;
----: Query done;
C->S [weak]: Next Request;



Similar to this, using the ‘...’ element can express the passage of time by making the vertical lines dotted.

```
C: Client;  
S: Server;  
B: Backend;  
->C: Hit;  
C=>S: Request;  
S>>B: Query;  
S<<B: Response;  
C<=S: Reply;  
...: \iSome time elapses;  
C->S [weak]: Next Request;
```




Sometimes one merely wants to add some text to a chart. In that case the empty element can be used either like ': text;'. The symbol '|||' simply inserts an empty row.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
C=>S: Request;
S>>B: Query;
: the backend is very busy here;
S<<B: Response;
C<=S: Reply;
|||;
C->S [weak]: Next Request;

```




The above construct always places the text in the middle of the chart. Smaller chart explanations are better done via the `text at` command.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
C=>S: Request;
S>>B: Query;
text at B: The backend is
            very busy here.;
S<<B: Response;
C<=S: Reply;
text at C+: The client now
            finishes.;
C->S [weak]: Next Request;

```



More options to comment and annotate your chart can be found in Section 4.6 [Annotating the Chart], page 33.


4.4 Drawing Boxes

A box is a line around one part of the chart. It can be used to add textual comments, group a set of arrows or describe alternative behavior. In their simplest form they only contain text, but they can also encompass arrows. A box spans between two entities, or alternatively around only one.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
box C--C: Generate\ncrequest;
C=>S: Request;
box S--B: Server gets info\nfrom Back
C<=>S: Reply;

```




The line around boxes can be dotted, dashed and double line, too, by using ‘..’, ‘++’ or ‘==’ instead of ‘--’. Boxes can also be used to group a set of arrows. To do this, simply insert the arrow definitions enclosed in curled braces just before the semicolon terminating the definition of the box.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
box C--C: Generate\ncrequest;
C=>S: Request;
box S..B: Server gets info
{
    S>>B: Query;
    S<<B: Response;
}
C<=>S: Reply;


```



When a box contains arrows, it is not necessary to specify which entities it shall span between, it will be calculated automatically. Also boxes can be nested arbitrarily deep.

```


C: Client;
S: Server;
B: Backend;
->C: Hit;
box ...: Server query
{
  box C==C: Generate\nrequest;
  C=>S: Request;
  box S..B: Server gets info
  {
    S>>B: Query;
    S<<B: Response;
  };
  C=<S: Reply;
}
;
```



You can shade boxes, by specifying the color attribute. For a full list of box attributes and color definitions, See Section 9.4 [Boxes], page 114, and see Section 8.5 [Specifying Colors], page 83. It is also possible to make a box ‘weak’ or ‘strong’.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
box ...: Server query
{
  box C==C: Generate\nrequest [strong];
  C=>S: Request;
  box S..B: Server gets info
  [color=lgray]
  {
    S>>B: Query;
    S<<B: Response;
  };
  C=<S: Reply;
}
;
```




A number of box contours are available via the ‘line.corner’ attribute.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
box ...: Server query
    [line.corner=round]
{
    C==C: Generate\ncrequest
        [strong, line.corner=note];
    C=>S: Request;
    box S..B: Server gets info
        [color=lgray,
         line.corner=bevel]
    {
        S>>B: Query;
        S<<B: Response;
    };
    C<=S: Reply;
}

```




Boxes can express alternatives. To do this, simply concatenate multiple box definition without adding semicolons. These will be drawn with no spaces between. Changing the line style in subsequent boxes impacts the line separating the boxes, otherwise all attributes of the first box are inherited by the subsequent ones.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
box C==C: Generate\ncrequest;
C=>S: Request;
box S--S: Check cache;
box S--B: Alt\#1: cache miss
    [color=lgray]
{
    S->B: Query;
    S<-B: Response;
}
...: Alt\#2: cache hit
{
    S->S: Read\ncache;
};
C<=S: Reply;

```




You can use *tags* to label boxes (both standalone or in a series). This can be used to indicate alternatives, loops or optional parts of the sequence.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
box C==C: Generate\ne\nrequest;
C=>S: Request;
box S--S: Check cache;
box S--B: cache miss
  [tag="Alt\#1", color=lgray]
{
  S->B: Query;
  S<-B: Response;
}
...: cache hit [tag="Altt\#2"]
{
  S->S: Read\ncache;
};
C<=S: Reply;

```



You can observe in the previous example that the ‘\#’ sequence inserts a ‘#’ character into a label. The ‘\’ is needed to differentiate from a comment.


Finally, similar to entity groups, boxes can also be collapsed, if they are not empty. Standalone boxes can be collapsed to an empty box or block arrow by specifying the ‘collapsed’ attribute (or via the GUI on Windows). This feature is useful to hide or summarize irrelevant parts of the chart and enables quick working with large processes.

```
hscale=auto;
C: Client;
S: Server;
B: Backend;
C=>S: Request;
box S -> B: Server gets info {
    S -> B: Query;
    S <- B: Response;
};

C <= S: Reply;
---: Again...;
C=>S: Request;
box S -> B: Server gets info [collapsed=yes]
    S -> B: Query;
    S <- B: Response;
};

C <= S: Reply;
---: And again...;
C=>S: Request;
box S -> B: Server gets info [collapsed=arrow]
    S -> B: Query;
    S <- B: Response;
};


C <= S: Reply;
```



4.5 Drawing Things in Parallel


Sometimes it is desired to express that two separate process happen side-by-side. The easiest way to do so is to write ‘parallel’ before any arrow, box or other element. As a result the elements after it will be drawn in parallel with it.

```
C: Client;  
S: Server;  
B: Backend;  
  
parallel B--B: State;  
C->S: Remove Req;  
S->B: Remove Req;  
S<-B: Ack;  
parallel B--B: No State;  
C<-S: Ack;  
C--C: Now we have it;
```



It is also possible to have bigger blocks of action in parallel using *Parallel blocks*. Consider the following example.

```
Left_MN, Left_AR, Server, Right_AR, Right_MN;
{
    Server->Left_AR: Query;
    Left_AR->Left_MN: Query;
    Left_AR<-Left_MN: Response;
    Server<-Left_AR: Response;
} {
    Server->Right_AR: Query;
    Right_AR->Right_MN: Query;
    Right_AR<-Right_MN: Response;
    Server<-Right_AR: Response;
};
box Server--Server: Now I have both;
```



In the above example a central sever is querying two AR entities, which, in turn query MN entities further. The query on both sides happen simultaneously. To display parallel actions side by side, simply enclose the two set of arrows between braces '{ }' and write them one after the other. Use only a single semicolon after the last block. You can have as many flows in parallel as you want. It is possible to place anything in a parallel block, arrows, boxes, or other parallel blocks, as well. You can even define new entities or turn them on or off inside parallel boxes.

The top of each block will be drawn at the same vertical position. The next element below the series of parallel blocks (the "Now I have it" box in our example) will be drawn after the longest of the parallel blocks.


4.6 Annotating the Chart

Often it is important to make annotations to the chart detailing what is going on. Msc-generator supports several types of annotations, let's start with *notes* and *comments*. Both have a *target element* to which the note or comment is made. Notes appear as small callouts in the chart and should preferably contain short text. Comments, on the other hand appear on the side and allow for more elaborate explanations.

```

defstyle note, comment [text.color=red];
defstyle note [line.color=red];
C: Client;
S: Server;
B: Backend;
->C: Hit;
box ..: Server query
{
    box C==C: Generate\nrequest;
    C=>S: Request;
    note: This must\nbe very fast;
    box S..B: Server gets info
    {
        S>>B: Query;
        S<<B: Response;
    };
    comment:
        An important part of this
        process is that it runs
        entirely inside the Backend
        infrastructure and hence
        does not impact the client.;


    C<=S: Reply;
}
;
```



Another way to explain what happens is by annotating a series of events. This can be done by *verticals*, which are elements spanning vertically usually besides an entity line. Here we show two of them, the *range* and *brace* verticals, but there are more.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
mark top;
C==C: Generate\nrequest;
C=>S: Request;
note: This must\nbe very fast;
S..B: Server gets info
{
    S>>B: Query;
    S<<B: Response;
};
vertical brace at S- :Query;
C<=S: Reply;
vertical range top<-> at C-: Whole process;
;
```




4.7 Other Features

There are a few more features that are easy to use and can help in certain situations.

One useful feature is the numbering of labels. This is useful if you want to insert your chart into some documentation and later refer to individual arrows by number. By specifying the `numbering=yes` chart option all labels will get an auto-incremented number. This includes boxes and dividers, as well. You can individually turn numbering on or off by specifying the `number` attribute. You can set it to `yes` or `no`, or to a specific integer number. In the latter case the arrow will take the specified number and subsequent arrows

will be numbered from this value. On the example below, we can observe that in case of parallel blocks the order of numbering corresponds to the order of the arrows in the source file.


```
numbering=yes;
Left_MN, Left_AR, Server, Right_AR, Right_MN;
{
    Server->Left_AR: Query;
    Left_AR->Left_MN: Query;
    Left_AR<-Left_MN: Response;
    Server<-Left_AR: Response;
} {
    Server->Right_AR: Query;
    Right_AR->Right_MN: Query;
    Right_AR<-Right_MN: Response;
    Server<-Right_AR: Response;
};
Server--Server: Now I have both [number=no];
```



Sometimes a block of actions would be best summarized by a block arrow. This can be achieved by typing ‘block’ in front of any arrow declaration.

```
C: Client;
R1: Router;
R2: Router;
S: Server;

->C: Hit;
C<->S: Query/Response\n- (normal);
block C<->S: Query/Response\n- (block);
```




Similar, many cases you want to express a tunnel between two entities and messages travelling through it. To achieve this, just type ‘pipe’ in front of any box definition. You can define a series of connected or disconnected pipe segment each with its own visual style or even encapsulate pipes. More on this in Section 9.5 [Pipes], page 117.

```

C: Client;
R1: Router;
R2: Router;
S: Server;

->C: Hit;
C==C: Generate\request;
pipe R1--R2: Tunnel {
    C=>S: Request;
};
S--S: Set up\more tunnels;
pipe R1--R2: Segment 1 []
    R2==S: Segment 2
{
    C<=S: Response;
};
pipe R1--R2: Outer
    [solid=255, color=green] {
        pipe C++S: \p1Inner
            [color=red] {
                C<=>: \prSome message;
            };
    };
}

```




Adding a title to the chart is easy. Just type **title:** followed by the title text.

```

title: This is the title;
a.b.c;
a->b: message 1;
b->c: message 2;

```

This is the title




Another handy feature is multi-page support. This is useful when describing a single procedure in a document in multiple chunks. By inserting the **newpage;** command, the rest of the chart will be drawn to a separate file. You can specify as many pages, as you want. In order to display the entity headings again at the top of the new page, add the **auto_heading=yes** attribute. Breaking a page is possible even in the middle of a box, see the following example.

```


C: Client;
S: Server;
B: Backend;
->C: Hit;
box C==C: Generate\request;
C->S: Request;
box S--S: Check cache;
box S--B: cache miss
  [tag="Alt\#1", color=lgray]
{
  S->B: Query;
#break here
newpage [auto_heading=yes];
  S<-B: Response;
}
...: cache hit [tag="Alt\#2"]
{
  S->S: Read\ncache;
};
C<=S: Reply;

```

Chunk one:



Chunk two:




From version 3.3 you can draw arbitrary circles and rectangles onto the chart. Their syntax is quite rich to allow free placement. You can even specify to draw below the entity lines or over other drawn elements. More detailed description can be found in Section 9.13 [Free Drawing], page 143, but here are a few examples.

```

mark top;
source ,middle1, middle2, destination;
vspace 10;
source->destination: \p1message \#1;
source->destination: \p1message \#2;
mark a_top;
source->destination: \p1message \#3;
mark a_bottom [offset=10];
symbol ... center at source-middle1;
source->destination: \p1message \#n;
mark bottom;
symbol rectangle top-bottom left at middle1 -40 right at middle2 +40
  [fill.color=lgray, line.type=none,
   draw_time=before_entity_lines];
symbol arc a_top-a_bottom center at destination
  [xsize=60, line.color=red, line.width=3,
   fill.color=none];

```



Finally, an easy way to make charts visually more appealing is through the use of *Chart Designs*. A chart design is a collection of colors and visual style for arrows, boxes, entities and dividers. The design can be specified either on the command line after double dashes, or at the beginning of the chart by the `msc=<design>` line.


Currently several designs are supported. ‘plain’ was used as demonstration so far. Below we give an example of the others.

```


ms c=qsd;
C: Client;
S: Server;
B: Backend;
->C: Hit [strong];
box C==C: Generate\request;
C=>S: Request;
box S--S: Check cache;
box S--B: cache miss [tag="Alt\#1"]
{
    S->B: Query;
    S<-B: Response;
}
...: cache hit [tag="Alt\#2"]
{
    S->S: Read\ncache [weak];
};
C<=S: Reply;
---: All done;

```


The ‘qsd’ design:




The ‘rose’ design:




The ‘mild_yellow’ design:




The ‘omegapple’ design:




The ‘modern_blue’ design:




The ‘round_green’ design:




The ‘green_earth’ design:




The ‘colores’ design:



The ‘black_on_white’ design:



And the the ‘norton_commander’ tribute design.




5 Graph Language Tutorial

The language of the graphs (the DOT language) is documented well on the graphviz pages at <http://www.graphviz.org/> (<http://www.graphviz.org/>). Here we give a brief introduction (using Msc-generator additions).

You need to start with either the **graph** or **strict graph** to specify a graph, the second version does not allow multiple edges between the same two nodes.

After this, you can list nodes and edges enclosed between curly braces. Although there is no need to terminate the lines with a semicolon, we suggest to do so for better readability. Node names are underlined the first time they are used (as with signalling chart entities).


```
graph {
    rankdir = LR;
    a_b;
    a->_c;
    a->_e;
    a->b->c=>d<->e
    a<<b;
    c<e;
}
```



You can use various arrow symbols (\rightarrow or \leftarrow , for example) to generate edges with different style. The **rankdir=LR** line instructs the DOT layout algorithm to lay out from left to right as opposed to from top to bottom.

You can group nodes using curly braces. This is useful to add edges to all of them in one go or to instruct the dot algorithm to lay them out on the same level. The latter can be observed with node ‘e1’ being on the same rank (left-right position) as ‘D and d’. This is because both of them are mentioned inside a set of curly braces with **rank=same**. Without this, ‘e1’ would have been laid out right of ‘D and d’ as it happens with ‘e2’.

```
graph {
    rankdir = LR;
    _a;_b;
    c [label="C and c"];
    d::D and d;
    other [color=red];
    {a, b} -> other -> {c, d};
    b:sw == d;
    {rank=same; d->e1;};
    d->e2;
}
```




On the example above you can also see how to assign attributes. The **label** attribute can be substituted with the colon-syntax, but note that for graphs, you need to use two colons to indicate the label (as opposed to signalling charts), since single colons are used to indicate ports and compass direction. The latter is useful to impact at which angle an edge

arrives/leaves the node. In the example above `sw` represents southwest, and the `==` arrow symbol results in an edge with no arrows, but double-lined.


Graphviz supports many attributes for nodes, edges and layout algorithms, try experimenting. You can also use the text formatting escapes for labels used by the signalling chart syntax, but only if you specify your label with the double colon syntax. Note that due to the peculiar design of graphviz the `style` attribute can have values that affect the line, fill or the shape of nodes. Multiple such values can be specified separated using the attribute multiple times.

```
graph {
    a::Thick [penwidth=3];
    b::Fill [fillcolor="red:blue"
               style=filled];
    c::Shape [shape=triangle];
    d::Shadow [shadow_offset=7,
                  shadow_blur=5];
    e::Style [style=dashed];
    f::Styles [style=dotted,
                  style=filled,
                  shape=box,
                  style=rounded];
    a->b::\i Formatted\n bLabel;
    c>>d::Color [color=green];
    e>f [arrowhead=onormal];
}
```



You can use the `subgraph` keyword to make a part of the graph separate. By graphviz convention, if you select a name that begins with ‘`cluster_`’ the subgraph will be visibly shown. This convention is only honoured by some of the layout algorithms. As a syntactic sugar, Msc-generator defines the `cluster` keyword, that prepends ‘`cluster_`’ to the name you supply and automatically sets the label¹

```
graph {
    rankdir = LR;
    subgraph cluster_one {
        label="Cluster one";
        style=rounded;
        style=filled;
        one1->one2;
    } => two;
    cluster Three {
        style="rounded,filled";
        three1->three2;
    } => two;
}
```



¹ If you do not specify any name, a numeric one is generated, but is not used as label.

More details on the syntax of graphviz and the Msc-generator extensions can be found in Chapter 10 [Graph Language Reference], page 154.

6 Block Diagram Language Tutorial

Very often the DOT language is used to lay out blocks in a certain geographical arrangement, e.g., for organizational charts, or similar. This is hard, since graphviz aims to generate a layout automatically and one has to tweak and workaround the layout algorithm by, for example, inserting invisible nodes, etc. This is the primary motivation for the Block Diagram language: to be able to control, how nodes of a diagram are laid out. This language targets network diagrams, architecture figures, software stack figures and any figure, where blocks are arranged and this arrangement conveys meaning. Of course, we have many mouse-based variants, but sometimes it is just easier to do it via text. It also integrates well to command-line document toolchains, such as tex or doxygen.

6.1 Defining and Arranging Blocks


The simplest diagram is (from Andrew Tannenbaum's famous book) two connected blocks.

A ->B;



This is essentially an arrow specification with the participating blocks automatically created (using default shape, label, layout and other attributes). A more detailed version can be seen below.

box_A: First;
box_B: Second;
A ->B;



Here each block is first defined to be a *box*, which essentially dictates its shape. After the **box** keyword comes the *name* of the block that is used later to specify the arrow. These names also become the label of the block if you do not specify any other label. In this example, we do - after the colon comes the label. Note that you do not have to name a block, names can be omitted. They are needed only if you want to refer to the block later (such as for an arrow).

Note that you can omit the **box** keyword. In that case the block is created only if it has not been created previously. If it has, only its attributes are changed, such as in the case of block B below. Blocks that are created take their attributes from the *running style*, which can be set by the **use** command. Any block or arrow created after that is impacted by any attributes setting after **use**. There is in fact two running styles, one for blocks and one for arrows. They can be set via the **blocks use** and **arrows use** command. Writing just **use** will impact both.

```

box_A: First;
box_B: Second;
B,_C [fill.color=lightgray];
blocks use line.type=double;
arrows use line.width=3;
C->D;
use color=red;
D->E;

```



If just listed one after the other, the blocks are laid left-to-right. This is called a *row*. You can make blocks go top-to-bottom by inserting them into a *column*, or *col* for short. On the figure below, blocks B and C are part of an unnamed column, where block A and the unnamed col are in a row (that is laid left-to-right). The `space` command is used to insert a bit of extra blank space between block A and the unnamed column.

```

box_A;
space;
col {B,_C;}
A -> B,C;

```




You can also see that it is easy to specify several arrows at once, by listing multiple destinations (or sources) in a comma-separated list.

If you want to enclose two blocks in another block, add them as *content*. In the example below the unnamed col has been replaced to an unnamed box. We use the `boxcol` keyword, this is the same as `box`, but arranges its content in a column instead of a row by default. After the `boxcol` keyword, attributes can be specified in square brackets, similar to other languages of Msc-generator. In particular, we set the corners of the unnamed box to round.

```

box_A;
boxcol [line.corner=round] {
    box_B;
    box_C;
}
A -> B,C;

```




Next, we can add a label to a *container block*, that is, one which has content. This, however, moves blocks B and C downwards messing up the symmetry of the arrows. This is because in a row elements are vertically centered by default, thus block A is vertically centered to the unnamed container block right to it. (Note also that blocks in a column are also centered horizontally by default. Note also, that you can define multiple block with one `box` command.)

```

box_A;
boxcol: Container [line.corner=round] {
    box_B,_C;
}
A -> B,C;

```




To fix the arrow asymmetry problem, we can add a named column (BC) inside the container box and make block A align to that vertically. For this we use the `middle` attribute on the block A. This makes the block's middle to vertically align to the middle of another block. This way we can override the default row alignment between block A and the unnamed container block. Note that you can forward reference BC even before it was defined.

```

box_A [middle=BC];
box: Container [line.corner=round] {
    col_BC {
        box_B,_C;
    }
}
A -> B,C;

```




Adding this extra column is a bit cumbersome just to group blocks for alignment. You can shortcut it by listing a group of blocks directly in the `middle` attribute. Msc-generator will calculate the middle (or top, bottom or any other part) of the blocks combined and will align there.

```

box_A [middle=B+C];
boxcol: Container [line.corner=round] {
    box_B,_C;
}
A -> B,C;

```




Continuing the example, we could add two more blocks with different alignment. The attributes `top` and `bottom` make the top and the bottom of the block to align to that of another block. Note that the blocks A and E are laid out further apart from one another - this is the only way to fulfill the alignment requirements we have specified via the attributes.

```

col {
    box_A [top=B];
    box_E [bottom=D];
}
boxcol: Container [line.corner=round] {
    box_B,_C,_D;
}
A,E -> B,C,D;

```




You can use the `below`, `above`, `leftof` and `rightof` alignment modifiers in front of any block to place it in relation to the previous block. In the example below, `below` refer to the big container `Cont`, so block `N` is placed below it. In turn, `leftof` refers to block `N`, so block `M` is placed left of it.¹

```

col {
    box_A [top=B];
    box_E [bottom=D];
}
boxcol Cont: Container [line.corner=round] {
    box_B,_C,_D;
}
A,E -> B,C,D;

below box_N;
leftof box_M;
M->E;
N->D;

```



In addition to `middle`, `top` and `bottom` there are three more alignment attributes that can be used to govern horizontal alignment. They are called `center`, `left` and `right`. Since I always confuse `middle` (the vertical one) and `center` (the horizontal one), Msc-generator also has four aliases to these two: `xcenter`, `xmiddle`, `ycenter` and `ymiddle` to disambiguate which axis is along with we seek the center position.

As value for these attributes, you can specify not only a block name, but also a part of the block after the at sign @. Furthermore, alignment modifiers can be followed by a block so that they can work no blocks other than the previous ones.


¹ Alignment modifiers also result in side-by-side placement (subject to the respective margins). Thus `below` means that the Y coordinate of the bottom of the previous block (plus its bottom margin) equals to the Y coordinate of the current block (minus its top margin).

```

box_A {
    box_A, _B, _C;
}

box_D [bottom=A@top];
leftof A box_E [middle=A@bottom];

```




It is possible to have not only boxes, but any arbitrary shape for a block. This can be achieved via using the `shape` keyword followed by the name of the shape instead of the `box` keyword. As a shorthand, the asterisk symbol (*) can also be used. A handful of basic shapes are defined in the default design library and can be used. You can also add text without any borders or background. This is useful as explanatory notes. Those can be added with the `text` command as seen below.

```

boxcol_Cont: Container
    [line.corner=round] {
        box_a: box;
        *oval_b: oval;
        shape trapezoid_tr: trapezoid;
    }
    text_t: \-a strange\ncshape [middle=a+b];
    t->tr;

```




You can define your own shapes via the `defshape` command, but defining a separate shape for something simple is often too complicated. To this end, Msc-generator allows you to *join* blocks that are overlapping or touching. Joing blocks will remove any internal lines and keep a single contour. Below you can see 4 blocks without and with an additional `join` command.

```

box_X1 [width=100];
box_X2 [height=100,
    top=X1@bottom, left=X1,
    label.orientation=left, fill.color=lgray];
box_X3 [height=100,
    top=X1@bottom, right=X1,
    label.orientation=right, fill.color=lgray];
box_X4 [width=100,
    left=X1,
    top=X3@bottom,
    label.orientation=upside_down];

join X1+X2+X3+X4 [shadow.offset=3, shadow.blur=3];

```



Note that the fill of the joined blocks is kept, together with any label. You can assign any line, fill or shadow attributes to the joined block, which override the line, fill and shadow


attributes of the original blocks. By default, only line attributes are on, thus the joined block has no fill or shadow.

Another way to define a new block is to *encircle* a set of existing blocks. This can be done by the `around` keyword. In the middle figure of the below example you can see that we use two boxes (`A1` and `A2`) and an oval (`A3`) to encircle various boxes. On the right figure, we have added the `join` command at the last line of the example, which combined the two boxes. Note that even if `A3` does not touch or overlap with `A1` or `A2`, it can still be part of the join.

```
use label="";
col_x {box_a,_b,_c,_d;}
col_y {box_a,_b,_c,_d;}
col_z {box_a,_b,_c,_d;}
col_v {box_a,_b,_c,_d; }

use line.radius=5, line.type=dotted;
around x.a+v.a box_A1;
around z.a+z.c box_A2 [imargin=10];
around x.c+x.d *oval_A3;


join A1+A2+A3 [fill.color=lgray+50,
               draw_before=x];
```



Such `around` blocks have no fill by default. Note that we have used *compound names* when specifying the boxes to encircle. Each of the four column had a name (`x`, `y`, `z` and `v`), but inside the columns the box names were re-used (`a`, `b`, `c` and `d`). To name a specific block you can use the name of the containing block to disambiguate before the block name.

When drawing diagrams, a technique is often used to indicate that we have several instances of a block. This technique can be expressed by prepending the `multi` keyword in front of any block definition.

```
multi 2 box_A: From;
multi box_B: To {
    use margin=5;
    *oval_o: My;
    box: Yours;
    below prev+first *oval: Hers;
}
rightof box_C: Beyond;
A.front->B;
A->B.o;
B.front, B.back -> C;
B->C [color=red];
```



You can apply a number after the `multi` keyword to specify how many copies do you want, the default is 3. You can still refer to any blocks contained within (`B.o` in the example) and two more blocks: `front` and `back` to refer to the first and last in the series of copies.


You can also copy an already existing block. This is a quick way to define larger diagrams. You can give the copy a name, change its attributes or even replace its content (or add to it).

```

box_A: [line.type=double] {
  use fill.color=red+80;
  box_a: My;
  *trapezoid_b: Love;
  *oval_c: is;
  below a+b+c *oval_d: Ethernal;
}

use line.color=red;
copy A as_B [line.type=dashed];
use line.color=;
below A+B copy A as_C {
  use fill.color=red+80;
  replace b box_b: Love [text.color=red];
  add before d box: \bNOT [fill.color=green+80];
}

```




6.2 Arrows and Lines

You can use various arrow symbols (as in other languages) to impact the appearance of the arrow. like `>>` or `=>`. You can even use symbols without arrowheads (`--`, `..`, `++` and `==`) to get a simple line.. In addition, you can specify additional attributes in square brackets after the arrow. These include line, arrowhead and text attributes.

```

box_A;
boxcol_Cont: Container
  [line.corner=round] {
    box_B,_C;
  }
box_D;
A->B<<D;
A<=>C<=>D [arrow.endtype=double_empty];
A>>D [color=red, via=Cont@top];

```




You can add labels to arrows simply using the colon label syntax. You can also add further labels after the arrow, specifying their position along the arrow as percentage of length. By default the labels are placed above the arrow, you can influence it via the `label.pos` attribute. If you indicated the `mark` keyword before the label, a small mark will be added to the arrow. This mark is an arrowhead defaulting to `gvstyle=tick`.

```

box_A;
boxcol_Cont: Container
  [line.corner=round] {
    box_B,_C;
}
box_D;
A->B<<D;
A<=>C<D [arrow.endtype=double_empty];
A<>D [color=red, via=Cont@top]: middle;
10%: head [color=red];
mark 90%: tail [color=red];

```




Now, it might be that you would like to route the long connection below the container block. In such cases you can use the `via` attribute for an arrow. You can name a block and any of `top`, `left`, `right` or `bottom` after it. The arrow will try to go that way.

```

box_A;
boxcol_Cont: Container
  [line.corner=round] {
    box_B,_C;
}
box_D;
A->B<<D;
A<=>C<D [arrow.endtype=double_empty];
A<>D [color=red, via=Cont@bottom];

```




Arrows are routed around blocks in-between the two ends of the arrow. This can be prevented by setting the `allow_arrows` attribute.

```


box_A;
boxcol_Cont: Container
  [line.corner=round,
   allow_arrows=yes] {
    box_B,_C;
}
box_D;
A->B<<D;
A<=>C<D [arrow.endtype=double_empty];
A<>D [color=red];

```




Arrows can start and end not only at blocks. Based on the block layout, you can specify 2D coordinates on the diagram, which can serve as start, end and waypoints for arrows.

To specify a coordinate, use two block sides, separated by a comma and enclosed in parenthesis, for example: `(a@top, b@right)`. You can also add offsets in pixels, like `(a@top+10, b@right-5)`.




Arrows can be laid out using several algorithms. Use the `routing` attribute and choose one between `horizontal`, `vertical`, `grid`, `straight`, `manhattan`, `polygon` or `curvy`. (`Grid` means either horizontal or vertical.) In case of `polygon` and `manhattan`, you can set `line.corner` to influence the appearance and make the vertices smoother via `line.corner=round`. The default is `curvy`.



In case of `vertical`, `horizontal` or `grid` the line will be drawn in the middle of the vertical or horizontal overlap of the two blocks, respectively.


```
row {
    box_a,_b,_c,_d;
    text: ...; box_e;
}
below box_BIG: BIG [xpos=prev];
a,b,c,d,e->BIG [routing=vertical];
```




You can govern where arrows end via *ports*, similar to graphviz. These should be specified after the block name and the `at` symbol (@). Boxes have one port on each side and vertex. You can also use a *compass point* as port. Each of these also influence the direction in which the arrow shall leave or arrive at the block. With the `distance` attribute you can govern, how wide a certain block (or all blocks) are routed around or - in case of starting and ending blocks - how wide an arc the arrow makes at its start or end. The default is 5 pixels for the former and 10 pixels for the latter.

```
box_A [size=40];
space 60;
box_B [size=40];



A@top->B@top: top [color=red, distance=20];
A->B@topleft;
A@bottom->B@bottom: bottom [color=green,
distance=B@20, text.bgcolor=none];
```





Each shape may define its own ports (besides the compass points, which are valid for all shapes). For example, the shape `actor` has one compass point for each hand and leg, plus its head. In the below example, we also use a *direction marker* in addition to the port to explicitly govern the direction the arrow start in or ends from. (The attribute `label.align` determines where on the arrow the label is placed with 0 being its start and 100 being its end (default is 50). `label.pos` on the other hand specifies which side of the arrow the label is (default is above).)



If two arrows would overlap at their start or end, they are automatically shifted a bit to show up as separate arrows. (Note, as well, that the arrowheads follow the curvature of the arrow. The `xmul` attribute simply makes them longer.)




You can use a syntax similar to arrows to specify separator lines. Note that the default routing for lines is `straight`, which does not go around blocks in the way, but just crosses them. If you set `routing=curvy`, the line will route around blocks in the way arrows do, see the red line on the example below.



Finally, a more elaborate version of lines, labels and markers. Note that the plain numbers used in the coordinates are interpreted relative to the inner margin of their container block. (If there is no outer block then to the diagram canvas.) The numbers after the `mark` commands are not followed by a percentage sign, so they are interpreted as pixels from the beginning of the line.

```
box [imargin.left=45] {
  (0,190)->(200,);
  mark 30; mark 50; mark 70; mark 90;
  mark 110; mark 130; mark 150; mark 170;
  180:time [label.pos=below, text.bgcolor=none];

  (10, 0)<- (, 200);
  mark 30; mark 50; mark 70; mark 90;
  mark 110; mark 130; mark 150; mark 170;
  180:effort [label.pos=left, text.bgcolor=none];
}
```



7 Usage Reference

In this section we explain details of the GUIs. Throughout the chapter we will talk about the MFC and CLI GUIs mixed.

7.1 Multiple Chart Types

From version 5.0 Msc-generator supports multiple chart types. Each chart type has its own textual language, which are similar to some extent (use attributes in square brackets, options and curly braces to structure the chart).


Both GUIs supports color syntax highlighting, hinting, auto-complete, small indent, element controls and tracking for all (most) languages. The GUIs ask what type of chart do you want to create every time it is started or when a new chart is created.

Each language has a set of associated file extensions. By default, the file extension is used to determine the type of chart when a file is opened in the GUI or processed on the command line. This can be overridden on the command line with `-S`. There is always a *primary* extension for each language, this is used to name the language (on the command line) and also for chart designs, see below.

The currently supported languages and their extensions is listed below (the first extension is the primary one).

Chart type	Extensions	Comment
Signalling Chart	.signalling, .msc	This chart type is the original chart type of Msc-generator.
Graphviz Graph Block Diagram (experimental)	.gv, .graph, .dot .block	This language is the superset of the DOT language. This language describes block diagrams for architecture or software stack figures.

The first panel on the Home tab of the MFC GUI displays the current file type and the controls relevant for that given chart type. For example, on the picture below, the controls for graphs are shown. Compared to signalling charts there are two differences. First, you can also select the graphviz layout algorithm to apply. This is equivalent to using `layout=<layout>`; inside a graph. Second, you can collapse and expand all cluster subgraphs with one click. (Applicable only when the ‘dot’ layout algorithm is used.)



In the CLI GUI the design and layout can be selected in **Settings**.

7.2 Design Library

For each of its languages Msc-generator comes with a bunch of Chart Designs (and Shapes, see Section 8.9 [Chart Designs], page 91, and Section 9.2.6 [Entity Shapes], page 104). The designs available can be viewed in the design selector combo box on the ribbon (MFC GUI) or in **Settings** (CLI GUI).

The files describing the chart designs (and shapes) are of the same language that the language of the chart they provide designs for. These files shall have the primary extension of the chart type.

Note that you should define only designs and shapes in design libraries (for languages that support them). Colors, styles and procedures shall be defined as part of a design. If you want to define one color, style or procedure to be available without applying a design, add them to the `plain` design. Any styles, colors or procedures defined in a design library outside a design will probably be unavailable later - however, the behaviour is undefined and Msc-generator may generate errors for this in the future. Similar, avoid commands in the design library that actually generate chart elements.

On Windows the file describing the default designs and shapes are installed besides the executable. On Linux and the Mac these files are incorporated into the executable so that `msc-gen` can be a standalone file. Specifically on Windows at startup for each supported chart type Msc-generator looks for a file called `designlib.xxx` in the directory where the executable is located (`xxx` is the primary extension of the chart type). If not found, the file `original_designlib.xxx` is searched (of which a default one is placed there by the installer)¹. If any found, the content is parsed as regular chart text before any user chart text.

You are free to create and modify your own design files². Designs (and shapes) defined in these will be added to the list of available designs in the GUI and can be referred to from files processed via the command line. On Linux and the Mac place your design files to the `.msc-genrc/` folder, that will be searched for under your HOME directory or your folder of preference specified by the `MSC_GEN_RC` environment variable. On Windows place these to the roaming appdata folder, which is under `Users\<user name>\AppData\Roaming\Msc-generator` on Windows 7 and 8. Any file in the above directories with the primary extension of a supported chart type will be read by Msc-generator in no specified order before processing your chart. All other files are silently ignored.

7.3 External Editor

Although there is a built-in editor in Msc-generator GUIs, you can also use an external text editor of your choice with the MFC GUI. When you press `Ctrl+E` or click on ‘Text Editors... | External Editor...’ button on the ribbon, an external text editor is started, where you can edit the chart description. If you perform save in the text editor, the chart drawing is updated, so you can follow your changes visually. Also, if there were errors or warnings, they are displayed in the usual manner. If you select an error, Msc-generator will instruct the external editor to jump to the location of the error (if the external editor supports this functionality.)

During the time you are working with an external editor, the built-in text editor becomes read-only. You can exit the external editor any time to return to the built-in one. By

¹ This mechanism was provided to enable the user to (re)define chart designs and is retained for backwards compatibility only. (The idea was that `original_designlib.xxx` is overwritten, when a newer version of Msc-generator is installed, whereas `designlib.xxx` is not. The current recommended practice is to add your own designs to new files in the roaming AppData folder, see below.

² However, please avoid any construct in design, which result in visual elements. Also try not to create files that result in warnings, errors.

pressing Ctrl+E or clicking on the ‘Text Editors...|External Editor...’ ribbon button again, Msc-generator attempts to close the external editor (which will probably prompt you to save outstanding changes).

You can select the text editor to start in **Preferences|External Editor**. You can select between the Windows Notepad, Notepad++ or any editor of your preference. The author finds Notepad++ a very good editor, so I included specific support³.

Since version 4.6 Msc-generator supports unicde characters via both UTF-8 and UTF-16 input, thus any of these can be used by the external editor to save the file. See Section 7.12 [International Text Support], page 67, for more.

7.4 Internal Editor

Since version 2.3 (MFC GUI) or 7.0 (CLI GUI) Msc-generator has an internal text editor integrated with the GUI. Via tracking mode, you can quickly jump to a particular element in the text, but can also quickly see, what is the result of a specific line.

In the MFC GUI there is a separate Ribbon category governing options for the internal editor. On its first pane you can toggle if you want the current line highlighted and/or line numbers displayed. The other features of the internal editor are explained below. If you accidentally close the internal editor, use the ‘Text Editors...’ button on the Edit pane of the Home category to turn it back on.



The CLI GUI supports the same features, but has less customization options (in **Settings**) for simplicity.

7.4.1 Smart Indent

The internal editor supports automatic indentation for TAB, RETURN and BACKSPACE keys and braces. TAB and Shift+TAB works also with selections as in most programming editors. Selecting a block of text (or using Ctrl+A to select the entire file) and pressing TAB or Shift+TAB has the same effect as if applied to each line separately.

In addition Msc-generator detects the beginning of multi-line labels and attribute lists and can align all subsequent lines of the label or attribute list to the first.

The parameters of automatic indentation can be set on the **Indent** panel of the **Internal Editor** ribbon category. Checking the **Auto Smart Indent** check box makes the tool automatically indent after presing RETURN, BACKSPACE, opening or closing braces and square brackets. This way indentation is automatic as you type.

The **TAB auto-indents** checkbox governs what happens when you press the TAB key. If it is on, the whole line of the cursor (or the whole selection) is auto-indented. If it is off, simply a TAB character is inserted or the whole selection indednted uniformly by the TAB size, set below.

³ You can download Notepad++ from <http://notepad-plus.sourceforge.net/>

The amount of space added by smart indent can be set in the next four edits. **instructions** govern how much instructions inside braces are indented; **attributes** are used to indent attribute lists; **blocks** govern, how a subsequent block in a parallel block, box or pipe series is indented, while **inside** governs, how anything else inside an instruction is indented, such as options in an option list, entities in an entity list, etc.

Finally, the last two check-boxes govern, if special indentation for labels and attribute list are applied or not. See the below figure for explanation.

```

box a->b { #The 'instructions' value
    a->b; #governs the indent of the
}; #a->b line inside the box here.

compress=yes, #The 'inside' value is
angle=2, #used in general to set
hscale=auto; #indent inside an instruction

a->b: Label #Special indent for lines
    in two lines, #indents subsequent lines of
    or three.; #labels to the first line.
a->b: With no special
    indent it gets the
    'inside' indent; #..which is (2) here

a->b [arrow.type=dot, #Special indent for attributes
    line.width=2]; #do the same for attributes.
a->b [
    angle=0 #If the first line is in a separate
    ]; #line from the aquare bracket,
        #the 'attributes' value is used (2)

{
    a->b;
} #finally, subsequent blocks
{ #in a parallel block series
    a->b; #get indented by the 'block'
}; #value (=1 here).

```

7.4.2 Color Syntax Highlighting

The internal editor also supports Color Syntax Highlighting. On the **Internal Editor** pane you can select one of four color schemes. There are four pre-defined schemes: Minimal, Standard, Colorful and Error oriented. The first three applies increasing amount of color, while the last is a minimalist scheme but with potential errors heavily highlighted⁴. At the moment you can not customize individual colors in the schemes. The examples in this document were colored using the ‘Standard’ color scheme.

In the preferences it is also possible to select to underline parse error locations (see **Hint Errors as You Type**). In this case you get instant feedback on syntax problems. Finally, it is also possible to request error messages for any error that has been underlined in the internal editor (see **Hint Errors in Window**). These explanatory messages appear in the same window as compilation errors, but they are prefixed with ‘Hint’. If the error they refer to is corrected, they disappear.


Note that during text edit Msc-generator does not perform a full parsing of the text to enhance performance. For example, correctness of attribute names and values is not verified, merely syntax.

⁴ We note here that all four schemes underline entities at their first use. This is to help you avoid a mis-typed entity name.


7.4.3 Typing Hints and Autocompletion

When turned on, the internal editor can also provide suggestions on how to complete the phrase you started typing, when pressing Ctrl+Space. You can use the up/down arrow keys to select between the offered alternatives and press enter or TAB to select it. Alternatively, you can continue typing the keyword or hit any non-alphanumeric character, which will automatically select the highlighted hint and continue after.

The hints provided are associated with a small icon showing the type of the symbol. On the example below, an entity name ('a'), an option name, a keyword, a design name and a style name is shown.



Various attribute values offer a graphic representation to ease selection.⁵ The items in italics do not represent actual text to be inserted into the chart, so you cannot select them. They are more like descriptions of what you can write there.



On the Internal Editor category of the ribbon you can control how much suggestions you will get and how they are displayed. You can turn hints entirely off. The two bottom checkboxes in the Auto Completion panel govern if the list of hints is grouped along dots (to reduce the length of the list) and if hints that are not matching what you have typed so far shall be removed or not. If grouping is on, attributes starting with the same text, such as `line.color` and `line.width` appear as a combined entry as `line.*`. Pressing the dot '.' key will automatically auto-complete the common part. If filtering is turned on, only those hints are displayed which begin the same as the word under the cursor. If you continue typing, the list is narrowed by every character. If filtering is off all values valid at the location of the cursor are shown.

Msc-generator can provide you hints even without pressing the Ctrl-Space. In the subsequent panel you can govern in what language contexts do you want to receive such automatic hints. In general it is best to experiment with these settings and see what you like.


⁵ Shapes defined in the current file do not appear with a thumbnail representation. Only shapes defined in design libraries do.

If you press **Ctrl+Space** and there are meaningful suggestions, the hint box pops up even if automatic hints are turned off. If there is only one possible way to finish what you have started typing, that ending is automatically inserted (word auto-completion).

7.5 Example Library

Msc-generator languages have quite many features and are not easy to learn. Hints provided by **Ctrl+Space** helps you to explore what attributes an element has and what potential values can these attributes take.

In addition to hints, Msc-generator contains an interactive example library demonstrating the features of each of its languages⁶. Just press the **Examples** button on the Chart panel of the Home category of the ribbon. A dialog box appears (close with the Esc key) listing the examples, the chart text for the selected example and the resulting chart.



The example list can be narrowed by keyword search, just start typing the keyword you are interested in.

The examples are editable. When you make a change to an example, the result of it is immediately reflected in the chart to the right. This allows you to quickly experiment with the language without changing the chart you work on. Any changes you make to the examples will be forgotten the moment you exit Msc-generator.

7.6 Renaming elements

Right-clicking an entity name (for signalling charts), a block name (for block diagrams) or a node name (for graphs) brings up the context menu for the internal editor. In that there is a **Rename...** menu item, which replaces all occurrence of this name in the source file to something else. Note that this does not replace the *label* of the entity (if you have specified

⁶ For the dot language, we only demonstrate the Msc-generator additions.

one), block or node (which will be shown in the compiled graphics), but how it is referred to from the source file. The resulting chart, diagram or graph remains visually the same after such a rename. (Unless you use some entity names as label.) However, it enables you to quickly change how you refer to these elements.


In case of signalling charts the value of the `relative` entity attribute (Section 9.2.3 [Entity Attributes], page 101) is the name of another entity. These attribute values are not renamed - you need to manually adjust the name.

In case of block diagrams, names are hierarchical. That is a `box B` inside a `box A` is called `A.B` in its full name - even if it can be referenced as just `B` from within `A`. (More than two levels of nesting is possible.) When you click on a name with multiple levels (such as `A.B`) only one part of it will be renamed - based on where you click. Only the actual block specified by the clicked name is renamed (even if there are blocks with the same name elsewhere), but that block is renamed at every place it is mentioned (even in attribute names and alignment modifiers, Section 11.4.2 [Alignment modifiers], page 171).

It is possible to perform the replace only in one part of the source text. Simply select the desired part and click on the element to rename (may be outside the selection). Only the mentions of this entity inside the selection will be replaced. This allows you switch from using one element to another (spilt a single element into two). Using the name of an already existing element will make the renamed ones refer to the already existing element (merge two elements into one).

7.7 Options

In the MFC GUI selecting the **Preferences** category on the ribbon allows you to set a few options of Msc-generator.



In the first category you can specify what is the text of the chart that pops up when a new chart of the current type is created. Just press the button and the current text will become the default for the current chart type. You can place your frequently used constructs here to be readily available when you start a new chart; or just delete everything here to start real empty.

Under ‘Options’ you can set a few compilation options. When pedantic is set Msc-generator enforces stricter language interpretation. For signalling chart, it means generating a warning if an entity is not declared explicitly before use. For graphs, it means generating a warning for each use of directed edges in undirected graphs and vice versa (graphviz does not allow such mixing), but when turned off, mixing directed and undirected edges in a graph becomes possible. For Block Diagrams, it prevents auto-generating an block when mentioned without a ‘box’ or similar keyword. This allows forward referencing blocks in arrow definitions (by avoiding their creation instead of referencing a later definition). Turning the second option on will supress the generation of warning messages altogether (including the ones generated due to the pedantic option). Checking the third option will

show the full path of the filename in error messages. This is useful when using `include` and suspecting the precise identity of the included file with an error. When ‘Ask char type’ is set, Msc-generator displays the list of available chart types at startup and when creating an empty document and you can select one. If unchecked, the last file type is used always. When ‘Technical Info’ is ticked, a summary of the compilation is shown among the errors and during compilation colored progress indicators show the various stages of compilation. This is mostly for debugging. In the third column you can control automatic behaviour. ‘Auto paste clipboard’ makes Msc-generator to automatically paste the content of the clipboard, whenever an alt-text embedded chart is copied to it. (Except when you have unsaved changes in the currently edited chart. In that case you can use the ‘Paste Entire Chart’ button in the ‘Home’ ribbon category.) Checking ‘Auto Save’ will save the chart text every time you make changes (except for OLE embedded charts). Finally, ‘Instant compilation’ will recompile the chart every time you make changes.

On the ‘Appearance’ panel you can first select the color of the tracking overlay (what flashes when you click on a chart element). ‘Show Page breaks’ governs if a dashed line is drawn to show where page breaks are when watching all of the pages. See Section 9.12 [Multiple Pages], page 142, for more information. Lastly you can also set the maximum zoom factor selected by the *Overview* and *Fit to Width* automatic zoom modes. You may want to increase this for very large screens. See Section 3.1.2 [Zooming], page 6.

The ‘Mscgen compatibility’ panel is applicable only to signalling charts and governs how Msc-generator switches to backwards compatibility mode with mscgen. See Section 9.15 [Mscgen Backwards Compatibility], page 152, for more.

On the last panel you can specify which external text editor to use. You can select any editor using the first option. In this case you have to give a command-line to start the editor and one to invoke to jump to a certain line by pressing the button to the right. The latter can be omitted if the editor does not provide a command line option to jump to a certain location in an existing editor window. Use ‘%n’ for the filename and ‘%l’ for the line number; these will be replaced to the actual filename and linenumber at invocation.

7.8 Working with Multi-page Charts


Msc-generator supports multi-page charts. These may be useful when you want, e.g., to print a long signalling chart. For signalling charts you can manually start a new page by typing the `newpage;` command or let Msc-generator automatically paginate for you. See more in Section 9.12 [Multiple Pages], page 142. In case of graphs, each graph is laid out to a separate page. Block Diagrams have no pagination at present.

You can select on the ribbon/menu bar which page to view. This setting is also saved with embedded charts, and of course only the selected page is shown in the container document. You can also select to view all pages. When viewing all pages, Msc-generator marks page breaks with a dashed or dotted line for manual and automatic page breaks, respectively, and also prints page numbers to the left. This behaviour can be turned off in the preferences. (See Section 7.7 [Options], page 62.)

When you embed a multi-page chart (either as OLE object or via alt-text embedding) to a document, you can chose to show only one page of it. The text of the whole chart will always be saved in the embedding, so later you can easily change the page displayed. Simply select the page before copying to the clipboard (or saving an already embedded chart).

In the MFC GUI the last pane on the Home category of the ribbon governs automatic pagination. The first checkbox turns it on. The paper size can be selected in ‘Print|Print Setup...’, whereas margins, page alignment and scaling can be selected in Print Preview. Ticking the second checkbox will result in a heading to be displayed for the active entities at the top of every page.


In the MFC GUI you can preview a multi-page chart before printing using the Print Preview option from the main menu. It behaves similar to Print Preview in other programs.



Here you can print or export to PDF⁷. You can also set Automatic Pagination and Auto Headings from here, but you can also set, how the chart pages are sized; how they are aligned within the physical pages and also what margins to apply. The former option can take *Fit width* and *Fit page*⁸. In the former case the scaling factor is selected to fit the width of the chart to the page, this is perhaps the most useful selection to print long charts combined with automatic pagination. In the latter case, the scaling is set uniformly for all pages to fit the longest one. With automatic pagination this is in effect equivalent to *Fit width*.


7.9 Scaling Options

If the chart is exported to a PNG image then after selecting the filename an additional dialog box appears where you can set scaling options. In all but the last option the original aspect ratio of the chart is kept. After the ‘No scaling’ option the native size of the chart is shown.



⁷ When you export a multi-page chart from Print Preview, Msc-generator automatically exports to a multi-page PDF file. When you use the Export option from the main menu, Msc-generator asks what you want.

⁸ They correspond to the `-s=width` and `-s=auto` command-line options, respectively.



7.10 Advanced OLE Considerations


7.10.1 Graphics of Embedded Charts

The technology used to embed charts into other document, called OLE, has certain limitations on graphics⁹. To work around these, Msc-generator employs a few simplifications.

- Due to clipping limitations, certain arrowheads, like ‘line’ draw differently on signalling charts.
- Due to font limitations, the label of slanted arrows on signalling charts are drawn with limited resolution and looks somewhat different than non-slanted text.
- Due to missing gradient fill support, gradient fills and shadows are approximated. At large magnification this becomes visible.
- Due to the limited size of the coordinate space, placement of elements in very large charts appear imprecise.
- Due to lack of transparency support, translucent areas (such as pipes) are drawn on a bitmap, a *fallback image* and then inserted, see below.

If a chart contains a lot of fallback images, the size of the embedded object can become large, several megabytes for a chart. To control the size of the embedded chart and eventually that of your container document, for embedded charts a new category ("Embedded Object") appears on the ribbon, allowing you to adjust the quality of the fallback images.

Note that this issue is fixed in newer versions of Microsoft Office, which are able to compress the images in embedded objects.



When this category is selected, Msc-generator shows the chart as it will appear in the container document. Fallback image locations are briefly highlighted when switching to this category. The ribbon category shows how large the embedded object will be (if not

⁹ Only drawing operations permitted in old-style Windows Metafiles (developed in 16-bit Windows times) are permitted by design.

compressed by the container applications) and what percentage of the chart is drawn on fallback images (if any). There is a slider allowing you to set the resolution of the fallback images. You can observe the resulting image size and visual quality immediately.

7.10.2 Linking

You can also choose to insert a Link to a copied chart instead of embedding it into the document. In this case updating the source chart will get reflected in the document, as well. You can also insert a link to only a page of a chart, by copying that page to the clipboard via the drop-down menu of the *Copy Entire Chart* button.

Note, however, that you cannot insert a link to a chart that is not saved on disk, but is yet ‘Untitled’. In addition, not all container applications implement the full range of linking features.

- LibreOffice and OpenOffice do not allow links to be inserted into documents. You can only embed charts in their documents.
- Microsoft PowerPoint allows links to be inserted into a slide, but does not allow other programs to link to a chart embedded in a slidepack. This includes the case when you want to insert a link into a slidepack that points to a chart embedded in the very same slidepack.
- Microsoft Excel implements full linking features, that is it allows you to insert links into worksheets, but also allows you to insert a link pointing to a chart embedded in a worksheet into other documents (or the same worksheet). You can even insert links that point to a single page of a chart embedded in a worksheet. (You can do this by opening the embedded object in Msc-generator and select ‘Copy Page #1’, and use Paste Special to insert a link.)
- Microsoft Word allows you to insert link to charts that are saved in files or are embedded in some other container (such as an Excel worksheet). It also allows others (including Word itself) to link to a full chart embedded in a Word document, but does not allow linking to a page of a chart embedded in a word document. If you invoke ‘Copy Entire Chart’ or ‘Copy Page #x’ from within Msc-generator for a chart embedded in a Word document, the link will not work. However, if you copy the chart to the clipboard from Word (and then you can copy all of it) then if you insert a link via Paste Special, you will get a valid link.

There is a suspected bug in Word 2003 that fails linking to a single page of a chart embedded in a Word document.¹⁰

7.11 Autosave and Recovery

Since version 5.0 the MFC GUI supports the Restart Manager introduced to Windows in Vista. That is, if Msc-generator crashes (or needs to be stopped due to an installation), its status is saved and is automatically restarted.

¹⁰ To link to a full chart embedded in Word, make sure you place chart to the clipboard using Word and not using Msc-generator.


The documents you are working on are autosaved.¹¹ Clicking on the ‘Autosave and Recovery’ category on the ribbon opens a list of autosaved files on the side.¹² Double-clicking a file will recover it either by overwriting the current chart text or by inserting it at the cursor (you can select after double-clicking). You are also offered the option to delete the autosaved file after recovery. On the ribbon you can filter the list of autosaved file to the current chart type and make Msc-generator to open the list of autosaved files at startup. You can also delete all autosaved files at once.¹³

Note that in the rare event of an Msc-generator crash¹⁴ the restart function kicks in only if the application was open for more than 60 seconds (to prevent restart loops). Your documents are autosaved even before this time, nevertheless.

The CLI GUI has a simple autosave and recovery mechanism. The chart text (and context) is saved roughly every second or if you exit the program without saving your work. When Msc-generator is opened next time, you can recover the last autosave. Note that if you start Msc-generator with a filename, your autosave will be overwritten and lost. Also note that multiple running instances of the CLI GUI perform autosave to the same location overwriting each other constantly.

7.12 International Text Support

Msc-generator supports international text via both UTF-8 and UTF-16¹⁵. The type of encoding is determined automatically when the file is loaded. If it is ASCII or valid UTF-8, it is treated as so. Else it is treated as UTF-16. The Windows GUI displays the detected encoding on the status bar. You can change it by double clicking the indicator on the status bar - next time you save the file, it will be saved using the encoding displayed there. Below is an example of a simple chart shown in the language tutorial in my native language, Hungarian.



¹¹ The autosave location is C:\Users\<username>\AppData\Local, you can manually recover autosaved files from there.

¹² Any autosaved version of the currently open document is omitted from the list.

¹³ If you filter the list of autosaved files to the current chart type then clicking the ‘Delete autosaved files currently showing’ button will not delete the autosaved files of other chart types.

¹⁴ The graphviz library crashes on some chart input. This is something I cannot easily fix. Hence the addition of the autosave and recovery functions.

¹⁵ UTF-8 is a superset of ASCII, where international characters are encoded in multiple bytes all larger than 127 (which is the largest ASCII character). In contrast, in UTF-16 all characters are encoded on two bytes, ASCII and international characters alike. The UTF-16 file usually starts by the characters 0xff and 0xfe and their order determines in which order the two bytes of a character comes. UTF-8 is widely used in Linux and other Unix systems, whereas Windows has adopted UTF-16 as the system encoding. Msc-generator supports both file formats and both byte orders of UTF-16.

The CLI GUI only supports UTF-8.

Note that not all font faces contain characters for all possible unicode codepoints. This has to be considered when selecting a font¹⁶. If you mix characters of different languages, you need to ensure that they are displayed via a font that has coverage for all those characters.

In labels, you can use the `\$` escape followed by exactly four hexadecimal digits, to insert any unicode character to a label, e.g., `\$00a9` inserts the copyright symbol.

Note that if you plan to use non-ASCII characters in the chart text you need to select a font that can display them not only in the chart, but also for the internal editor. See Section 7.14 [Fonts], page 74.

7.13 Command-Line Referece

The syntax of the command-line version is the same on all platforms¹⁷.

Note that the command line syntax below is a superset of the command-line options of the mscgen tool. This means that by renaming `msc-gen` to `mscgen` you can use Msc-generator's extra features and rich language in every tool that is integrated with mscgen. These tools include Doxygen, Sphinx and Msctexen¹⁸. For more on the compatibility of the languages see Section 9.15 [Mscgen Backwards Compatibility], page 152.

```
Usage: msc-gen [OPTIONS] infile
       msc-gen --gui [OPTIONS] [infile]
       msc-gen --tei-mode -S <lang> [OPTIONS]
       msc-gen -l
       msc-gen --help
       msc-gen --version
```

Use `msc-gen --help` for detailed command-line summary.

7.13.1 Label Maps

When you specify `lmap` as output file format, Msc-generator creates a text file with one line for each text label in the chart¹⁹ (and no graphics output). The default extension will be `.map`. The lines in the output file contains the followin information separated by space.

`<type> <page> <x1> <y1> <x2> <y2> <first line>`

The `type` character tells, what chart element contained this label. The following characters are possible

- A Arrow, including block arrows (including boxes collapsed to arrows)
- E Entity heading. Each appeareance of the entities will result in one line.

¹⁶ In signalling charts you can select a font via the `text.font.face` chart option or attribute, in graphs use the `fontname` attribute, whereas in Block Diagrams you need use `text.font.face=<fontname>;`. On Linux/Mac, where `fontconfig` is used to select the fonts, specifying the `text.font.lang` attribute along with the font face (see Section 7.14 [Fonts], page 74), will try to select a font having glyphs for the specified language.

¹⁷ The only two exceptions are in how pathnames are written on Windows and Linux/Mac and where the design libraries (e.g., `designlib.signalling`) will be searched.

¹⁸ This is thanks to the original mscgen author Michael McTernan and many others

¹⁹ Note that box tags are not included in the label map.

B	Box that has content (unless collapsed)
b	Box that contains just a label (or collapsed)
P	Pipes
V	Verticals (all forms, including boxes, block arrows, ranges, braces and brackets)
D	Divides, titles, subtitles, discontinuity lines and plain text (like [label="aaa"] ;).
N	Floating notes
C	Comments (on the side or at the end)

The second item `page` gives which page the label is on. One label is mentioned only once even if it spans multiple pages.

The following four numbers give the upper left and lower right corner of the bounding box of the label (and not the corresponding element). It is given in pixels for bitmap output and in logical coordinates matching the logical size of the output image for vector graphics output. The coordinates are relative to the top left corner of the page origin and are rounded to integers for ease of use. So if you run Msc-generator twice, once with a graphics output format and once with label map output (leaving all scaling and other switches the same), the coordinates of the label map shall match the graphics output perfectly.

Finally the line ends with the first line of the label (which may contain spaces), potentially with the number prepended (if any) in the number format used in the chart (e.g., roman numbers). Note that the coordinates specify the bounding box of the entire label, not just the first line given here.

In this release of Msc-generator, label maps are not emitted for graphs.

7.13.2 Coloring Input Files

Using `-T csh` produces a colorized version of the input file text. Specifically, the output is a text file according to the signalling chart language (irrespective what the input language was). The output contains a single colon-label divider using the formatting escapes described in Section 8.2 [Text Formatting], page 76. Running `msc-gen` on the resulting text output again (interpreting it as a signalling chart), will produce a colorized version of the original input text as graphical output. This is how the color text version of the examples in this document were produced.

7.13.3 Embedding Charts

You can embed chart text in PNG files by using the `-e`. Then you can use this PNG file as an input file (and open it in the GUIs, too). You can extract the source text from a PNG file by using `-T src`.

You can also embed charts in PPT files, by using `-T embed` and specifying an existing `pptx` file as output. This will append a slide and place the chart in the middle of it. Using `-T embed:<slide>` will place the chart on an existing slide, identified by its number. Numbers start from 1 and you get an error if the slideshow does not have as many slides. Finally, specifying `-T embed:<slide>:<chart>` will *replace* an existing chart on the given slide. Charts on a slide are numbered from one in Z-order, where chart #1 is the one on the back.

You can extract the source text from a PPT file, by appending the slide and chart number after the input filename, like `msc-gen my.pptx:2:1`. Adding `-T src` will save the chart text, but other output graphics formats are also possible. The following example updates the chart with a given design.

```
msc-gen my.pptx:2:1 -T embed:2:1 -o my.pptx --design=round_green
```

In order to see what slides contain what charts, use `msc-gen --list-embedded my.pptx`.

7.13.4 Text Editor Integration

The options here are intended for integration with text editors. They may change in the future. Please contact the author if you want to integrate with a text editor.

Using `--tei-mode` will put msc-gen into *Text Editor Integration mode*, where it expects a series of input files on its standard input and produces color syntax highlighting, quick error, text indentation or hints/autocomplete information on the standard output. These are produced via a quick parsing of the input text and are intended to be fast enough to be run for every keystroke pressed. (At least for input files of reasonable sizes, say 64K.) Parsing in this mode is limited, thus not all errors are caught, but most are. Msc-gen reads and parses all the design libraries at startup and remains running for several input files - again to save time. (Reading design files can be avoided via option `--nodeligns`, as usual.)

On the standard input msc-gen expects a text line with a series of space separated letters, each optionally followed by a number, then a new line character. After this shall come the text of an input chart, terminated by a zero byte. The letter in the first line determines what kind of output is expected on the standard output (see below). The output is also terminated by a zero byte. All the above can then be repeated until the standard input is closed. In that case, msc-gen closes the standard output and exits with success.

Msc-gen uses the standard error to produce critical error messages (ones that do not pertain to the correctness of input text, such as bad combination of options, etc.).

The type of the chart has to be specified via option `-S` and cannot be changed later. (Similar, if the design libraries change, msc-gen need to be restarted.) The text of the input can only be ASCII or UTF-8, but not UTF-16. Input or output filenames are invalid, so are the options `-i` and `-o`. A host of other options to manipulate graphics output (sizing, pagination, etc.) are simply ignored.

The following table summarizes, what letters can be used in the first line of the input and what effect they have. If you specify more than one letters in the first line separated by spaces, all output for them will be returned separated by an empty line. Thus if you want to get both the coloring information and the quick errors in one pass, use `C E` in the first line. The order of the letters in the first line do not matter, output is ordered as below in the table. You can specify one letter several times, but it has no effect. In case the letter requires parameters, that of the last mention will take effect.

<code>C</code>	In response to <code>C</code> , msc-gen emits a description of how to color the input. Each line describes a particular contiguous area of the input file colored the same way in the format below.
----------------	---

```
<first_pos>-<last_pos>: <color_id> <red> <green> <blue> <bold>
<italics> <underline>
```

The first two numbers mark the first and the last position of the color range starting from 1. If the color range consists of a single character, the two values

are equal. In case of UTF-8 encoding, multi-byte characters are counted as one thus the position is not a byte index. The different line endings (CR, LF or CRLF) are all understood, with CRLF counted as one character.

The `color_id` is a number between 0 and 36 denoting what kind of color do we talk about. The table further below summarizes the values. The rest of the numbers on the line give actual colors and font flags (associated with the default color scheme).

`D<pos>,<insert>`

Similar to `C`, this emits coloring information, but only a delta compared to a previous run. This may be beneficial to reduce the amount of coloring results generated by msc-gen. Also, if the text editor supports incremental updates in coloring, this may also reduce the work of updating coloring after simple changes.

The delta can only be compared if a single (but perhaps multi-character) insertion or deletion were made to the input file. `<pos>` represents the caret position *after* which the characters were inserted or deleted, while `<insert>` tells us how many were inserted (negative number for deletes). If you have specified the `N` letter, you can omit the chart text entirely from after the first line for deletions. For insertions, you just need to specify the characters inserted (which have to be as many) as the `<insert>` you specified. By default (no `N` added), msc-gen still expects the full chart text on its input. If the chart text given to msc-gen was modified in a different way compared to the original than the parameters of the `D` command indicate the resulting delta will likely be wrong. The returned coloring info will contain `color_id` of 0 for texts that were colored before, but now coloring has to be removed.

Note that the delta computation algorithm assumes that the text editor maintains coloring before `<pos>` and shifts coloring information after the insertion. That is, if a single character was deleted at position 5, a blue colored character at position 3 will remain blue (and hence this can be omitted from the delta) and a red character originally at position 7 (now at position 6) will remain red at its new position. This is how any sane text editor having text formatting shall work.

You can only specify one of `C` or `D` at the line, the latter takes effect.

`E`

Msc-gen will emit quick errors, for highlighting during editing. (Not all errors are captured, such as attributes not applying, etc. This is to speed up the process.) The output contains a character range `<first_pos>-<last_pos>` followed by a textual description. The text editor may underline these locations and assign balloon tooltips for them.

`H<pos>`

If the letter in the first line is `H` a nonnegative number shall be appended (with no space). In this case msc-gen will return the possible list of auto-completions, when the cursor is at the number specified. This position is a caret position, value zero is before the first character in the file, value 1 is after it. In case of UTF-8 input each character counts as 1, even if multi-byte. Thus if the first character takes two bytes, then the position 1 will start at byte #2. The list returned consists of one line per hint and the hints are neither grouped nor

filtered (see Section 7.4.3 [Typing Hints and Autocompletion], page 60). There are seven fields separated by the \x01 character.

1. The hint text. This is to be shown to the user.
2. The string to insert. If empty then use the hint text above.
3. 1 if the hint is selectable (can actually be inserted to the chart) and 0, if it is only explanatory.
4. The RGB color of the hint: three, comma separated integers between [0..255].
5. 1 if the hint text shall be bold, 0 otherwise.
6. 1 if the hint text shall be italics, 0 otherwise.
7. The description of the hint in english to be used as a tooltip. Newline characters are replaced to character \x02 in order to avoid confusion with newlines separating the hints.

If the cursor is at a location where the user can type a new identifier (and not just one of the values among the hints), the first hint will have its hint text equal to the asterisk ('*'). This shall be honoured by not expanding to one of the hints when the user presses the space or a non alphanumeric character.

I<pos> In this case the position is also a character position as for H, but smart indentation information is returned. Text editors can use it when the user presses TAB or RETURN keys, or when the user has typed a brace or square bracket to the beginning of a line. Msc-gen will return four numbers separated by a space and terminated by a newline.

- The character index of the beginning of the line containing <pos>.
- The byte index of the first character of the line containing <pos>. In case of a pure ASCII input this equals the first number.
- The current number of whitespace (tabs count as a single whitespace) at the beginning of this line. If the line contains only whitespace this number is the total number of characters in that line.
- The suggested number of whitespace at the beginning of the line (which is the complete line if there is only whitespace). Note that for empty lines the correct indentation may change after the first character is typed.

Note that if you would like to indent several lines at once, you need to actually indent the first one (by adding/removing spaces) and then re-parse the resulting file and do it for the second line, etc. This is because the calculated indent depends on lines above.

N If this letter is added to the line, the chart text following the first line can be missing. (It is ignored if present all the way to the terminating zero character.) The text of the previous run is re-used, but parsed again. It is useful, when the text has not been changed, but we want hints at a different location.

Note that C, D, E and H will result in a re-parse of the chart text, so it is a good strategy to send them all in once to be done with a single re-parsing as opposed to send them one-by-one with N, which results in a fresh re-parse again. I, on the other hand can work from a previous parse.

And the code for the `color_id` is as follows.

- 0 Regular text (if a character is not listed, its color defaults to this).
- 1 A keyword, like commands or `parallel`
- 2 A partially typed keyword
- 3 An msc-gen only keyword
- 4 The equal sign
- 5 The semicolon sign
- 6 The colon character
- 7 The comma character
- 8 A symbol, e.g., arrows, box symbols, etc.
- 9 An mscgen symbol, e.g., arrows, box symbols, etc.
- 10 An opening or closing brace { }
- 11 An opening or closing bracket []
- 12 An opening or closing parenthesis ()
- 13 The name of a design
- 14 The name of a style
- 15 The name of an entity
- 16 The name of an entity when first occurs in the file
- 17 The name of a color
- 18 The description of a color other than a name
- 19 The name of a marker
- 20 The name of a marker partially written
- 21 The name of a procedure
- 22 The name of a file to include
- 23 The name of a procedure parameter (\$xxx)
- 24 The name of a chart option
- 25 The name of a chart option partially written
- 26 The name of an mscgen-only chart option
- 27 The name of an attribute
- 28 The name of a attribute partially written
- 29 The name of an mscgen-only attribute
- 30 The value of an attribute
- 31 Emphasized part of the attribute value (usually variables)

32	The text of a label, either after a colon or between quotation marks
33	The text of a label escape, such as '\n'
34	A place where an error was detected.
35	The inverse of error formatting (used only internally)
36	A comment

7.14 Fonts

For signalling charts the fonts used for labels, comments, etc. can be selected in 5 ways (listed in decreasing order of preference).

1. Using the `\f(font name)` and/or the `\l(language code)` text formatting escape sequences²⁰. These can be applied even in the middle of a label.
2. Using the `text.font.face` and/or the `text.font.lang` attributes, in which case they apply to the whole label. The value of these attributes can also be set by styles.
3. Using the `text.font.face` and/or `text.font.lang` chart options. They affect all subsequent labels (until the next closing brace, if any).
4. Using the `-F` command-line option.
5. By setting the `MSCGEN_FONT` and/or `MSCGEN_FONT_LANG` environment variables.

The fonts available are system dependent. On Windows, you can use all the Windows fonts available, but only OpenType and TrueType fonts provide correct alignment. On Linux or Mac, FontConfig is used and you can use fonts installed with it. Use the `fc-list` command to list available fonts. Pick the family name in the list for use in Msc-generator. You can also specify the language you want (see above), this impacts the font selected so that it contains the glyphs needed by the language you specify.

For graphs, the selection is similar, but the name of the attribute is `fontname` and there you need to use the `node [fontname=<face>]`; to impact all subsequent nodes. There is no way to set the language except from the `\l()` text formatting escape.

For Block Diagrams the precedence is similar to that of the Signalling Charts, but instead of the `text.font.face` chart option, you can change the running style via the `use text.font.face=;` command. (And similar for the language.)

On the Windows GUI you can use the last panel of the ‘Internal Editor’ ribbon category to select the font of the internal text editor. The two checkboxes can be used to filter the list of font families offered by the lisbox below them. Note that if you plan to use non-ASCII characters in the chart text you need to select a font that can display them not only in the chart, but also for the internal editor.

Finally, when you start Msc-generator with the `MSCGEN_NOFONT` environment variable set (to any non-empty value other than ‘0’), no fonts will be drawn to the output. The area occupied by any text will be left empty (but included in layout calculation). This is useful for testing.

²⁰ Use the RFC 3066 (<https://www.ietf.org/rfc/rfc3066.txt>) style language names. See a list of examples here (<http://www.i18nguy.com/unicode/language-identifiers.html>).

8 Reference for Common Language Elements

Msc-generator languages all share a common underlying logic and design.

- They all define elements and relations between elements, such as entities and arrows for signalling charts; nodes and edges for graphs; and blocks and arrows for block diagrams.
- These elements and relations usually can have a text label, the format of which (such as bold or italics) can be influenced via *text formatting escapes*.
- These elements and the relations all have *attributes* to govern appearance, which can be specified in square brackets.
- Attributes can be collected in *styles*, which can be applied to elements and relations in a simple way.
- Each element and relation type has a *default style*, which can be changed to change the appearance for all such elements or relations.
- Curly braces can be used to group elements and to isolate a *scope* - any changes to styles or colors are limited to the current scope.
- You can define *designs*, which contain styles, colors and can be applied to the whole chart to change its appearance in one go. These are collected into *design libraries*, which are read at program startup and are available by default.
- You can also include other files in the middle of the any document.
- You can define and later replay *procedures* to avoid frequent copy/paste and reuse frequent constructs instead.

This chapter contains all the language elements that are common to all or most chart types supported by Msc-generator. Variations specific to each language are individually described.

8.1 Labels

Entities, arrows, boxes, pipes, dividers (signalling charts), nodes, edges and clusters (graphs) and blocks (block diagrams) all have a `label` attribute, which specifies the text to be displayed for the element. Each element displays it at a different place, but the syntax to describe a label is the same for all. For entities (signalling charts), nodes and clusters (graphs) the label defaults to the name of the elements, while for the rest it defaults to the empty string. Labels have to be quoted if they contain any character other than letters, numbers, underscores and the dot, or if they start with a dot or number or end with a dot. You can use all character formatting features in labels, see Section 8.2 [Text Formatting], page 76.

To avoid typing `[label="..."]` many times it is possible to specify the label attribute in a simpler way. After the definition of the element, just type a colon (two colons for graphs), the text of the label unquoted and terminate with a semicolon (or opening brace '{' or bracket '['). You can write attributes before or after such a *colon-label*. Thus all lines below result in the same text.


```
a->b [label="This is a label", line.width=2];
a->b: This is a label [line.width=2];
a->b [line.width=2]: This is a label;
```

If the label needs to contain an opening bracket ('['), opening brace ('{'), hashmark ('#') or a semicolon (';') use quotations or precede these characters by a backslash '\'. This is needed since these characters would otherwise signal the end of the label (or the beginning of a comment). If you want a real backspace, just type '\\'.

When using the colon notation, heading and trailing spaces are removed from the label. If these are needed, place the entire label between two quotation marks '\"'².


```
hscale=auto;
a->b: Label with a semicolon(";\") in it
a->b: " Label with a semicolon(\";\") in i
box a---b: Escapes: \{ \[ \; \# and \\. ;
---: Can escape these, too: \] \} \";
: but not needed: ] } ";

```



Labels can span multiple lines. You can insert a line break by adding the '\n' escape sequence. Alternatively you can simply break a label and continue in the next line. In this case leading and trailing whitespace is removed from each line.

```
compress=yes;
a->b: First line
        Second line #comment
        Third line;
b->c: First line
        \ -Smaller text
        And here, too!;
box a---c: \prAll lines
        right aligned... {
box a...c: ... or only \prthe
        second one;
}
```



8.2 Text Formatting

Any text displayed in the chart may contain *formatting escapes*.³ Each formatting escape begins with the backslash '\' character. You can also use the backslash to place special characters into the label. Below is the list of escape sequences available.

¹ This character is often called the *escape character* making an *escape sequence* together with the character it follows.

² In this case there is no need to escape the opening bracket or brace, the hashmark or the semicolon, since the end of the label is clearly indicated by the terminating quotation mark. If, on the other hand you need quotation marks in the label use '\"'. Also, you cannot break the text in multiple lines in the input file, you have to use the '\n' escape to insert line breaks. This *colon-quotation-mark* mode is available only for signalling charts and is provided only for backwards compatibility with mscgen.

³ For graphs only labels specified via the double-colon notation interpret escape sequences this way. If you assign a label via the `label=` attribute, the escape sequences will be interpreted as graphviz escape sequences, see <http://www.graphviz.org/doc/info/attrs.html#k:escString>.

<code>\n</code>	Inserts a line break.
<code>\-</code>	Switches to small font.
<code>\+</code>	Switches to normal (large) font.
<code>\^</code>	Switches to superscript.
<code>_</code>	Switches to subscript.
<code>\b</code>	Toggles bold font.
<code>\B</code>	Sets font to bold.
<code>\i</code>	Toggles italics font.
<code>\I</code>	Sets font to italics.
<code>\u</code>	Toggles font underline.
<code>\U</code>	Sets font to underlined.
<code>\f(font face name)</code>	Changes the font face. Available font face names depend on the operating system you use. (See Section 7.14 [Fonts], page 74.) If you specify no font, just <code>f()</code> , the font used at the beginning of the label is restored.
<code>\l(font language code)</code>	Changes the language of the font. On Windows it is ignored. On FontConfig-based systems (Linux, Mac) you can use an RFC3066-style language specifier, typically a 2-letter ISO language code, such as <code>ja</code> for Japanese or <code>zh</code> for Chinese. This code is used to select the font file which includes the characters of the given language. (See Section 7.14 [Fonts], page 74.) If you specify no code, just <code>l()</code> , the font used at the beginning of the label is restored.
<code>\0..\9</code>	Inserts the specified number of pixels as line spacing below the current line.
<code>\c(color definition)</code>	Changes the color of the text. Color names or direct rgb definitions can both be used, as described in Section 8.5 [Specifying Colors], page 83. No quotation is needed. You can also omit the color and just use <code>\c()</code> , which resets the color back to the one at the beginning of the label.
<code>\s(style name)</code>	Applies the specified style to the text ⁴ . Naturally only the <code>text.*</code> attributes of the style are applied. You can omit the style name and specify only <code>\s()</code> , which resets the entire text format to the one at the beginning of the label ⁵ . See Section 8.6.4 [Styles], page 88, for more information on styles.

⁴ Note that the `\s` formatting escape was used to switch to small font in 1.x versions of Msc-generator (since 2.0 `\-` is used for that). In order to work with old format charts, if the style name is not recognized, Msc-generator will give a warning but fall back to using small font.

⁵ Any formatting escapes strictly at the beginning of a label (up to the first non-formatting escape or literal character) are included in the text format, so if you start a label with '`\b`' then '`\s()`' will restore a bold font. To prevent this use the '`\|`' escape to create an invisible non-formatting character.

\S(*shape|height|fillcolor*)

Inserts a shape into the text as a single character. Shapes can be arbitrary forms described by polygons or bezier curves and are defined using the `defshape` command, see Section 8.10 [Defining Shapes], page 92, for more information. This escape allows drawing arbitrary things everywhere, where a label is permitted. Within the escape sequence, you first need to specify the name of the shape, then optionally its height (default is the font size) and optionally its fill color (default is transparent fill). The latter option takes effect only for shapes having a fill section, see Section 8.10 [Defining Shapes], page 92. You can omit the height (`\S(shape|fillcolor)`) or the fill color (`\S(shape|height)`) or both (`\S(shape)`).

\mu(*num*)

\md(*num*)

\ml(*num*)

\mr(*num*)

\mi(*num*)

Change the margin of the text or the inter-line spacing. The second character stands for up, down, left, right and internal, respectively. ‘*num*’ can be any nonnegative integer and is interpreted in pixels. Intra-line spacing comes in addition to the line-specific spacing inserted by `\0..\9`. Defaults are zero. You can also omit the number, which restores that particular value to the one in effect at the beginning of the label, such as `\mu()`. Note that for signalling charts, Msc-generator always adds enough left and right margins to arrow labels to avoid overlapping the label with the arrowhead. Thus if you specify less margin, it may not show as you expect.

\mn(*num*)

\ms(*num*)

Changes the size of the normal or small font. This applies only to the label, where used, not globally for the entire chart. Defaults are `\mn(16)\ms(10)`. You can also omit the number, which restores that particular value to the one at the beginning of the label.

\pl \pc \pr

Changes the indentation to left, centered or right. Applying at the beginning of a line (t.i., before any literal character) will apply new indentation to that line and all following lines within the label. Applying after the beginning of a line will only impact subsequent lines.

\{ \[\\" \; \# \} \]

These produce a literal ‘{’, ‘[’, “”, ‘;’, ‘#’, ‘}’ or ‘]’, respectively, since these are characters with special meaning and would, otherwise signal the end of a label. The last two can actually be used without the backslash, but result in a warning.

\\$xxxx

This escape can be used to insert any unicode character into a label, via its codepoint. Specify the codepoint in hexadecimal using exactly four hex digits. E.g., `\$00A9` will result in the copyright symbol. See more on international characters in Section 7.12 [International Text Support], page 67.

\|

This escape is a non-formatting escape that generates no output. It can be used at the beginning of a label to delimit those formatting escapes that are included

in the default formatting restored by the ‘\s()’ escape and used to format the label number, from those which are just to be applied at the beginning of the label.

\N This escape marks the position of the label number within the label. If omitted the number is prepended to the beginning of the label (after the initial formatting escapes). If no number is specified for the label, this escape has no effect. You can specify ‘\N’ multiple times, with each occurrence being replaced by the number. Note that if you omit ‘\N’, the number inserted at the beginning of the label is augmented by the value of the `numbering.pre` and `numbering.post` options, whereas with the ‘\N’ option, those are not used.

\r(*refname*)

This escape inserts the number of the referenced element. Use the `refname` attribute to name elements. Similar to the ‘\N’ escape, the value of the `numbering.pre` and `numbering.post` options are ignored. When no name is given (that is ‘\r()’) the escape is equivalent to ‘\N’.

\L(*link target*)

\L() This escape shall be used in pairs to assign a hyperlink to the text in between the two. The target of the link shall be specified in the first escape, whereas the second shall be empty. For example: `a->b: A \L(http://abc.com)link\L() to abc.com.`; The link target cannot contain closing parenthesis. If it contains opening square brackets ([), opening curly braces ({), semicolons or hash marks, which normally terminates colon labels, use quotation marks around the label, such as `a->b: "A \L(http://abc.com/#x)link\L() to abc.com/\#x.`; Note that *tag labels* of signalling chart boxes can also contain links. See more in Section 8.3 [Links], page 79, below.

***** In Block Diagrams, adding this to the label of a block, will insert the name of the block there. Using this as a default label (e.g., via `use_label="*";`) makes the label default to the name of the block.

Font size commands (including superscript or subscript) last until the next font size formatting command. For example in order to specify a subscript index, use `label="A_i\+ value".`

Any unrecognized escape characters in a label are removed with a warning. Unrecognized escapes and plain text in `text.format` attributes is ignored with a warning.

Note that the `text.*` chart options can be used to set the default text formatting for signalling charts and block diagrams. See Section 8.6.3 [Text Formatting Attributes], page 87, for details.

8.3 Links

Hyperlink info can be added to labels (or parts of labels). This is useful when the resulting image is embedded into web pages. Hyperlinks are not (yet) exported into a PDF or SVG file. Currently link information can be extracted via the `-T ismap` command-line option, which provides an NCSA formatted ismap file. Such files are used by Doxygen, for example, so this feature is most applicable to Doxygen integration. (If you specify a link target as

\ref Name then Doxygen will point to a function or class named Name in the documentation created by Doxygen.)


In the MFC GUI, links are visible and if they represent a URL they are clickable. In the CLI GUI they are not.

To add a hyperlink to a label, use two \L() escape in a pair, making the text between them to point to the link target. The target of the link shall be specified as the parameter to the first \L() escape. In this case the link target may not contain closing parenthesis⁶. Using this method it is possible to add several links to (different parts of) the same label.

In signalling charts, you can also use the url attribute. This makes the whole label a hyperlink. The value of the attribute is the target of the link. You cannot use both mechanisms to the same label. Note that you can only use the first method to add hyperlinks to box tags.

Links also change formatting. In the plain design, they became blue and underlined. In signalling charts, this is governed by the text.link_format chart option and attribute. Any formatting escape sequence you specify as value to this chart option or attribute will be applied at the beginning of the link text and de-applied at the end.

```
hscale=1.2;
text.wrap=yes;
numbering=yes;
a->b: Whole label is link
[url="http://abc.com"];
a->b: One \L(http://abc.com)word\L() is link;
a->b: "One \L(http://abc.com/#x)word\L() is link"
a->b: \L(\ref note)At the beginning\L() only;
a->b: \| \L(\ref note)At the beginning\L(), but
        not the number;
a->b: Different \L(target)style\L() of link
        [text.link_format="\c(red)\I"];
```



8.4 Numbering

Signalling chart arrows, boxes and dividers (any element with a label, except entities) and block diagram blocks can be auto-numbered. It is a useful feature that allows easier reference to certain steps or blocks from explanatory text. At present numbering is not implemented for graphs. To auto assign a number to an element, simply set its number attribute to yes. You can also assign a specific number, in that case the element will get that number and subsequent elements will be numbered (if they have number set to yes) from that number upwards.

⁶ If the link target contains opening brace, hashmark, semicolon or symbols, which terminate a colon label, use quotation marks around the label - it is not possible to use escapes such as \[inside a link target.

Notes and comments (signalling charts) will not increase numbering, instead they carry the number of the element they are referring to. If the target element had no number comments will have none, even if numbering is turned on for them.


Styles can also control numbering. If a style has its `number` attribute set to `yes` or `no`, any element that you assign the style to will have its attribute set likewise. See Section 8.6.4 [Styles], page 88, for more.

In order to minimize typing, the `numbering` chart option can be used. It can be set to `yes` or `no` and serves as the default for freshly defined elements. You can set the value of `numbering` at any time and impact elements defined thereafter. You can use scoping to enable or disable numbering for only blocks of the chart, see Section 8.7 [Scoping], page 89.

Most of the time you just declare `numbering=yes` at the beginning of the chart and be done with it. However, if you want to control that only some parts of the elements (e.g., only concrete messages and not boxes, for example) got a number, you may need the other alternatives.

```
hscale=auto;

a->b: Not numbered;
a->b: Numbered [number=yes];
a->b: Numbered [number=10];
a->b: Not numbered;
numbering=yes;
a->b: \bNumber is bold;
a->b: \|\bNumber is not bold;
a->b: Our number is: \N;
a->b: (\N) This may be optional;
numbering.pre="Step #";
numbering.post=" is: ";
a->b: Some action;
a->b: Some action;
numbering.pre="\c(red)\b";
numbering.post=": \s()";
a->b: Some action;
a->b: \c(blue)\uSome action;
a->b: \|\c(blue)\uSome action;
```



If numbering is turned on for a label, the number is inserted at the beginning of the label and is followed by a semicolon and a space by default. More precisely, the number is inserted after any initial text formatting sequences, so that it has the same formatting as the label itself (see Section 8.2 [Text Formatting], page 76)⁷. The above default can be changed by inserting the `\N` escape sequence into a label. This causes the number appear where the `\N` is inserted, as opposed to the beginning of the label. In this case, the colon and the space is omitted, only the number itself is inserted.

The colon and space can be changed to some other value by setting the `numbering.post` chart option to the string you want to append to the number. Similar, any string the

⁷ You can use the `\|` formatting escape to insert a non-visible break into a stream of formatting escapes. The number will be inserted there.

`numbering.pre` option is set to will be prepended to the number (empty by default). Both options are ignored when using the `\N` escape sequence to set the label position.

Note that for the last two arrows formatting escapes were added to the `numbering.pre` option. These are reversed by the '`\s()`' escape in the `numbering.post` option. See Section 8.2 [Text Formatting], page 76, for more details.

The format of the number can be set with the `numbering.format` chart option. You can specify any of '123', 'iii', 'III', 'abc', or 'ABC' for arabic, lowercase and uppercase roman numbers or lowercase and uppercase letters, respectively⁸. You can also prepend or append any text before or after the above strings, those will be prepended or appended to the number (and will be included also when the number is inserted via the '`\N`' escape).

If you want to specify an increment other than one, use the `numbering.increment` option with a nonzero integer. You can even count backwards using a negative increment. In case of multi-level numbering (see below) this applies to the last level,

Note that the value of the `numbering` options is subject to scoping, that is any change lasts only up to the next closing brace.

Note also, that when using roman numbers or letters, you can use such numbers as the value of the `number` attribute, as shown below for '7c'.

It is also possible to have multi-level numbering (such as 1.1). To achieve this, use the `numbering.append` chart option and specify the format of the second level including any separator. Use the same format as for `numbering.format` above.

It is possible to change the format of a multi-level label via the `numbering.format` option. Simply use multiple of the number format strings (such as '123' or 'roman') as in the 'Exotic format' line of the example above. If you use less number format strings than the current number of levels (as in the 'Only the last number' line of the example), Msc-generator displays only the end of the number, omitting levels from the top. Those levels, however, are still maintained, just are not displayed.

The `numbering.append` option can only be used to add levels. There is no explicit way to decrease the number of levels, you have to use scoping to achieve that. On the example above, the second level appended in the scope of 'Alternative #1' is cancelled at the end of the scope, so we need to append a second level also in 'Alternative #2', which then restarts from 'a'.

⁸ Using 'arabic', 'letters' or 'roman' is also valid (both uppercase or lowercase).


```

hscale=auto, numbering=yes;

numbering.format = "roman";
a->b: Lowercase roman;
numbering.format = "ABC)";
a->b: Uppercase letters;
numbering.format = "123";
a->b: Arabic numbers;
{
    numbering.append = ".123";
    a->b: Two-level numbers;
    a->b: Set last number [number=5];
};

a->b: Back to one level;
box a--b: Alternative \#1 {
    numbering.append = "abc";
    a->b: Step \N;
    b->a: Step \N;
}
a..b: Alternative \#2 {
    numbering.append = "abc";
    a->b: Step \N;
    b->a: Step \N [number=c];
    numbering.format = "roman-roman";
    a->b: Exotic format;
    numbering.format = "roman";
    a->b: Only the last number;
    numbering.format = "123abc";
    a->b: Back to normal;
};
a->b: Back to one level;

```



Finally, if an element is named, you can reference the number of that element in another label using the `\r(name)` escape sequence. Signalling chart elements can be named using the `refname` attribute, while block diagram blocks can simply have their name at their definition. Note that the value of the `numbering.pre` and `numbering.post` options are ignored when inserting the number of a referenced element, similar to how the `\N` escape inserts numbers. Specifying an empty `\r()` escape inserts the number of the current element and is thus equivalent to `\N`.


8.5 Specifying Colors

Msc-generator has a number of color names defined initially, among them: `none`, `white`, `black`, `red`, `green`, `blue`, `gray` and `lgray`, the first for completely transparent color, and the last for light gray. Graphviz has a different set of default color names, so we adopted those for graphs. When you specify a color by name, no quotation marks are needed.

Color names in signalling charts and block diagrams can be appended with a '+' or '-' sign and a number between [0..100] to make a color lighter or darker, respectively, by the

percentage indicated. Any color +100 equals white and any color-100 equals black. Aliases can be further appended with a comma and a value between [0..255] (or [0..1.0] similar to RGB values). This specify color opaqueness: 0 means fully transparent and 255 means fully opaque.

```
a,_b;
--: Overall [fill.color = blue+50,
            fill.gradient=up] {
    a--b [fill.color=red]:      solid;
    a--b [fill.color=red,128]:  transparent;
    a--b [fill.color=red-50]:   dark red;
    a--b [fill.color=red+50]:   light red;
    a--b [fill.color=none]:    none;
    defstyle arrow [line.color=red,
                    arrow.color=green,
                    text.color=blue];
    a->b: Unmodified;
    a->b: Yellow overcast [color=++yellow,150];
}
```



You can specify colors giving the red, green and blue components. For graphviz graphs precede them by hash-marks (#) and use hexadecimal numbers, six digits in total.

For signalling charts and block diagrams separate the R, G and B components by commas. An optional fourth value can be added for the alpha channel to control transparency. Values can be either between zero and 1.0 or between 0 and 255. If all values are less than or equal to 1, the former range is assumed⁹. If any value is negative or above 255 the definition is invalid. Note that you must not enter spaces between the color name, its lighter/darker or transparency modifier or between the RGB values. You can mark a color a overlay color by prepending the ‘++’ symbol. An overlay color applied to an attribute is not set as the color value. Instead, it is overlaid over and thus is combined with the existing color. (See the arrows in the example above.) You can also assign an overlay color to a style or a color name (see below). (This is how their `weak` style is implemented.)

For charts other than graphviz graphs, it is possible to define your own color names using the `defcolor` command as below.

```
defcolor alias=color definition, ... ;
```

Color names are case-sensitive and can only contain letters, numbers, underscores and dots, but can not start with a number or a dot and can not end with a dot. Aliases can also be later re-defined using the `defcolor` command, by simply using an existing alias with a different color definition.

Msc-generator honors scoping for colors, as well. Color definitions (or re-definitions) are valid only until the next closing brace ‘}’. This makes it possible to override a color only for parts of the chart, returning to the default later. Note that you can start a new scope any time by placing an opening brace. See Section 8.7 [Scoping], page 89, for more on scopes.

⁹ This mechanism allows both people thinking in range [0..1] and in [0..255] to conveniently specify values. (Internally values are stored on 8 bits.)

8.6 Common Attributes

Attributes can influence how chart elements look like and how they are placed. There is a set of attributes that apply to several types of elements, so we describe them collectively here.

Attribute names are case-insensitive. Attributes can take string, number or boolean values. String values shall be quoted in double quotes ("") if they contain non-literal characters or spaces¹⁰. Quoted strings themselves can contain quotation marks by preceding them with a backslash, like '\\"'. Numeric values can, in general be floating point numbers (no exponents, though), but for some attributes these are rounded to integers. Boolean values can be specified via yes or no. The syntax of color attributes is explained in Section 8.5 [Specifying Colors], page 83.

The attributes below can be part of a *style*, see Section 8.6.4 [Styles], page 88.

8.6.1 Line and Fill Attributes

Currently line and fill attributes apply only in signalling charts and block diagrams. Graphviz has a totally different way of specifying line appearance via the `style=` attribute. Further languages planned to be added will use the below attributes.

`line.color`

Specifies the color of the lines, see Section 8.5 [Specifying Colors], page 83, for the syntax.

`line.width`

Specifies the width of the line in pixels. In case of double or triple lines this is not the total width of the compound line, but only that of an individual components. E.g., a double line with `line.width=1` is three pixels wide in total.

`line.type`

Specifies the type of the line. Its value can be `solid`, `dotted`, `dashed`, `long-dashed`, `dash_dotted`, `double`, `triple`, `triple_thick` or `none`.

`line.corner`

For boxes-like elements this attribute specifies how the corners of the box are drawn. Its value can be `none`, `round`, `bevel`, `note`. It has sometimes effect on other elements. Signalling chart arrows pointing to the same entity use this. Also arrows in block diagrams are impacted.

`line.radius`

Specifies the size of the corner above. In some cases it is also used to specify the turning radius of turning arrows (see `line.corner` above).


¹⁰ Specifically for signalling charts and block diagrams: strings that contain characters other than letters, numbers, underscores or dots, must be quoted. If the string starts with a number or a dot or it ends with a dot, it must also be quoted. The only exception to this are built-in style names, see Section 8.8 [Defining Styles], page 90, or color definitions containing commas and + or - signs, see Section 8.5 [Specifying Colors], page 83. For graphviz graphs strings can not start with number and can only contain alphanumeric characters - else they have to be quoted.

```

a--b: Squarish, radius=0;
a--b: Round, radius=5 [line.radius= 5];
a--b: Round, radius=10 [line.radius=10];
a--b: Bevel, radius=10 [line.corner=bevel];
a--b: Note, radius=15 [line.radius=15,
                        line.corner=note];

pipe a--b: Radius=5;
pipe a--b: Radius=10 [line.radius=10];
pipe a--b: Radius=15 [line.radius=15];
pipe a--b: Left side [side=left];

```

**fill.color**

Defines the background of an element covering an area. Specifying **none** results in no fill at all (transparent).

fill.color2

If this attribute is specified then the fill gradient will not be between **fill.color** and a lighter variant, but between **fill.color** and the value specified here. If no gradient specified or **button** is used, this attribute has no effect.

fill.gradient

Defines the gradient of the fill. It can take five textual values **up**, **down**, **in**, **out** and **button**. The first two result in linear gradients getting darker in the direction indicated. The second two result in circular gradients with darker shades towards the center or edge of the entity box, respectively. The last one mimics light on a button. In addition, it can also be set to a numerical value, which results in a linear gradient and specifies the angle of it.

```

hscale = auto;
defstyle entity
[fill.color="yellow-25",
 text.format= "\mu(10)\md(10)\ml(10)\mr(10)" ];
Up      [fill.gradient=up],
Down    [fill.gradient=down],
In      [fill.gradient=in],
Out     [fill.gradient=out],
Button  [fill.gradient=button];
Slant   [fill.gradient=45];

```



8.6.2 Shadow Attributes

`shadow.offset`

If not set to zero, then the element will have a shadow (default is 0). The value of this attribute then determines, how much the shadow is offset (in pixels), in other words how "deep" the shadow is below the entity or box. In graphviz, this attribute is called `shadow_offset`.

`shadow.color`

The color of the shadow. This attribute is ignored if `shadow.offset` is 0. In graphviz, this attribute is called `shadow_color`.

`shadow.blur`

Specifies how much the shadow edge is blurred (in pixels). E.g., if `shadow.offset` is 10 and `shadow.blur` is 5, then half of the visible shadow will be blurred. Blurring is implemented by gradually changing the shadow color's transparency towards fully transparent. This attribute is ignored if `shadow.offset` is 0. In graphviz, this attribute is called `shadow_blur`.

```
hscale = 0.5;
One  [shadow.offset= 5],
Two  [shadow.offset= 5, shadow.blur= 2],
Three [shadow.offset=10, shadow.blur= 5],
Four  [shadow.offset=10, shadow.blur=10];
    
```

8.6.3 Text Formatting Attributes

These attributes can be used to change the appearance of text labels. They are basically alternatives to *text formatting escapes*, see Section 8.2 [Text Formatting], page 76. They do not apply to graphs: Graphviz labels can only be styled using text formatting escapes in a colon-label. Other languages have them.

`text.ident`

This can be `left`, `center` or `right` and specifies the line alignment of the label. The default is centering, except for non-empty boxes in signalling charts, where the default is left. It can be abbreviated as simply `ident`.

`text.color`

Sets the color of the label. See Section 8.5 [Specifying Colors], page 83, for the syntax of colors.

`text.bgcolor`

Sets the background color of the label, none by default.

`text.font.face`

Specify the font face family using this attribute, such as `Arial` or `Helvetica`. The fonts available depend on the platform. See Section 7.14 [Fonts], page 74.

`text.font.lang`

Specify the language for which you seek glyphs. This has effect only on Mac/Linux, see Section 7.14 [Fonts], page 74.

text.font.type

Select between normal or small font, superscript or subscript using the values **normal**, **small**, **superscript** and **subscript**, respectively.

text.bold**text.italic****text.underline**

You can set them to **yes** or **no**.

text.gap**text.gap.up****text.gap.down****text.gap.left****text.gap.right**

These four attributes can be used to set the margins around the label in pixels.

If you set **text.gap**, it sets all four to the same value.

text.gap.spacing

This sets the line spacing in pixels.

text.size.normal**text.size.small**

These sets the height of the normal font type (see **text.font.type** above) or the height of small, superscript and subscript, respectively.

text.format

Takes a (quoted) string as its value. Here you can specify any of the text formatting escapes that will govern the style of the label, see Section 8.2 [Text Formatting], page 76. Specifying them here or directly at the beginning of the label has the same effect, so having this attribute is more useful for styles.

text.link_format

Similar to the above, you can specify the formatting applied to links. See Section 8.3 [Links], page 79.

text.wrap

Used in signalling charts only. Can be set to **yes** or **no**. If disabled (default), the label will follow the line breaks inserted by the user. If enabled, these line breaks are ignored and the line is typeset to fill available space, see Section 8.2 [Text Formatting], page 76.

8.6.4 Styles


Styles are packages of attribute definitions with a name. Applying a style to any element can be easily done by simply stating the name of the style whereever an attribute is allowed. See in the example below, how the **strong** and **weak** styles are applied to entities.

```

B: Backend [shape=def.oval];
S: Server\n\-(Frontend) [strong];
C: Client [weak];

C<-: Hit;
C=>S: Request;
S>>B: Query;
S<<B: Response;
C<=S: Reply;

```



For signalling charts styles can contain any of the attributes listed in Section 9.10 [Signalling Chart Attributes and Styles], page 133. For graphs any graphviz attribute can be part of a style. For block diagrams, the applicable attributes are listed in Section 11.3 [Block Attributes], page 162, and Section 11.5 [Arrows in Block Diagrams], page 175. If a style contains an attribute not applicable for the element that you apply the style to, that attribute is simply ignored. For example, applying a style with `fill.color=red` attribute setting to an arrow, will ignore this attribute since arrows take no fill attributes.

You can define your own styles or redefine existing ones. See Section 8.8 [Defining Styles], page 90, for more on this.

8.7 Scoping

Each time an opening brace is put into the file, a new scope begins. Scopes behave similar as in programming languages, meaning that any color name, style or procedure definition take their effect only within the scope, up to the closing brace. Thus if you redefine a style just after an opening brace, the style returns to its original definition after the closing brace. (See Section 8.8 [Defining Styles], page 90.)

In signalling charts scoping also applies to the `numbering` (including `pre`, `post`, `format` and `append`), `compress`, `vspacing`, `indicator`, `angle` and `text.*` chart options. Any changes to these take effect only until the next closing brace. Scoping explicitly does not apply to `background.*` and `comment.*` options. Those take effect until the next such option or all the way to the bottom of the chart. The option `hscale` is global, the last value you set will be used, irrespective of scope.

Similar, with block diagrams scoping applies to all chart options, except `background`.

You can nest scopes arbitrarily deep and can also manually open a new scope, such as below. This practically has no effect in any of the languages other than limiting the effect of an attribute or color/style definition.

```

...numbering is off here...
{
    #number only in this scope
    numbering=yes;
    ...various elements with numbers...
};

```

...other elements with no numbers...

For graphs, this opens an unnamed subgraph - but that has no effect on layout. For block diagrams there is actually no effect. For signalling charts the layout will be the same as in case when they are not enclosed in braces (including the handling of `compress`, `vspacing`, `keep_with_next` and `keep_together` attributes and the use of `parallel` and `overlap` keywords). This is true only if you do not change the `layout` attribute of the block, but use the default, see Section 9.9 [Parallel Blocks], page 127. Thus if you mark an element between the braces with `parallel` elements after the closing brace can be laid out besides it. Marking the entire block with `overlap` or `parallel` will make elements after the block to be laid over or besides the entry block, respectively. See Section 9.9.1 [Parallel Keyword], page 131.

8.8 Defining Styles

It is possible to define a group of attributes as a style and later apply them collectively, see Section 8.6.4 [Styles], page 88.

Styles can be defined using the `defstyle` command, as below.

```
defstyle stylename, ... [ attribute=value | style, ... ], ... ;
```

First you list the name of the style(s) to define then the attributes and their intended values. Similar to color names, style names are case-sensitive and can only contain letters, numbers, underscores and dots, but can not start with a number or a dot and can not end with a dot. You do not have to specify all possible attributes, just those you want to modify with the style. The rest of the attributes will remain unspecified. When you apply the style to an element, attributes of the element that are unspecified in the style are left unchanged.

You can also enlist styles among the attributes. In this case the newly defined style inherits all the attributes specified in that style. If you apply a style to an element, those attributes of the style, which not applicable to that particular element type are simply ignored. For example, applying a style including `fill.color` to an arrow will silently ignore the value of the `fill.color` attribute.

The same syntax above can be used to extend and modify styles. You can add new attributes to an existing style or modify existing attributes. This is when listing multiple styles comes in handy. You can set attributes to the same value in multiple styles in a single command. Re-defining an existing style do not erase the attributes previously set in the style. Only the new attribute definition is added - changing the value of the attribute if already set in the style.

It is also possible to unset an attribute by specifying the attribute name, followed by the equal sign, but no value.

There are a number of default, built-in styles that govern the default appearance of elements. By modifying these you can impact, e.g., all the arrows in a chart or all edges in a graph. This is how chart designs operate: by modifying the built-in styles.

If you want to change a set of attributes for multiple elements (such as both for arrows and dividers) simply list these separated by commas before the attributes.

```
defstyle arrow, divider [line.width=2] ;
```

It will apply to both.

In addition to default styles, most languages have a number of *refinement styles*. Such styles may contain some attributes set and are applied after and in addition to default styles. For example, after applying the default `arrow` style to a message in a signalling chart or an arrow of a block diagram, specified as `a=>b;`, an additional refinement style, named `=>` is applied, making the arrow double-lined.

Redefining refinement enables you to quickly define, e.g., various arrow styles and use the various symbols as shorthand for these.

Thus, in summary the actual attributes of an element are set using the following logic. (There are minor variances for each language.)

1. If you specify an attribute directly at the element (perhaps via applying a style), the specified value is used¹¹.
2. Otherwise, if the attribute is set in the refinement style (at the point and in the scope of where the element is defined), the value there is used.
3. Otherwise, if the attribute is set in the default style of the element, the value there is used.
4. Otherwise (if they exist in the language), the value of the applicable chart option is used, e.g., `text.*` for signalling charts. In order for these chart options to be effective default styles usually have no value specified for these attributes. You can set these attributes in styles, e.g., to set font type for empty boxes, which will take precedence over chart options.

8.9 Chart Designs

A chart design is a collection of color, style and procedure definitions and values of certain chart options. For signalling charts the value of the `hscale`, `numbering`, `compress`, `vspacing`, `text`, `background` and `comment` options are included. For numbering you can turn it on or off and specify the format of the top level number - but you cannot specify multiple levels. For block diagrams the `background` and the `conflict_report` chart options can be made part of designs.

Note that color, style and procedure definitions are local to a scope (take effect only within) and can be stored in designs. Shape and design definitions on the other hand are global and as such can not be added to designs.

Except for graphs, we differentiate are *full designs* and *partial designs*. A full design contains a value for all the chart options, default colors and styles. A partial design contains values only for some of these. E.g., the `thick_lines` design is a partial one - it merely makes all lines of width 2 in all the default styles, but leaves color, line type, fill or any other attribute or chart option unchanged.

In signalling charts you can apply a full style using the `msc = <design_name>` chart option and a partial style using `msc += <style_name>`. For graphs and block diagrams use `usedesign <design_name>` (for all designs). Note that it is best to apply a full design at the first line of the chart, as any chart option specified before it will likely be overwritten by applying a full design. You can apply a full or partial design inside any scope, the effect will last only within the scope.

¹¹ If you specify the attribute several times, the last one is used.

Currently the following partial designs ship with Msc-generator for signalling charts: `hcn`, `thick_lines`, `all_blue`, `feng_shui_notes`. The first one simply sets `hscale` to auto and turns on compression and numbering. The second one makes lines of all default styles of width 2. The third makes the color of lines in all default styles blue. The last makes notes rounded and red on yellow background. Block diagrams have the same, except `hcn`. Try them.

You can define or re-define chart designs by using the syntax below.

```
defdesign designname {
    [ msc=parent design | usedesign parent design ]
    [ msc+=partial design | usedesign partial design]
    options, ...
    color definitions, ...
    style definitions, ...
    procedure definitions, ...
}
```

First you can name an existing full design to inherit from using the '`msc=`' option or via '`usedesign parent design`'. If specified the design will become a full design, too. Thus in each such design definition the default styles are always present and fully specified. If omitted, the style will become a partial style. Then you can specify optional multiple '`msc+=`' or '`usedesign`' options to bring in partial designs. Finally, you can define colors, styles and procedures in any order and/or set one or more of the attributes mentioned above.

It is possible to add your design definitions to a file having the `.signalling`, `.graph` or `.block` extension and making them available in all charts for use. See Section 7.2 [Design Library], page 56.

8.10 Defining Shapes

This section describes how you can define new shapes that can be used in entity headings of signalling charts or as block shapes in block diagrams. Note that shapes are not (yet) available in graphs.

To define a new shape use the `defshape` command.

```
defshape shape_name {
    defining lines;
    ...
};
```

The `shape_name` is the name of the shape to define. You can use letters, numbers, underscores and dots; and you must start with a letter. Each *defining line* shall start with a capital letter, one of the list below. After this letter comes a space delimited list of arguments, terminated by a semicolon. Comments can start with a hashmark '#' as usual and empty lines are also permitted.

S Starts a new *section* in the new file. Each shape may contain three optional sections. You shall include the number of the section after the S separated by a space. Sections contain paths, which describe the shape. Anything that comes

after this line will belong to this section all the way until the next **S** or the end of the shape definition.

- 0 Section 0 specifies the background of the shape. This shall be a (set of) closed path(s), which will be filled by Msc-generator using the value of **fill.color** attribute.
 - 1 Section 1 specifies the foreground of the shape. This shall be a (set of) closed path(s), which will be filled by Msc-generator using the value of the **line.color** attribute.
 - 2 Section 2 specifies an alternative, additional way to specify the foreground. It shall be a set of potentially open paths, which will be drawn (stroked) using the value of the **line.color** attribute. Contrary to section 1, other line attributes will also be applied, such as **line.width** and **line.type**. You can specify both Section 1 and 2, in this case both will be drawn.
 - 3 Section 3 is a combination of Sections 0 and 2. Whatever shape is specified will be filled with **fill.color**, as well as have its path stroked with **line.width** and **line.type**. Only closed subpaths will be filled. Using Section 3 is equivalent to specifying the same path twice, once in a Section 0 then in a Section 2.
- M** This command is used within the specification of a path, and moves the (imaginary) cursor to a given point. Two space-separated numbers after the **M** specifies the x and y coordinates. Note that you can use floating-point numbers and any scaling or position—Msc-generator will normalize the size of each shape based on their height.
- L** This command is used within the specification of a path, and draws a straight line from the current position of the (imaginary) cursor to a given point. Two space-separated numbers after the **L** specifies the x and y coordinates.
- C** This command is used within the specification of a path, and draws a curved line from the current position of the (imaginary) cursor to a given point. Six space-separated numbers after the **C** specifies the x and y coordinates of the target point and two control points, respectively. The curved line is computed as a cubic bezier curve.
- E** This command is used within the specification of a path, and closes a path. It takes no arguments.
- T** This optional item specifies where the entity label shall be drawn if the shape is used in a signalling chart entity heading. This place will also be used if the shape is applied to a block diagram block, which has content. It takes two points (four numbers) as arguments specifying the opposing corners of a rectangle. If this line is omitted, the label will be displayed below the shape. This line can be specified before, inside, between or after sections.
- H** This optional item specifies which portion of the shape shall be shown in the hints popup box on Windows in the internal editor (see Section 7.4.3 [Typing

Hints and Autocompletion], page 60). This also takes 4 numbers specifying a rectangle as above. If omitted the whole shape is shown (on a miniature scale). This line can be specified before, inside, between or after sections.

- P This optional item can be used to define a *port*. Ports are currently used by Block Diagrams, but we plan to extend support to Graphs, as well. After the letter you need to specify an X and an Y coordinate and a name (containing letters, numbers, underscores and dots). Optionally a *direction* can also be associated with a port by adding a number at the end. This number signifies at what angle arrows leave or arrive at the port. It is interpreted in degrees clockwise from the north (top) direction. So east is 90, left is 270.

See Section 7.2 [Design Library], page 56, to see where to find and put your own files defining Shapes.

8.11 Procedures

Procedures allow you to store some chart text and replay it later. They are useful to save typing, promote re-use instead of copy/paste. Since procedures can take parameters their result can be customized.

To define a procedure, use the `defproc` command.


```
defproc proc_name {
    ...chart text for procedure...
};

defproc proc_name($param_name[=default value, ...] {
    ...chart text for procedure...
};
```

You can list parameters and their (optional) default value, as well. Parameter names should begin with the dollar sign (\$). They can be used anywhere a string or number is needed, such as for attribute values, attribute names, entity and node names, etc. You can not invoke a keyword using a parameter.

```
defproc addgroup($group_name, $target, $color=black) {
    edge [color=$color];
    A ~$group_name->$target:: \c($color) 1st\n(\l0($group_name));
    B ~$group_name->$target:: \c($color) 2nd\n(\l0($group_name));
    C ~$group_name->$target:: \c($color) 3rd\n(\l0($group_name));
}

graph {
    composite;
    replay addgroup(iron, composite);
    replay addgroup(steel, composite, blue);
    replay addgroup(gold, composite, yellow);
};
```




The `\Q()` text formatting escape can be used to substitute a parameter value into a label. The tilde (~) character can be used to append or prepend string constants to a parameter, enabling string composition. In graphviz the + character can also be used. The double-dollar special symbol (`$$`) can be used to denote a string that is always unique for each procedure replay but when printed it is empty. This can be used, when you want to define a node (or entity or marker for signalling charts) for each time the procedure is replayed, but you do not want to pass a procedure parameter to compose with (such a `$group_name` above).

```

defproc addgroup($color=black, $target) {
    edge [color=$color];
    A~$$->$target:: \c($color) 1st;
    B~$$->$target:: \c($color) 2nd;
    C~$$->$target:: \c($color) 3rd;
}

graph {
    composite;
    replay addgroup(), composite;
    replay addgroup(blue, composite);
    replay addgroup(yellow, composite);
}

```



As you can see you can fall back to the default value of parameters, even if they are not the last one.

It is possible to replay a procedure within another procedure. Do not use recursion. It is not allowed to define a procedure in another procedure, however. If you replay a procedure inside another one, the parameters of the outer procedure will not be visible in the inner procedure.

When you refer to a style, color or procedure name, or an entity, node or marker from within a procedure definition, their value will be substituted at the place of the replay. Thus their existence cannot be verified at the time of definition. A full parsing and verification is performed at replay. In order to avoid repeating error messages as much as possible, if an error is already caught at procedure definition time, the procedure cannot be replayed.

Color Syntax Highlighting is also limited for procedures, for example, styles cannot be determined to exist, so they are never marked for error.

Styles and colors defined inside a procedure are local to that procedure (since the procedure has its own scope). If you want to export these after a replay, you can use the [export=yes] attribute at the procedure definition. This enables you to create procedures that define styles based on some parameters.

```
defproc adjust_line($color, $width, $weak_width=1) [export=yes] {
    defstyle arrow, blockarrow, vertical, pipe, box, emptybox, divider,
        symbol, indicator, note, vertical_pointer, vertical_range,
        vertical_bracket, vertical_brace
        [line.width=$width, line.color=$color];
    defstyle box_collapsed, box_collapsed_arrow, entity, entitygroup,
        entitygroup_collapsed
        [line.width=$width, line.color=$color];
    defstyle box, box_collapsed, emptybox
        [tag.line.width=$width, tag.line.color=$color];
    defstyle entity, entitygroup_collapsed
        [vline.width=$width, vline.color=$color];
    defstyle weak [line.width=$weak_width, vline.width=$weak_width];
};

hscale=auto;
a->b: before;
replay adjust_line(black, 2);
b->c: after black and 2;
replay adjust_line(blue, 3);
c->d: after blue and 3;
```


Finally, inside procedures you may use conditional statements that execute or don't based on the value of some parameters. After the if statement you can use either a single parameter or two strings (potentially composed of multiple substrings with ~) compared with one of the following operators: <=, <, ==, >, >= and <> (the latter for non-equal). String comparison is UTF-8 and not numeric (thus "2" is larger than "10"). If you use only a single parameter, the condition is true if the parameter is any non-empty string.

```

defproc threeway($e1, $e2, $note="")
{
    $e1->$e2: SYN;
    if $note then {
        $e1<-*$e2: SYNACK;
        note [label=$note];
    } else {
        $e1<-$e2: SYNACK;
        $e1->$e2: ACK;
    };
}

replay threeway(a, b);
replay threeway(b, c, "Lost, so we fail.");

```



Note that you can use braces to include more than one statement for the then or else branches.

8.12 Variables

In addition to procedure parameters strings preceded by a dollar sign (\$) can also denote variables. Variables, just like procedure parameters can contain strings or numbers and can be used anywhere where an attribute, entity, marker, block or node name or attribute value is needed. Similar to procedure parameters, you can use the tilde (~) character to prepend or append strings with variables.

One big difference between variables and procedure parameters is that variables can be used outside and independent of procedures. This allows conditional statements (**if**) to be used outside parameters.

You may want to use variables for two reasons. One is to save typing. A long and often repeated attribute value may be shortened (and potentially replaced everywhere in one go). The second reason is to influence the operation of procedures without needing to pass a lot of parameters.

You can define or set the value of a variable using the **set** command. The **<value>** must either be a number or an alphanumeric string - any other string must be enclosed in-between quotation marks.

```

set $<name> = <value>;

```

```

hscale=auto;
set $a="Very long";
set $b=" text you don't want to copy-
a->b [label=$a~$b];
b->a [label=$a~$b];
set $a="Even longer";
a--b [label=$a~$b];

```

```

Very long text you don't want to copy-paste
Very long text you don't want to copy-paste
Even longer text you don't want to copy-paste

```

```

sequenceDiagram
    participant a
    participant b
    a->>b: $a
    activate b
    b-->a: $a
    deactivate b
    a->>b: $a
    deactivate b
    note over a: Even longer text you don't want to copy-paste

```

Note that you can use the `set` command to change the value of parameters inside a procedure.

Variable definitions honour scoping, thus a variable defined in a scope is not available after the scope has been closed (after the closing brace). It is, however, available in any nested scope. This includes procedures replayed (to any depth). Variable lookup inside procedures happens during replay, thus you can refer to variables inside procedures (even if the variables are defined later in the input file). Any value change made to a variable remains in effect after a scope closure.

8.13 File Inclusion

It is possible to include the text of another file in the current chart, similar to how it works in the C preprocessor. Use the `include` command followed by a filename in quotation marks. Currently the file is searched only in the directory of the file containing the `include` command. You can use relative or absolute paths to include files from other directories.

On Windows, you can use both forward and backward slashes to separate directories, but due to text escape sequences you need to double backslashes. Thus on Windows both are valid syntax:

```
include "../file.signalling";  
include "..\\file.signalling";
```

You can include other files from included files, in arbitrary depth.

When you issue an `include` command, parsing continues in the specified file and returns to the original one - just after the `include` command. Certain errors (such as opening braces without closing pairs) may cause errors in the original file - try including only totally correct files.

Note that Color Syntax Highlighting does not parse include files. Any styles, colors or procedures defined there may be marked as undefined in the original file.

Note that if you want to define a set of procedures, designs or shapes you frequently use, in the Windows GUI it is better to place them in a design library (Section 7.2 [Design Library], page 56) as opposed to including them at the beginning of your file. Design libraries are pre-parsed when the GUI starts up and are not re-parsed again at compilation. In addition any styles or designs defined in them are taken into account by Color Syntax Highlighting.

9 Signalling Chart Language Reference

9.1 Titles

The `title` and `subtitle` commands can be used to specify titles for the chart. You must supply a label, perhaps using the colon syntax.

```
title: This is a title;
subtitle: This is a subtitle;
```

The title and subtitle include text and a box around the text - the latter being omitted in the ‘plain’ design. You can turn it on by setting the ‘`line.*`’ and ‘`fill.*`’ attributes. The default attributes are taken from the default styles `title` and `subtitle`, changing these will affect all titles in the chart. Entity lines are not drawn behind the titles by default, this can be changed by setting the `vline.*` attribute.

9.2 Specifying Entities

Entities can be defined at any place in the chart, not only at the beginning.

Entity names can contain upper or lowercase characters, numbers, dots and underscores. They are case sensitive and must start with a letter or underscore and cannot end in a dot. If you want other characters, you have to put the entity name between quotation marks every time it is mentioned. This, however, makes little sense: you can set the label of the entity to influence how the entity is called on the drawn chart.

It is also possible to define entities without attributes (having all attributes set to default) by typing

```
entityname, ...;
```

It is also possible to change some of the attributes later in the chart, well after the definition of the entity. The syntax is the same as for definition — obviously the name identifies an already defined entity.

Note that typing several entity definition commands one after the other is the same as if all entity definitions were given on a single line. Thus

```
a;
b;
c;
```

is equivalent to

```
a, b, c;
```

Also, `heading` commands are combined with the definitions into a single visual line of entity headings.

9.2.1 Entity Positioning

Entities are placed on the chart from left to right in the order of definition. This can be influenced by the `pos` and `relative` attributes.

Specifying `pos` will place the entity left or right from its default location. E.g., specifying `pos=-0.25` for entity B makes B to be 25% closer to its left neighbour. Thus `pos` shall be


specified in terms of the unit distance between entities. (Which is 130 points - a historic value kept for backwards compatibility.)

The next entity C, however, will always be from a unit distance from the entity defined just before it, so in order to specify a 25% larger space, on the right side of entity B, one needs to specify pos=0.25 for C.

```
A, B, C, D;  
A->B-C-D;
```



```
A, B [pos=-0.25], C, D;  
A->B-C-D;
```




```
A, B [pos=-0.25], C [pos=+0.25], D;  
A->B-C-D;
```



The attribute **relative** can be used to specify the base of the **pos** attribute. Take the following input, for example. In this case C will be placed halfway between A and B.

```
A, B;  
A->B;  
C [pos=0.5, relative=A];
```



Note that specifying the **hscale=auto** chart option makes entity positining automatic. This setting overrides **pos** values with the exception that it maintains the order of the entities that can be influenced by setting their **pos** attribute. See Section 9.11 [Chart Options], page 139. In most cases it is simpler to use **hscale=auto**, you need **pos** only to fine-tune a chart, if automatic layout is not doing a good job.

9.2.2 Group Entities

A group entity can contain other entities. Groups can be nested arbitrary deep. To specify a group entity, use curly braces after an entity definition (but before the colon or comma). Between the braces you can list entity definitions, style/color definitions or chart options¹. The curly braces open a new scope, so any style or color definition or chart option takes its effect only within the group of entities between the curly braces. See Section 8.7 [Scoping], page 89, for more information.

Any entity you specify in the group must be a newly defined entity. It is not possible to place already defined entities into a group. Similar, an already defined entity cannot be

¹ Only some of the chart options can be used, the ones that merely change the context and do not draw. E.g. the ‘background’ options cannot be used. Practically only the ‘indicator’ chart option makes any sense.

made a group entity later by adding entities to it. Nor can a group be later extended with additional entities.

The position of a group entity is derived from its members so the ‘pos’ and ‘relative’ attributes cannot be used.


Group entities can be *collapsed*, by setting the ‘collapsed’ attribute to yes (or via the GUI on Windows). A collapsed group entity does not show its member entities, but is displayed as a non-grouped entity. Arrows and boxes in the chart are modified (or even removed) to reflect the collapse. If the ‘indicator’ attribute of the entity is set to yes, a small indicator is shown both inside the collapsed entity and for each arrow or box removed.

Setting the `large` attribute makes the group entity span the entire height of the chart in the background as opposed to drawing a group entity heading. See the difference below. This is in effect only if the group entity is not collapsed.

```
hscale=auto;
CI: Client Infra {
    T: Terminal;
    C: Client;
};

SI: Server Infra [large=yes] {
    S: Server;
    B: Backend;
};

->T: Hit;
T->C: Internal Req;
C=>S: Request to Server;
S>>B: Internal Query;
S<<B: Internal Response;
C<=S: Reply from Server;
T<-C: Internal Reply;
```



9.2.3 Entity Attributes

The following entity attributes can only be set at the definition of the entity.

- | | |
|----------|---|
| label | This specifies the text to be displayed for the entity. It can contain multiple lines or any text formatting character. See Section 8.2 [Text Formatting], page 76. If the label contains non alphanumeric characters, it must be quoted between double quotation marks. The default is the name of the entity. |
| pos | This attribute takes a floating point number as value and defaults to zero. It specifies the relative horizontal offset from the entity specified by the <code>relative</code> attribute or by the default position of the entity. The value of 1 corresponds to the default distance between entities. See a previous section for an example. Grouped entities cannot have this attribute. |
| relative | This attribute takes the name of another entity and specifies the horizontal position used as a base for the <code>pos</code> attribute. Grouped entities cannot have this attribute. |

shape This attribute takes the name of the shape you want for the entity headings. See Section 9.2.6 [Entity Shapes], page 104.

shape.size

This attribute specifies the size of the shape to use for the entity headings. Only has effect if a valid shape is specified via the **shape** attribute. It takes one of **tiny**, **small**, **normal**, **big** or **huge** with **small** as default. You can also use **auto**, which makes the shape large enough to contain the label. This setting is valid only for shapes that define a position for the label. (See the **T** command in Section 8.10 [Defining Shapes], page 92.) For shapes that define a label position if you set a specific size, the label's size is scaled to fit the shape (well, the label position in it). In contrast, if you set **auto**, the size of the shape is scaled so that it fits the size of the label.

collapsed

This attribute can be used to collapse a group entity. Only group entities can have this attribute.

indicator

If set to yes (default) a small indicator will be displayed in a collapsed entity and also for any arcs that disappeared because of the collapse of this entity. On non-collapsed group entities it has no effect. Only grouped entities can have this attribute.

The following attributes can be changed at any location and have their effect downwards from that location.

show This is a binary attribute, defaulting to yes. If set to no, the entity is not shown at all, including its vertical line. This is useful to omit certain entities from parts of the chart where their vertical line would just crowd the image visually. The commands ‘**show**’ and ‘**hide**’ followed by entity names are shorthands for setting or clearing these attributes. See more on entity headings in Section 9.2.5 [Entity Headings], page 103.

active This is a binary attribute, defaulting to no. If set the entity line becomes a thin long rectangle indicating that the entity is active. You can set the fill of the rectangle via the ‘**vfill.***’ attributes. The commands ‘**activate**’ and ‘**deactivate**’ are shorthand for setting or clearing this attribute. Using the keywords is equivalent to setting the attributes, except that when the keywords are used just after an arrow, the activation/deactivation will take place immediately at the tip of the arrow, and not after.

color This sets the color of the entity text, the box around the text and the vertical line to the same color. It is a shorthand to specify **text.color**, **line.color** and **vline.color** to the same value.

```

line.*
vline.*
fill.*
vfill.*
text.*

shadow.* See Section 8.6 [Common Attributes], page 85, for the description of these
           attributes.

aline.*
atext.*
arrow.* These attributes can be used to set the line, text and arrowhead attributes of
           arrows starting from this entity. Their use is described in Section 9.3.3 [Arrow
           Appearance], page 112. Note they do not apply to block arrows.

```

9.2.4 Implicit Entity Definition

It is not required to explicitly define an entity before it is used. Just typing the arrow definition `a->b;` will automatically define entities ‘`a`’ and ‘`b`’ if not yet defined. This behaviour can be disabled by specifying the `--pedantic` command-line option or specifying `pedantic=yes` chart option. See Section 9.11 [Chart Options], page 139. Disabling implicit definition is useful to generate warnings for mis-typed entity names².

Implicitly defined entities always appear at the very top of the chart. If you want an entity to appear only later, define it explicitly.

9.2.5 Entity Headings

By default, when an entity is defined, its heading is drawn at that location. If the entity name is preceeded by the `hide` keyword or the `show=no` attribute is specified at the entity definition then the entity heading is not drawn at the location of the definition. It is drawn later, if/when the entity is turned on by using `show` followed by the entity name or by setting `show=yes`. Note that multiple entities can be listed after both `show` and `hide`. It is also possible to specify other attributes for entities after these keywords.

Mentioning an entity after its definition either preceeded by `show` or with `show=yes` will cause an entity heading to be drawn into the chart even if the entity is already shown. This can be useful for long charts, see Section 4.2 [Defining Entities], page 19, for examples.

You can display all of the entity headings using the `heading;` command, as well. This command displays an entity heading for all (currently showing) entities. This may be useful after a `newpage;` command, see Section 9.14 [Commands], page 150. However, the best practice is to use ‘`newpage [auto_heading=yes];`’ instead, since it only shows the heading when the chart is viewed per-page (which is the same for page breaks inserted by automatic pagination).

Using the `show` and `hide` commands in themselves with no entity specified after, will make *all* entities show or hide. (Same with `activate` and `deactivate`.) See the example below.

² To this end, color syntax highlighting underlines an entity name appearing the first time. This allows quickly realizing if the name of an entity is misspelled.

```

msc+=feng_shui_notes;
B->C->D;
D->>A;
B->A;
heading;
note at A: All currently showing
entities are displayed.;
note at E: 'E' included, since it is
implicitly declared later
and shown already here.;


B->A;
hide C, B;
note: Turning B & C off;
---;
show;
note: Showing again entities
that were turned off;
A->D;
hide;
show A, D;
note: Turning all, but A and D off;
A->E;
note at E: 'E' is implicitly defined here.;

F;
note: 'F' appears here due to being
explicitly defined here.;

show E;
E->F;
show B;
nudge;
show C;
note: 'nudge' made B and C show separately.;

B
C
D
A
E
F

```



Note that entity related commands are merged. Thus saying `show; hide A;` will show all entities, except A in a single line of entity headings. If you want to manipulate entities in separate visual lines, insert a `nudge;` command in between the entity manipulating commands (as at the end of the example above).

9.2.6 Entity Shapes

The shape of an entity heading can be altered from the default box-like appearance to something custom using the `shape` attribute. Its value is a string, the name of the shape. The actual appearance of shapes is defined in separate files. Msc-generator comes with a few default shapes (their name all start with `def.`), but you can define your own shapes or add third-party Defining Shapes. See Section 8.10 [Defining Shapes], page 92, for more.

For some shapes, the label of the entity is written inside the shape, for some it is written below. This is decided by the author of the shape. If the label is written inside, it is scaled to fit. You can influence the size of the shape via the `shape.size` attribute, which takes the values `tiny`, `small`, `normal`, `big`, `huge` or `auto` with `small` as default.

Note that the above two attributes can be set in the `entity` and the `entitygroup_collapsed` style, which will influence all entities in a chart at once.

9.3 Specifying Arrows

Arrows are probably the most important elements in a message sequence chart. They represent the actual messages. Arrows can be specified using the following syntax.

```
entityname arrowsymbol entityname [attr = value | style, ...];
```

`arrowsymbol` can be any of '`->`', '`<-`' or '`<->`', the latter for bidirectional arrows. `a->b` is equivalent to `b<-a`. This produces an arrow between the two entities specified using a solid line. Using '`>`/`<>`', '`>>`/`<<>>`' or '`=>`/`<=>`', will result in dotted, dashed or double line arrows, respectively. These settings can be redefined using styles, see Section 8.8 [Defining Styles], page 90. There are two more arrow symbols for you to re-define: '`==>`' and '`=>>`' (with reverse and bi-directional variants, too: '`<==>`' and '`<<=>>`'). These default to double lines, with the latter having sharp arrowheads, as well.


It is possible to omit one of the entity names, e.g., `a->;`. In this case the arrow will expand to/from the chart edge, as if going to/coming from an external entity. Using the pipe symbol '`|`' as start or endpoint instead will still make the arrow go to (or come from)

no specific entity, but not at the very end of the chart, but immediately after the last entity (or before the first one).

```

hscale=auto;
a,b,c,d;
a->b-|: Goes to out there;
|->c-d: Comes from out there;
->d start before b:
  Start further before
  [arrow.starttype=dot];
---;
a->b: Trigger Msg;
b-> end after b;
b-> end after b +10;
b-> end after b +20;
---;
b<- start before b;
b<- start before b +10;
b<- start before b +20;
b->a: We got them all;

```



You can also use the `start before` and `end after` keywords after the arrow specification to precisely indicate where the arrow should start and/or end. You need to specify an entity afterwards and optionally an offset. The terms ‘before’ and ‘after’ should be understood in the direction of the arrow, so if an arrow is right-to-left, then `start before X` makes the arrow start right of entity X. A positive offset moves the arrow endpoint even further from the entity specified. (In the example above `start before b +10` would move the start of the arrow even more to the right. The offset is specified in pixels. The default distance from the entity (to which offset is added) is 30 pixels times the value of the `hscale` (and exactly 30 for `hscale=auto`).

It is possible to specify multi-segment arrows, such as `a->b->c` in which case the the arrow will expand from ‘a’ to ‘c’, but an arrow head will be drawn at ‘b’, as well. This is used to indicate that ‘b’ also processes the message indicated by the arrow. The arrow may contain any number of segments, and may also start and end without an entity, e.g., `->a->b->c->d->;`. As a syntax relaxation, additional line segments can be abbreviated with a dash (‘-’), such as `a<=>b-c-d;;`. Subsequent segments inherit the line type and direction of the first one. This enables quick changes to these attributes with minimal typing, as only the first arrow symbol needs to be changed. As a further possibility, different arrow symbols can also be used for different segments, such as `a->b=>c>>d-e;;`, but all the arrow symbols must be of the same direction. It is therefore not possible to mix arrows of different directions, such as `a->b<-c;;` or `a->b<->c;;`. Note that specifying different arrow symbols affect only the line attributes of the segments by default, but they can also impact the arrowhead, text or other attributes, see Section 9.3.3 [Arrow Appearance], page 112, below.

If the entities in a multi-segment arrow are not listed in the same (or exact reverse) order as in the chart, Msc-generator gives an error and ignores the arrow. This is to protect against unwanted output after rearranging entity order.


Arrows can also be defined starting and ending at the same entity, e.g., `a->a;`. In this case the arrow will start at the vertical line of the entity and curve back to the very same line. Such arrows cannot be multi-segmented.

Only non-grouped entities can be used in an arrow definition. If an entity used to define an arrow is not shown due to the collapse of its group entity, Msc-generator will automatically use the collapsed group entity when drawing the arrow, instead. If the arrow becomes degenerate (spanning between only a single collapsed group entity) or disappears entirely, an indicator will be shown instead, if the ‘`indicator`’ attribute of the collapsed group entity was set to yes (default).

9.3.1 Lost Messages

Sometimes one want to express that a message is lost. You can do this in Msc-generator in two ways. Either, you can add an asterisk between the two entities where the message is lost; or you can add a ‘lost at’ clause after the arrow specification before the label or attributes. This causes a small x to be drawn at the place specified and the dimming of the remainder of the arrow.

```
a, b, c, d;
a*->b;
a->*b;
note: Lost between
neighbouring entities;
a->b*->d: \p1Lost after b;
a->b->*d: \p1Lost before d;
a->b->d lost at b++:
\p1Lost just after b;
a->b->d lost at d -30:
\p1Lost 30 pixels left of d;
```



The first three ones are the quick one. The message lost will be indicated around the entity after or before the asterisk. Specifically, it will be between this and its neighbouring visible entity. If that visible entity is also part of the arrow, the loss will be at one third the distance between them, else it will be halfway. Using the second form, you can specify exactly where the loss happened. It can be placed onto an entity, left or right from it, or between two entities. These are specified as ‘`lost at <entity>`’, ‘`lost at <entity>-`’, ‘`lost at <entity>+`’ or ‘`lost at <entity1>-<entity2>`’, respectively. You can add two plus or minus symbols to increase distance. You can also specify any offset in addition by adding a number after, such as in ‘`lost at <entity> <number>`’. The number will be interpreted in pixels and shifts the vertical left or right depending on its sign.

The appearance of the loss symbol (the x) can be influenced using the `x.line.width`, `x.line.color` and the `x.size` attributes. The latter takes the same values as arrowhead sizes: `tiny`, `small`, `normal`, `big` or `huge`, with `small` as default.

The appearance of the lost portion of the message can also be influenced via the `lost.text.*`, `lost.line.*` and the `lost.arrow.*` attributes. Anything specified here will be added to the text, line and arrowhead format of the arrow. Currently only the color of

the line and the arrowhead is overlaid with `++white,128` (for plain designs), making them weaker, but you can change to dash lines, specify a narrower line or empty arrowheads.

9.3.2 Arrow Attributes

Arrows can have the following attributes.

label	This is the text associated with the arrow. See Section 8.1 [Labels], page 75, for more information on how to specify labels. In Msc-generator the first line of the label is written above the arrow, while subsequent lines are written under it. Future versions may make this behaviour more flexible.
text.*	All text formatting attributes described in Section 8.6 [Common Attributes], page 85, can be used to manipulate the appearance of the label ³ .
number	Can be set to <code>yes</code> , <code>no</code> or to a number, to turn numbering on or off, or to specify a number, respectively. See Section 8.4 [Numbering], page 80.
refname	Can be set to any string and is used to give a name to the arrow, which can be used to reference this arrow. Use the <code>\r(name)</code> escape in labels to insert the number of the referenced arrow. See Section 8.4 [Numbering], page 80.
compress	Can be set to <code>yes</code> or <code>no</code> to turn compressing of this arrow on or off. See Section 9.10.3 [Compression and Vertical Spacing], page 136.
vspacing	Can be set to a number interpreted in pixels or to the string <code>compress</code> . Governs how much vertical space is added before the arrow (can be negative). This attribute is another form (superset) of the <code>compress</code> attribute; <code>compress=yes</code> is equivalent to <code>vspacing=compress</code> , whereas <code>compress=no</code> is equivalent to <code>vspacing=0</code> .
angle	This takes a number in degrees and makes the arrow slanted. Arrows pointing to the same entity cannot have such an attribute. This attribute takes its default value from the <code>angle</code> chart option (or is zero in the absence of such an option, which corresponds to horizontal arrows).
slant_depth	This is similar in effect to the <code>angle</code> attribute, but instead of specifying the angle of the slant in degrees, you can use this attribute to specify how many pixels shall the end of the arrow be below the start of it. If you specify both the <code>angle</code> attribute and <code>slant_depth</code> the latter takes effect.
color	This specifies the color of the text, arrow and arrowheads. It is a shorthand to setting <code>text.color</code> , <code>line.color</code> and <code>arrow.color</code> to the same value.
line.color, line.width	Set the color and the width of the line, see Section 8.6 [Common Attributes], page 85.

³ A special note on left and right text margins (to be specified via `\ml()` and `\mr()` escapes). Msc-generator always adds enough text margins to prevent the label to overlap with the arrowhead. Thus, if you specify less margin, it will have no effect.


line.corner

This attribute specifies how the line shall be drawn at corners. It impacts boxes and entities drawn with this line, for arrows it is effective for arrows that start and end at the same entity. Its value can be `none`, `round`, `bevel` or `note`. See the example below. Setting `line.corner` without `line.radius` will result in the default radius of 10.

line.radius

For arrows starting and ending at the same entity, this specifies the roundness of the arrow corners. 0 is fully sharp (equivalent to `line.corner=none`, positive values are meant in pixels, a negative value will result in a single arc (for any corner setting). If only `line.radius` is set and not `line.corner` the result will be a round corner.

```
{
  A->A: Radius=10 [line.radius=10];
  A->A: Radius=5 [line.radius=5];
  A->A: Radius=0 [line.radius=0];
} {
  B->B: Bevel [line.corner=bevel];
  B->B: None [line.corner=none];
  B->B: Radius=-1 [line.radius=-1];
};
hspace A-B 150;
```

**arrow.size**

The size of the arrowheads. It can be `epsilon`, `spot`, `tiny`, `small`, `normal`, `big` or `huge`, with `small` as default.

arrow.color

The color of the arrowheads.

arrow.type

Specify the arrowhead type. The values can be `half`, `line`, `empty`, `solid`, which draw a single line, a two-line arrow, an empty triangle and a filled triangle, respectively. The above 4 types also exist in `double` and `triple` variants, which draw two or three of them. `sharp` and `empty_sharp` draws a bit more pointier arrowhead, filled or empty, respectively. Even more pointed are `vee` and `empty_vee`. `diamond` and `empty_diamond` draws a filled or empty diamond, while `dot` and `empty_dot` draws a filled or empty circle. The last four has non-symmetric versions, where the dot or diamond is not drawn on the entity line, but before it: `nsdiamond`, `empty_nsdiamond`, `nsdot` and `empty_nsdot`. Specifying `none` will result in no arrowhead at all. This attribute sets both the `endtype` and `midtype`, see below.

arrow.endtype

Sets the arrow type for arrow endings only. This refers to the end of the arrow, where it points to. In case of bidirectional arrows, both ends are drawn with this type. It defaults to a filled triangle.

arrow.midtype

This attribute sets the arrowhead type used for intermediate entities of a multi-segment arrow. It defaults to a filled triangle.

arrow.skiptype

This attribute specifies what to draw for entities the arrow passes over but does not stop at. E.g., if we have three entities ‘a’, ‘b’ and ‘c’, then an arrow specified as `a->c`, will pass over entity ‘b’. This attribute defaults to none, but another suitable value is `jmpover`, which draws a small half-circle.


arrow.starttype

This attribute sets the arrowhead type used at the starting point of an arrow. It defaults to no arrowhead.

arrow.xmul**arrow.ymul**

These attributes change the width or the height of the arrowhead. The default value is ‘1’. They are multipliers, thus the value of ‘1.1’ results in a 10% increase, for example.

```
hscale=auto, compress=yes;
{
  a->b: solid [arrow.type=solid];
  a->b: empty[arrow.type=empty];
  a->b: line[arrow.type=line];
  a->b: half[arrow.type=half];
  a->b: sharp[arrow.type=sharp];
  a->b: empty_sharp[arrow.type=empty_sharp];
  a->b: vee[arrow.type=vee];
}
{
  c->d: double[arrow.type=double];
  c->d: triple_empty[arrow.type=triple_empty];
  c->d: double_line[arrow.type=double_line];
  c->d: triple_half[arrow.type=triple_half];
  c->d: xmul=1.5 [arrow.type=empty, arrow.xmul=1.5];
  c->d: xmul=1.5, ymul=0.7 [arrow.type=double,
                                arrow.xmul=1.5, arrow.ymul=0.7];
  c->d: empty_vee[arrow.type=empty_vee];
};
a->b-c-d: \pldiamond [arrow.midtype=diamond];
a->b-c-d: \plempty_diamond [arrow.midtype=empty_diamond];
a->b-c-d: \pldot [arrow.midtype=dot];
a->b-c-d: \plempty_dot [arrow.midtype=empty_dot];
a->b-c-d: \plnsdiamond [arrow.type=nsdiamond];
a->b-c-d: \plnsdot [arrow.type=nsdot];
a->b-c-d: indicating skipped [arrow.skiptype=jumpover];
```



```
arrow.gvtype
arrow.gvendtype
arrow.gvmidtype
arrow.gvskiptype
arrow.gvstarttype
```

Use these attributes to set the corresponding arrowhead using the graphviz-style arrowhead specification syntax. This syntax is a superset of what graphviz supports and is not really meaningful for block arrows. Setting any of **arrow.*type** or **arrow.gv*type** will erase the other attribute, since these govern the same property of the arrow.

You have a set of basic forms **normal** (a regular filled triangle), **inv** (filled triangle in inverse direction), **vee** (a v shape), **crow** (v shape in the reverse direction), **tee** (a thick line perpendicular to the arrow line), **tick** (a thin line perpendicular to the arrow line), **box** (a square), **diamond**, **dot**, **curve** (a half circle line) and **icurve** (in the reverse direction). Msc-generator supports the following additional basic forms: **sharp**, **line** and **jumpover**, which are the same as for the ‘**arrow.type**’ attribute above; and **sbox**, **sdiamond** and **sdot**, which are symmetric versions of **box**, **diamond** and **dot**, respectively.⁴

You can combine the basic forms by appending them, such as **normalnormal**, which is a double filled triangle. You can create arbitrary long combinations. The first form specified will be the tip of the arrow closest to the entity line with the remaining forms following behind.

Some of the forms may be prepended by the letter **o** to create a non-filled version, such as **odot** or **onormal**. You can use the **l** or **r** letters to draw only half of the arrowhead which falls on one side of the arrow line, such as **lbox**, **rcurve** or even **olnormal**.

Inserting a size modifier will change the size of the subsequent forms, such as **spottedsmallnormal**, will draw a smaller dot and a larger triangle. Since the term ‘**normal**’ is both a size and a type, if you mean the size, use **sizenormal**. In addition to size modifiers, you can also insert any color name. Finally, you can also insert one of the following attributes: **lwidth** to change the line width of the arrowhead; **ltype** to change the line type of the arrowhead; **color** to change the color of the arrowhead (e.g., to an rgb value); **xmul**, **ymul** or **mul** to apply a (cumulative) multiplier to the x or y dimension of the arrowhead or to both. These attributes must be terminated with a pipe symbol (**|**) and when using them the attribute value must be enclosed in quotation marks, such as **arrow.gvtype="xmul=0.2|normal|color=12,34,56|dot"**. See the examples below.

Note that you can use these attributes to set arrowhead style for block arrows, but in that case only one form can be given (you can use **stripes** and **triangle_stripes**, but cannot use **line**, **curve**, **icurve**, **jumpover** and **crow**). Also, line/color attributes have no effect. All in all, you are better off not using the graphviz syntax for block arrows.

⁴ This is because **box**, **diamond** and **dot** in graphviz are drawn before the target, in conflict with Msc-generator behaviour for ‘**arrow.type**’. Thus, if you specify ‘**arrow.type=dot**’, you will get a dot centered on the arrow line; while if you use ‘**arrow.gvtype=dot**’, you will get one before the entity line.


```

hscale=auto;
{
    compress=yes;
    defstyle arrow [arrow.color=purple];
    a->b [arrow.gvtype=vee]: normal;
    b->c [arrow.gvtype=inv]: inv;
    a->b [arrow.gvtype=sharp]: sharp;
    b->c [arrow.gvtype=vee]: vee;
    a->b [arrow.gvtype=crow]: crow;
    b->c [arrow.gvtype=tee]: tee;
    a->b [arrow.gvtype=curve]: curve;
    b->c [arrow.gvtype=icurve]: icurve;
    a->b [arrow.gvtype=box]: box;
    b->c [arrow.gvtype=sbox]: sbox;
    a->b [arrow.gvtype=diamond]: diamond;
    b->c [arrow.gvtype=sdiamond]: sdiamond;
    a->b [arrow.gvtype=dot]: dot;
    b->c [arrow.gvtype=sdot]: sdot;
    nudge;
    block a->b [arrow.gvtype=vee]: vee;
    block b->c [arrow.gvtype=tee]: tee;
    block a->b [arrow.gvtype=inv]: inv;
    block b->c [arrow.gvtype=sbox]: sbox;
    block a->b [arrow.gvtype=diamond]:
        gv: diamond;
    block b->c [arrow.gvtype="xmul=0.5|odot"]:
        gv: odot, xmul=0.5;
};

---;

defcolor brown=139,69,19;
a<->b-c [arrow.gvtype=tinyreddotsizenormalblue];
a<->b-c [arrow.gvtype="smallredwidth=0.5|osdot"];
a<->b-c [arrow.gvtype=smallnormalbigredicurvered|
            arrow.gvstarttype=graytee];
a<->b [arrow.gvendtype="tiny|color=0,192,0|lineline"];
join b->c [arrow.size=normal, arrow.gvstarttype=normal];
a->b-c [arrow.size=normal, arrow.gvtype=normal];
a<->b-c [arrow.gvtype=bluetinynormalgreensmall];
a->b-c [arrow.gvendtype=llinegrayllinelgraylline|
            arrow.gvmidtype="lwidth=3|llinelline"];
a->b-c [arrow.size=normal, arrow.gvendtype=orange];
a<->b-c [arrow.size=normal, arrow.gvtype=black];

```



lost.text.*

The values specified here will be added to the values of `text.*` when drawing the label of the lost part of the message.

lost.line.*

The values specified here will be added to the values of `line.*` when drawing the line of the lost part of the message.

lost.arrow.*

The values specified here will be added to the values of **arrow.*** when drawing the arrowheads in the lost part of the message.

x.size The size of the loss symbol for lost messages. It can be **tiny**, **small**, **normal**, **big** or **huge**, with **normal** as default.

x.line.width

x.line.color

The linewidth and color of the loss symbol for lost messages.

Note that default values can be changed using styles, see Section 8.8 [Defining Styles], page 90.

9.3.3 Arrow Appearance

There are a number of factors influencing how an arrow appears. In this section we describe the order of precedence among them. See Section 8.8 [Defining Styles], page 90, for more information on styles.

1. Attributes specified after the arrow definition always take precedence.
2. If an attribute is not specified after the arrow definition, it takes its value from a refinement style, such as **->** or **=>**. In case of multi-segment arrows, which use multiple types of arrow symbols, such as **a->b=>c** each refinement style is applied in the order the arrow points. Thus in the case of **a->b=>c** any attribute specified in style **=>** takes precedence over the value (if any) specified in **->** for the same attribute. The same order (the order towards the destination end of the arrow) is applied even if the arrow is specified backwards, as **c<=b<-a**. For bidirectional arrows, the order is as written in the input file. Line attributes are special, because they are also recorded and applied to the respective segments. (This is why normally refinement styles only specify line attributes. Specifying the arrowhead could also make sense, but those are not applied to individual segments, but to the whole arrow.)
3. If a line, text or arrowhead attribute is not specified in any of the refinement styles applied it takes its value from the **aline.***, **atext.*** or **arrow.*** attribute of the source entity, respectively. (For bidirectional arrows, first entity written in the text file.) Since refinement styles usually overwrite the line type, it makes more sense to set e.g., line color in these attributes, if you want to automatically highlight arrows starting from a specific entity.
4. If an attribute is not specified via any of the ways above, it takes its default value from the **arrow** default style. (Chart designs change this to impact all arrows.) This style is guaranteed a value for all attributes, except **text.***.
5. Finally, unspecified **text.*** attributes take their values from the value set in the **text.*** chart options (used to globally set the font for example) or, if none set, a default font.


9.3.4 Block Arrows

When typing **block** in front of any arrow definition, it will become a *block arrow*. The label of a block arrow is displayed inside it. In addition to the attributes above, block arrows also have fill and shadow attributes, similar to entities.

All arrowheads explained above for regular arrows are supported, except the **double** and **triple** ones. In general, types with **empty** in them, draws a variant of the arrowhead

which is not taller than the body of the block arrow. The ones with `line` draw the same as the ones without. Three additional types ‘`empty_inv`’, ‘`stripes`’ and ‘`triangle_stripes`’ types are supported, as well. See the example below for a detailed list of all types supported for block arrows.


```
hscale=auto;
defstyle blockarrow [fill.color="green+80",
                     fill.gradient=right];
block a->b-c: solid [arrow.type=solid];
block a->b-c: empty(inv) [arrow.type=empty,
                           arrow.starttype=empty_inv];
block a->b-c: empty(inv) [arrow.endtype=empty,
                           arrow.starttype=empty_inv,
                           arrow.midtype=empty_inv];
block a->b-c: line [arrow.type=line];
block a->b-c: half [arrow.type=half];
block a->b-c: sharp [arrow.type=sharp];
block a->b-c: stripes
               [arrow.starttype=stripes];
block a->c: triangle_stripes
               [arrow.starttype=triangle_stripes];
block a->b-c: empty_sharp
               [arrow.type=empty_sharp];
block a->b-c: sharp, xmul=1.3
               [arrow.type=sharp,
                arrow.xmul=1.3];
block a->b-c: diamond [arrow.type=diamond];
block a->b-c: empty_diamond
               [arrow.type=empty_diamond];
block a->b-c: dot [arrow.type=dot];
block a->b-c: empty_dot
               [arrow.type=empty_dot];
block a->b-c: empty_dot, xmul=0.7
               [arrow.type=empty_dot,
                arrow.xmul=0.7];
```



If the arrow has multiple segments and the type of the inner arrowheads is either of `half`, `line`, `empty`, `solid` or `sharp` the block arrow is split into multiple smaller arrows. In this case the arrow label is placed into the leftmost, rightmost or middle one of the smaller arrows, depending on the value of the `text.ident` attribute.

It is also possible to use different arrow symbols leading to different line types, but only if the middle arrow type is such that the arrow is split into multiple contours. If not, the whole arrow is drawn with the line type of the first segment.

```
msc = round_green;
hscale=auto;
C [label="Client"], R1 [label="Router"],
R2 [label="Router"], S [label="Server"];
block C<->R1-R2-S: Query/Resp [arrow.midtype=dot];
block C<->R1-R2-S: Query/Resp\n- (block) [text.ident=left];
block C ->R1-R2-S: Query/Resp\n- (block) [text.ident=center];
block C<- R1-R2-S: Query/Resp\n- (block) [text.ident=right];
block C ->R1-R2-S: Query/Resp [arrow.midtype=diamond];
block C=>R1->R2>>S: Query/Resp;
block C=>R1->R2>>S: Query/Resp [arrow.midtype=diamond];
```



Block arrows do not accept the `arrow.skiptype` attribute and are not impacted by setting the `aline.*`, `atext.*` or the `arrow` attribute of their source entity. However, they can be slanted using the `angle` attribute.

9.4 Boxes

Boxes enable 1) to group a set of arrows by drawing a rectangle around them; 2) to express alternatives to the flow of the process; and 3) to add comments to the flow of the process. The first two uses is by adding a set of arrows to the box, while in the third case no such arrows are added, making the box *empty*.


The syntax definition for boxes is as follows.

```
[box] entityname boxesymbol entityname [attr = value | style, ...]
{ element; ... };
```

The `box` symbol is optional at the beginning of the line. The `boxsymbol` can be ‘..’, ‘++’, ‘--’ or ‘==’ for dotted, dashed, solid and double line boxes, respectively.

As with arrows the two entity names specify the horizontal span. These can be omitted (even both of them), making the box auto-adjusting to cover all the elements within. If there are no elements within and you omit one or both entities the default is to span to the edge of the chart. Specifying the entity names therefore, is useful if you want a deliberately larger or smaller box, or if you specify an *empty* box. Contrary to arrows, you can use group entities when specifying a box. The box will then cover all member entities in that group. Specifying the leftmost or rightmost member entity instead of the group entity makes a difference only if the group entity is collapsed. In the former case the box may disappear, in the latter case it will not. See the example below.

```
a, group [collapsed=yes] {
    b, c, d;
}, e;
a->b: Message;
box b--d: Box {
    b->c->d: Message;
};
d->e: Message;
d<-e: Message \#2;
box group--group: Box {
    b<-c-d: Message \#2;
};
a<-b: Message \#2;
```



Boxes take attributes, controlling colors, numbering, text indentation quite similar to arrows. Specifically boxes also have a `label` attribute that can also be shorthanded, as for arrows. For example: `... : Auto-adjusting empty box;` is a valid definition. The valid box attributes are `label`, `number`, `refname`, `compress`, `vspacing`, `color`, `text.*`, `line.*`, `shadow.*` and `fill.*`. The latter specifies the background color of the box, while `line.*` specifies the attributes of the line around. Note that `color` for boxes is equivalent to `fill.color`. `text.ident` defaults to centering for empty boxes and to left identification for ones having content.

After the (optional) attributes list, the content of the box can be specified between braces ‘{’ and ‘}’. Anything can be placed into an box, including arrows, dividers, other boxes or commands. If you omit the braces and specify no content, then you get an empty box, which is useful to make notes, comments or summarize larger processes into one visual element by omitting the details.


9.4.1 Box Series

If a box definition is not followed by a semicolon, but another box definition, then the second box will be drawn directly below the first one. This is useful to express alternatives, see the below example.

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
box C==C: Generate\ncrequest;
C=>S: Request;
box S--S: Check cache;
box S--B: Alt\#1: cache miss
    [color=1gray]
{
    S->B: Query;
    S<-B: Response;
}
...: Alt\#2: cache hit
{
    S->S: Read\ncache;
};
C<=S: Reply;

```



The subsequent boxes will inherit the fill, line and text attributes of the first one, but you can override them. The line type of subsequent boxes ('--' in the example) will determine the style separating the boxes — the border will be as specified in the first one. The horizontal size of the combined box is determined by the first definition, entity names in subsequent boxes are ignored.

Boxes can be collapsed, similar to group entities. The ‘`indicator`’ attribute governs if collapsed boxes show a small indicator to indicate that there is hidden content inside.

9.4.2 Box Tags


Boxes can also have *tags*, which are small enclosed labels at the top-left corner of the box. Tags are useful to label the content of the box, such as alternatives, loops or optional sections, while keeping the ability to add a regular box label, as well.

To specify a tag, use the `tag` attribute, and specify the label of the tag. If the label contains non-alphanumeric characters (or spaces), you need to put them in between quotation marks. You can specify the appearance of the tag via the `tag.line.*`, `tag.fill.*` and the `tag.text.*` attributes. Especially the `tag.line.corner` attribute can be used to influence the bottom right corner of the tag (but not the other corners).

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
box C==C: Generate\ncrequest;
C=>S: Request;
box S--S: Check cache;
box S--B: cache miss
  [tag="Alt\#1", color=lgray]
{
  S->B: Query;
  S<-B: Response;
}
...: cache hit [tag="Alt\#2"]
{
  S->S: Read\ncache;
};
C<=S: Reply;

```



9.5 Pipes

By typing `pipe` in front of a box definition, it is turned into a pipe. Pipes can represent tunnels, encapsulation or other associations (e.g., encryption) in networking technologies. Using them one can visually express as messages travel within the tunnels or along other associations.

Pipes take all the attributes of boxes, plus two extra ones, called `solid` and `side`. `solid` controls the transparency of the pipe. It can be set between 0 and 1 (or alternatively 0 and 255, similar to color RGB values). The value of 0 results in a totally transparent pipe: all its contents is drawn in front of it. The value of 1 results in a totally opaque pipe, all its content is "inside" the pipe, not visible. Values in between result in a semi-transparent pipe. `side` can be set to `left` or `right` and governs which side the pipe can be looked into from⁵.

For pipes the `line.radius` attribute governs, how wide the oval is at the two ends of the pipe. The default value is 5. Note that `line.corner` has no effect for pipes. Both `line.radius` and `side` can only be set on the first of the pipe segments, see below.

⁵ Beware that if you embed the chart in a Windows document, then using a lot of transparency can increase the size of the embedded object excessively.

```

msc=omegapple;
C: Client;
R1: Router;
R2: Router;
S: Server;

defstyle pipe [fill.color=rose];
defstyle pipe [fill.gradient=down];


pipe R1--R2: Tunnel [solid=0] {
    C->S: Request;
    C<-S: Response;
};

pipe R1--R2: Tunnel [solid=0.5] {
    C->S: Request;
    C<-S: Response;
};

pipe R1--R2: Tunnel [solid=1] {
    C->S: Request;
    C<-S: Response;
};

pipe R1--R2: Tunnel
    [solid=1, line.radius=10] {
    C->S: \plRequest;
    C<-S: \prResponse;
};

```



On the example above one can observe, that the last two pipes are smaller than the first two, even though they have exactly the same two arrows within. This is because in case of the first two arrows the label of the pipe itself is visible at together with the two arrows within. In contrast, the last two pipes are fully opaque so the pipe label can be drawn over its content.

Note the two `defstyle` commands before the pipes, as well. They are re-defining the default fill for pipes. You can read more about this in Section 8.8 [Defining Styles], page 90.


Similar to boxes multiple subsequent pipe definitions can be placed after each other without a semicolon. In case of boxes this results in a series of vertical connected boxes. In

case of pipes this results in a series of horizontal pipe segments besides each other. However, contrary to boxes only one set of content can be specified.

```

C: Client; R1: Router;
R2: Router; R3: Router;
R4: Router; S: Server;

pipe R1--R2: Tunnel 1 [color=red]
    R2==R3: Tunnel 2 [color=green]
    R3==R4: Tunnel 3 [color=blue, line.type=triple]
{
    C->S: \p1Request;
    C<-S: \prResponse;
}
;
```



9.6 Verticals

In contrast to arrows that span from left-to-right, a *vertical* is an element that spans up and down. It can be

- a box or a block arrow;
- a brace or bracket, to show grouping of things;
- a range to mark a time period and to comment on it; or
- a pointer which points from one point at an entity's timeline to another representing some cause/effect or a timer. This also has a version, where the pointer "lost" (e.g., to indicate a timer expiring).

```

vertical [shape] leftof|rightof [refname] [attributes...];
vertical [shape] [from] symbol [to] [at position] [attributes...];
```


First you have to specify what kind of a variant you want. The *shape* keyword can be one of **box**, **block**, **brace**, **bracket**, **range** or **pointer** or can be omitted (defaults to **block** for a block arrow). These variations result in different shapes, see the examples further down.

Then you have two methods to specify where the vertical should be placed. Using the first method the vertical is displayed left or right of another chart element. We call such verticals *by ref* verticals, as they are placed in reference to another element. That element can either be specified via its *reference name* as set by the **refname** attribute; or, if the **refname** is omitted, the element before the vertical is used. The preceding element may be a list of elements enclosed between curly braces, { and }, see the example below.⁶. Such

⁶ When you align a vertical (or symbol) to a parallel block series, which has **one_by_one_merge** layout and the vertical (or symbol) spans *above* the parallel blocks the parallel block cannot be shifted down to make

verticals (similar to *by ref symbols*, see Section 9.13.2 [Symbols], page 144) are attached to the element they refer to and layout considers them together.


```
hscale=auto;
a,_b,_c,_d;
{
    a->c: Message 1;
    c->d: Message 2;
    d->b: Message 3;
    b->c: Message 4;
    c->a: Message 5;
};
vertical leftof: A process of
                  5 messages;
---: Further procedure may follow;
```



The second method uses *markers* to specify the vertical's position. Markers can be placed with the `mark` command to mark a vertical position. (See more on marking at the end of this section.) In this case, the vertical is ignored during layout, and is placed at the position after all other elements are laid out. Horizontal space is reserved for it, unless the `makeroom` attribute is set to `no`.

The example below repeats the example above with the second method. The third line places a marker named `top` just below the entity headings. Then this marker is referenced by the vertical as the upper edge of it. The other marker is omitted in the example, it is then assumed to be the current vertical position. The `mark` command can have an `offset` attribute, which takes a number and shifts the position down by that many pixels (up for negative numbers).


```
hscale=auto;
a,_b,_c,_d;
mark top;
a->c: Message 1;
c->d: Message 2;
d->b: Message 3;
b->c: Message 4;
c->a: Message 5;
vertical top-- at a-:
    A process of\n5 messages;
---: Further procedure may follow;
```



room for the vertical as its elements are scattered around (being merged with other elements). Note that for this reason Msc-generator attempts to use `one_by_one` layout even if you specify `one_by_one_merge`, but there would be no difference - such as when the block in which the parallel block is has no sibling parallel block. If that is not possible, then the `one_by_one_merge` algorithm is used and if the attached vertical (or symbol) spans above the parallel block, a warning is emitted (which can be disabled by adding the `layout_warning=no` attribute to the vertical (or symbol)). See Section 9.9 [Parallel Blocks], page 127, for details of these two layout algorithms.

The second method is beneficial, if you cannot easily group the elements using curly braces, such as in the example below.

```
hscale=auto;
a,_b,_c,_d;
mark top;
a->c: Message 1;
c->d: Message 2;
mark mid1;
d->b: Message 3;
mark mid2;
b->c: Message 4;
c->a: Message 5;
vertical top->mid2 at a-: 3 msgs;
vertical mid1-> at d+: 3 msgs;
---: Further procedure may follow;
```




There are two built-in markers, that are available without the `mark` command: `chart.top` and `chart.bottom` referring top the top and bottom of the entire chart, respectively.

Between the two positions, one of the box or arrow symbols can be used: ‘--’, ‘..’, ‘++’, ‘==’, ‘->’, ‘=>’, ‘>’ or ‘>>’. If we omitted the *shape* specifier these symbols result in a box or a block arrow. The arrow symbols can be used in bidirectional or reverse variants, as well. For ranges and pointers the box symbols result in no arrowheads, for braces and brackets there is no difference between the box and arrow symbols, they only control the line style.

You can omit both markers. In this case the vertical spans besides the chart element before it. You can group a set of chart elements with curly braces and specify a vertical immediately after to make it span along the entire group. (This is a simpler way than to use the `mark` command.) You can even omit the *symbol* making it default to `->`⁷. The above chart can also be written as below.

```
hscale=auto;
a,_b,_c,_d;
{
    a->c: Message 1;
    c->d: Message 2;
    d->b: Message 3;
    b->c: Message 4;
    c->a: Message 5;
};
vertical box at a-:
    A process of 5 messages;
---: Further procedure may follow;
```




⁷ But you cannot omit both the *shape* and *symbol*.

Verticals can contain a label, which can be rotated 90 degrees compared to other elements. This can be set via the `side` attribute, which specifies from which direction the text is readable from (`left`, `right` or `end`, which means the regular horizontal typeset). If you sent `side=end` to typeset the label horizontally, you can use word wrapping by setting `text.wrap=yes`. If you do so (even if via the default text attribute), you must specify a text width for each such vertical to do the wrapping in. Use the `text.width` attribute for this purpose. Verticals with vertically typeset text ignore the `text.wrap` attribute and do no label word wrapping.

The text after the ‘at’ keyword determines the horizontal location of the vertical. The horizontal position is defined in relation to entity positions. It can be placed onto an entity, left or right from it, or between two entities. These are specified as ‘`<entity>`’, ‘`<entity>-`’, ‘`<entity>+`’ or ‘`<entity1>-<entity2>`’, respectively. You can also specify any distance from an entity by adding a number after the first form, such as in ‘`at <entity> <number>`’. The number will be interpreted in pixels and shifts the vertical left or right depending on its sign. Use a space before the number.

```
hscale=auto;
a, b, c, d;
...: Preceeding procedures;
{
    a->c: Message 1;
    c->d: Message 2;
    d->b: Message 3;
    b->c: Message 4;
    c->a: Message 5;
};
vertical -> at a-:
    This goes down...
vertical <->:
    ... both ways...
---: Further procedure may follow;
```



You can also omit the `at` clause, which results in the vertical being placed besides the entities it spans (by default on the right side of them using ‘`<entity>+`’). If, however, the `side` attribute is set to `right` (to mean that the text is readable from the right direction), the vertical is placed left of the entities it spans besides.

Verticals specified to be besides an entity (with ‘`<entity>-`’, ‘`<entity>--`’, ‘`<entity>+`’ or ‘`<entity>++`’ horizontal locations) are placed further from the entity line if there are boxes or elements in the way. Only those elements are considered, which are specified in the input file before vertical. If the vertical references markers below it, it may overlap with later elements, thus it is a good idea only to mark the top of the vertical and specify the vertical itself at the bottom location (as in all the examples)⁸. The `makeroom` attribute is a boolean value defaulting to yes. When it is turned off verticals are not considered when entity distances are calculated with `hscale=auto`. When `makeroom` is on, Msc-generator attempts to take the vertical into account when laying out entities. In a well-designed case


⁸ Note that Msc-generator does not lay out verticals entirely correctly in relation to parallel blocks.

you can even nest verticals, as a vertical specified earlier will be considered by subsequent verticals (but only if its `makeroom` attribute is set to yes).

```

C: Client;
S: Server;
B: Backend;
->C: Hit;
{
  box C==C: Generate\nrequest;
  {
    C=>S: Request;
    note: This must\nbe very fast;
    box S..B: Server gets info
    {
      S>>B: Query;
      S<<B: Response;
    };
    vertical brace at S-: Query;
    C=>S: Reply;
  };
  vertical bracket: Request/Response
    [line.corner=bevel];
};
vertical range: Whole process;

```




Below is a picture demonstrating all shapes of verticals. Here are a few tips on them.

The radius of the curves of the brace vertical can be adjusted with the `line.radius` attribute and defaults to 8. The width of the bracket vertical can also be influenced with the same attribuet. In addition you can set the `line.corner` attribute to `round` or `bevel` to influence the corners of the bracket. The range vertical can display either an arrow or just a simple line depending on whether you use the arrow or box symbols. In case you specify an arrow, you can adjust the arrowhead via the `arrow.*` attributes.

```

C: Client;
S: Server;
B: Backend;
{
  ->C: Hit;
  {
    box C==C: Generate\nrequest;
    vertical range -- [side=end]: Default\nrrange;
    mark top2;
    C=>S: Request;
    note: This must\nbe very fast;
    mark q_top;
    box S..B: Server gets info
    {
      S>>B: Query;
      mark middle;
      S<<B: Response;
    };
    vertical brace at S-: \-Default br;
    vertical brace at S-: \-Radius=3 b;
    [line.radius=3];
    mark bottom;
    C=>S: Reply;
  };
  vertical bracket top2=> at C-:
    Round, dbl bracket [line.corner=round];
  vertical bracket at C-: Default bracket;
  vertical pointer top2->*middle:
    Lost Pointer [lost.line.type=dashed];
  vertical pointer: Pointer [line.corner=round];
};
vertical range <>: Dotted arrow range
[arrow.endtype = empty_sharp];

```



The pointer vertical can be marked with an asterisk as being lost. If so it displays the same loss symbol that is used at lost messages (see Section 9.3.1 [Lost Messages], page 106) and you can use the `x.size` and `x.line.*` attributes to control its appearance. Note that you cannot control the exact location of the loss symbol (via the `lost at` construct), it is always at the bottom of the pointer.

Finally, about the `mark` command. You can mark a position between chart elements or the centerline of an arrow or box. This can be expressed via the `centerline` attribute (the `yes` and `destination` values are equivalent.) In order for this to work, the marker must be immediately after an arrow or box. Note that other elements do not have centerline. Distinguishing source and destination centerlines (you can also use `left` and `right` in addition to `source` and `destination`) makes sense only for slanted arrows or arrows starting and ending at the same entity. You can also add the `offset` attribute to specify certain number of pixels above (negative) or below (positive) the position. This can be used in combination with centerline, if needed.


```

hscale=auto;
compress=yes;
angle = 5;
a,_b,_c;

mark top1;
a->b: Message 1;
b->c: Message 2;
mark src2 [centerline=source];
mark dst2 [centerline=destination];
box c--c: Processing;
mark center_box [centerline=yes];
vspace 5;
c->c: Msg 3 [side=left];
mark src3 [centerline=source];
b<-c: Message 4;
mark dst3 [offset=+5];
parallel symbol arc at b
    [fill.color=aqua, draw_time=after_default];
text at b+ +6: exit condition;
a<-b: Message 5;

vertical pointer dst2>>src3:cause
    [line.radius=5, arrow.endType=empty];
vertical pointer dst3>src2 at b-:repeat
    [line.radius=5, arrow.endType=empty];
vertical range top1->center_box at c+:First part;

```



Note that in these complex cases the layout engine makes its best attempt at automatic horizontal spacing and compression, but sometimes it fails. Please report especially annoying cases.

9.7 Dividers

Dividers are called like this as they divide the chart to parts. Three types of dividers are defined. ‘---’ draws a horizontal line across the entire chart with potentially some text across it. ‘...’ draws no horizontal line, but makes all vertical entity lines dotted, thereby indicating the elapse of time.


The third type of divider ‘|||’ represents a simple vertical space. This can also be specified by entering just attributes in square brackets. The extreme ‘[] ;’ simply inserts a lines worth of vertical space. You can add text, too by specifying a label. See Section 4.3 [Dividers], page 25, for examples.

Dividers take the `label`, `color`, `text.*`, `line.*`, `compress`, `vspacing`, `number` and `refname` attributes with the same meaning as for arrows. In addition, the type of the vertical line can be specified with `vline.*`, with `vline.type` defaulting to `dotted` for ‘...’ dividers and to `solid` for ‘---’ dividers. Other values are `dashed`, `none` and `double`. Again, note that the default values can be changed by using styles, see Section 8.8 [Defining Styles], page 90.

9.8 Notes and Comments

The ‘note’, ‘comment’ and ‘endnote’ commands enable you to make annotations to the chart that are visible to the reader. Notes are placed onto the chart drawing area in a callout; comments are placed onto a column left or right from the chart; whereas endnotes are placed at the bottom of the chart. Notes are suitable for shorter comments, whereas the latter two fit longer explanations better.

```
msc += hcn;
C: Client;
S: Server;
->C: Hit;
box...: Server query
{
    box C==C: Generate\nrequest;
    note: This is\na NOTE;
    C=>S: Request;
    endnote: This is an ENDNOTE;
    C<=S: Reply;
    comment: This is a COMMENT;
};
```



Each note, command and endnote has a *target element*. The target element is the element preceding the ‘note’, ‘comment’ or ‘endnote’ command⁹. In case of notes the tip of the callout will point to the target element, whereas side notes will be typeset beside their target. You can issue multiple notes, comments and/or endnotes to the same target. If numbering is enabled for a note, comment or endnote, it inherits the numbering of its target (if any).

The syntax is simple, issue one of the three commands with attributes. You must specify a label, but similar to arrows or entities, the colon syntax can be used.

```
note: This is a note [attributes];
note at <tip>: Note pointing to <tip> [attribute];
comment: Comment text [attributes];
```

⁹ Note that some elements cannot be targets, such as chart options. In this case the preceding element becomes the target.

```
endnote: Endnote text [attribuest];
```

Note and comment text is typeset in a smaller font by default. You can change both of the above by changing the ‘note’, ‘comment’ or ‘comment’ styles.

9.8.1 Notes


For notes the tip of the callout can be guided using the `at` keyword. After it you can specify either an entity or a marker. This is useful if you want to make a note to a specific part of an arrow.

You can use the `note.pointer` attribute to define, what the tip looks like. It can take four values: `none`, `callout`, `arrow` or `blockarrow`.

The position of the note is selected automatically by Msc-generator, but you can influence the choice via the `note.pos` attribute. It can take one of the following values: `near`, `far`, `left`, `right`, `up`, `down`, `left_up`, `left_down`, `right_up` or `right_down`. The first two can be used to specify the distance from the element, whereas the rest dictate which direction the note shall be. You can set this attribute twice if needed, once for distance and a second time for direction.

The ‘note’ style contains text, fill and line attributes and also ‘`note.layout`’ and ‘`note.pos`’ to define default note layout.

```
a,b,c,d;
a->b-c-d [arrow.midtype=dot];
note at a: blockarrow [note.pointer=blockarrow];
note at b: arrow [note.pointer=arrow];
note at c: callout;
note at d: no pointer [note.pointer=none];
vspace 80;
box b--c: A nice\nlittle box;
note: right [note.pos = right];
note: left_up [note.pos = left_up];
note: far [note.pos = far];
note: near [note.pos = near];
note: down and far [note.pos=down, note.pos=far];
vspace 40;
```



9.8.2 Comments and Endnotes

Comments can be set either to the left or the right side of the chart as dictated by the `side` attribute. This attribute can also take the value `end`, which will turn the comment to an endnote. In fact endnotes are comments with their `side` attribute set to `end`. So you can convert all your comments to endnotes by redefining the `side` attribute of the ‘comment’ style, as below. For ease of use the `comment.text` and the `comment.side` chart options can also be used to set comment properties¹⁰.

```
defstyle note [text.size.normal=16, text.size.small=10];
defstyle comment [side=end];
comment.side=right;
comment.text.italics=yes;
```

When the chart contains comments on the side a line is drawn separating the comments from the chart text. You can change the properties of this line via the ‘comment.line.*’ chart options. Only the width, color and type of the line can be changed (not its radius or corner). You can turn this line off by selecting the ‘none’ line type. Similar, the background of the comments can be set via the ‘comment.fill.*’ chart options. These options can also be made part of designs. Finally, the space available on the side for comments can be adjusted with the `hspace left|right` `comment` command, see Section 9.13.1 [Spacing], page 143.


9.9 Parallel Blocks

Sometimes it is desired to express that two (or more) separate processes happen side-by-side. *Parallel blocks* allow this. Simply place the parallel blocks between ‘{}’ marks and write them one after the other, as in Section 4.5 [Drawing Things in Parallel], page 32. You can specify as many parallel blocks as you want. The last (and only the last) parallel block shall be followed by a semicolon. The order of the blocks is not much relevant, with the exception of numbering, which goes in the order the blocks are specified in the source file. It is possible to place anything in a parallel block, arrows, boxes, or other parallel blocks, as well. Below the series of parallel blocks the next element will be drawn after the longest of the parallel blocks.

¹⁰ These are equivalent to changing the ‘comment’ style. There is no such shortcut for endnotes, yet.

There are two ways to lay out parallel blocks. They differ in how they handle cases when elements from the individual blocks would overlap. For non-overlapping cases they function the same way. The first algorithm, called `one-by-one` places elements from blocks one by one always taking the next element from the block which is currently the shortest (has its bottom end the highest). Elements are placed so as to avoid overlap between them. The other algorithm, called `overlap` lays out the blocks independently and allows overlap. The algorithm to use can be selected by the `layout` attribute that can be specified for the entire parallel block series before the first block.

```
hscale=0.5;
a,_b,_c,_d;
[layout=overlap] {
    a->c: Message \#1;
    a<-c: Message \#2;
} {
    defstyle arrow [color=green];
    b->d; b->d; b->d; b->d;
};
---;
[layout=one_by_one] {
    a->c: Message \#1;
    a<-c: Message \#2;
} {
    defstyle arrow [color=green];
    b->d; b->d; b->d; b->d;
};
```




The `one_by_one` algorithm has a variant, called `one_by_one_merge`, which behaves differently in case of nested parallel blocks. If you apply this algorithm to the inner parallel block, they will be merged with the outer blocks (but not the one the nested block is inside). In contrast, `one_by_one` results in laying out the inner parallel blocks on their own as if they were a single element.

```

hscale=0.5;
a,_c,_b[pos=-0.7],_d[pos=-0.7],_e,_f[pos=-0.7];
{
  defstyle arrow [color=red];
  a->b; a->b; a->b; a->b;
  a->b; a->b; a->b; a->b;
} {
  c->e [color=blue];
  [layout=one_by_one_merge] {
    defstyle arrow [color=green+20];
    c->e; c->e; c->e;
  } {
    defstyle arrow [color=green-20];
    d->f; d->f; d->f;
  };
  c->e [color=blue];
};
vertical brace: \-Layout is\nbone_by_one_merge;
---;
{
  defstyle arrow [color=red];
  a->b; a->b; a->b;
  a->b; a->b; a->b;
} {
  c->e [color=blue];
  [layout=one_by_one] {
    defstyle arrow [color=green+20];
    c->e; c->e; c->e;
  } {
    defstyle arrow [color=green-20];
    d->f; d->f; d->f;
  };
  c->e [color=blue];
};
vertical brace: \-Layout is\nbone_by_one;

```




Note that with `one_by_one` and `one_by_one_merge` Msc-generator not only avoids overlap between elements, but in addition keeps a minimum distance between two elements. This means that arrows to/from the same entity cannot be drawn completely besides each other, since in that case they would touch. Thus one of them is drawn a little lower. If this is not intended, use `overlap`.

```

a,b,c;
...: layout=overlap {
  [layout=overlap] {
    a<-b [color=red];
  } {
    b->c [color=green];
  };
};
...: layout=one_by_one {
  [layout=one_by_one] {
    a<-b [color=red];
  } {
    b->c [color=green];
  };
};

```




If you use the `overlap` algorithm, you can also specify vertical alignment of the individual blocks via the `vertical_ident` attribute, which can be set to `top`, `middle` or `bottom`.

```

hscale=auto;
[layout=overlap] {
    a->b: msg; a->b: msg;
    a->b: msg; a->b: msg;
    a->b: msg; a->b: msg;
} [vertical_ident=bottom] {
    b->c [color=red];
    b->c [color=red];
} [vertical_ident=middle] {
    c->d [color=green];
    c->d [color=green];
} [vertical_ident=top] {
    d->e [color=blue];
    d->e [color=blue];
};

```



The default behaviour is `one_by_one_merge`, which allows fine parallelism, but avoids ugly overlaps¹¹.

You can mark entire parallel block series, with the `parallel` and `overlap` keywords (and attributes). This results in elements after the entire parallel block series to be laid out besides and over the elements in the parallel blocks, respectively. In addition, you can set the `keep_with_next` and `keep_together` attributes to influence automatic pagination.

Parallel block serieses also have `compress` and `vspacing` attributes. These govern, how they are laid out under the previous element. For block series with `layout=overlap` and `layout=one_by_one`, first the entire block series is laid out without regard to already placed elements. Then, if `compress` is on (or `vspace=compress` is set, which is equivalent) the whole block series is moved upwards as one until some parts of it bump into an already placed element. If a nonzero vertical spacing is used, the whole block series is shifted down such that the requested spacing appears between the top of the block series and the prior element.

For `layout=one_by_one_merge` with `compress=yes` the first elements of the parallel blocks are individually moved upwards until they hit one of the elements above. In case a positive `vspacing` is specified, the behaviour is the same as for `layout=one_by_one`. The `compress` and the `vspacing` attribute of the first element in each parallel block can modify this behaviour (e.g., by adding vertical spacing).

Note that if the `vspacing` chart option is set to a positive number, the `vspacing` attribute of parallel blocks is set to zero instead of this number by default. This is to avoid having this vertical space twice: once for the parallel block series and once for the first elements in the blocks. You can nevertheless assign a nonzero vertical space for the block series by manually specifying the `vspacing` attribute for the parallel block. If case of `vspacing=compress` or `compress=yes` chart options, the corresponding attribute of parallel block series is set according to the value of the chart option.

One design goal with `layout=one_by_one_merge` was that in case the parallel block series contains just one block, it should get laid out exactly as if its content were not enclosed between '{}' marks. This allows you to put '{}' marks around any set of elements, creating a new scope (see Section 8.7 [Scoping], page 89), where any changes to styles or options take effect only inside the scope.

¹¹ Setting the `classic_parallel_layout` chart option to `yes` causes the default to be `overlap`, because prior v3.6 the only algorithm available was `overlap`. This chart option is now deprecated and will be removed in future releases. Use `layout=overlap` instead.

9.9.1 Parallel Keyword

Specifying the keyword `parallel` in front of an element will make the rest of the chart be drawn in parallel with it. To be more precise the effect only lasts till the end of the scope, so elements after the next closing brace will be drawn sequentially under¹².


You can place `parallel` in front of really any element, including entity definitions or even series of parallel blocks. You can even combine several elements using braces.

```
hscale = 0.8;
parallel {
    a->b: first msg;
    a<-b: second msg;
};

box c--d: This will
            be besides;

parallel {
    c->d: third msg;
    c<-d: fourth msg;
};

box a--c: This cannot be besides;
```



9.9.2 Overlap Keyword

Sometime one explicitly wants two elements to overlap. One prime case is to show slanted messages to cross each other. You could do it via parallel blocks allowing overlap via '`layout=overlap`' but that is a bit cumbersome for overlapping short parts. A shorthand is offered by the `overlap` keyword.

Specifying this keyword in front of any element will result in ignoring this element at the layout of subsequent elements. The next element will be laid out exactly at the same vertical position as the element marked with `overlap` drawing on top of it. Subsequent are then laid out below and can still overlap the element marked with `overlap`. This effect is in place till the next closing brace (or the end of the file).


Thus this keyword is similar to the `parallel` keyword, but it allows direct overlap, not just side-by-side layout.

¹² This is how this works exactly: first, the element marked with `parallel` is placed. Then the rest of the elements in the scope are placed below it and are moved as one block up at most to the top of the element marked with `parallel`. The move stops if any element in the block being moved bumps into an already placed element, thus overlaps are avoided.

```

hscale=1.5;
angle=10;
H1: Host 1;
H2: Host 2;
overlap {
    H1->H2: \p1SYN;
    H2->H1: \prACK;
};
H2->H1: \prSYN;
H1->H2: \plACK;

```



9.9.3 Joining Arrows and Boxes

The `join` keyword can be used to connect arrows to other arrows or boxes. In its simplest form one arrow can be joined to another. If the second arrow is a continuation of the first arrow and they overlap, then a curly connecting segment is added as in case of the first example below. Note that this is done only if the two arrows have the same directionality; either the second comes after the first one or they are both bi-directional; they have a common end; and they overlap. Note that the first arrow can be an arrow starting and ending at the same entity, but the second cannot. If you want that use the `side` attribute, as illustrated on the second example below.

If the two arrows do not overlap, they will be simply laid out aligned by their centerline. It is possible to mix any type of arrow, including block arrows, irrespective of directionality. The arrows do not have to connect (but cannot overlap). This construct enables you to have a different label for each joined arrow part. Note that each joined arrow may be multi-segment itself.

Finally, you can also include boxes in a join series. You can leave the end of the arrows towards the box unspecified. In this case the arrow will connect to the side of the box. Note that you cannot join a box directly to another, an arrow must come in between.

```

A, B, C, D;
A->C: First segment;
join C->B: Tunnelerd [line.type=triple];
join B->D: Last Segment;


C->B: First;
join B->B: Second [side=left];
join B->A: Third;

A->B: Segment \#1;
join block ->: Segment \#2 [color=lgray];
join C->D->;

A->: Girl kiss;
join box B--B: Happy Boy;
join <-C: Girl kiss;

A->: Trigger;
join box B--B: State change;
join ->;
join box C--C: State change;
join ->: Message;

```



You cannot specify `parallel` or `overlap` for an element if you specify `join`. Any `parallel` or `overlap` keyword specified for a join series, on the other hand, will apply to all elements of the join series.

9.10 Signalling Chart Attributes and Styles

9.10.1 Appearance

- line.*** Specifies the appearance of the line for the element. For arrows and dividers it is the horizontal line. For block arrows, boxes, pipes and entities this is the line around the element. `line.radius` effects only on arrows starting and ending in the same entity (see Section 9.3.2 [Arrow Attributes], page 107); for entities and boxes, this specifies the size of the corners. For pipes, it specifies the width of the oval, in other words from how left we look at the pipe.
- vline.*** Specifies the color, width or type of the vertical line stemming from entities. This is useful to indicate some change of state for the entity. `vline.radius` and `vline.corner` has no effect. These attributes can be used for entities and dividers.
- fill.*** Defines the fill of the box, entity, block arrow or pipe.
- shadow.*** Defines the shadow for boxes, entities, block arrows and pipes. Currently no shadow can be specified for simple arrows.
- arrow.*** Arrowhead formatting attributes, described in detail in Section 9.3.2 [Arrow Attributes], page 107.
- lost.text.***
- lost.line.***
- lost.arrow.***
 - The values specified here will be added to the values of `text.*` `line.*` or `arrow.*` when drawing the text, line or arrowheads of the lost part of the message, see Section 9.3.1 [Lost Messages], page 106. Only applicable for arrows.
- x.size**
- x.line.*** The controls the appearance of the loss symbol for lost messages, see Section 9.3.1 [Lost Messages], page 106.
- tag.line.***
- tag.fill.***
- tag.text.***
 - These attributes apply only to boxes (applicable to the `box`, `emptybox` and `box_collapsed` style) and govern the style of tags, if the `tag` attribute of the box is set. (The `tag` attribute is not part of the style, you must set it individually on each box you want to have a tag.)
- shape**
 - This attribute takes the name of the shape you want for the entity headings. See Section 9.2.6 [Entity Shapes], page 104. They can be made part of style but have effect only on entities.

shape.size

This attribute specifies the size of the shape to use for the entity headings. Only has effect if a valid shape is specified via the `shape` attribute. It takes one of `tiny`, `small`, `normal`, `big` or `huge` with `small` as default. They can be made part of style but have effect only on entities.

note.layout

note.pos These govern how notes are laid out. See Section 9.8 [Notes and Comments], page 125, on how to use them. They can be made part of style but have effect only on notes.

side

This attribute can take either `left` or `right`. For pipes it specifies which side the pipe can be looked from into. For verticals it tells which side the text can be read from. For comments it specifies which side of the chart the comment is placed on. It has no effect on any other elements.

solid

This attribute can be used to set the transparency of a pipe. See Section 9.5 [Pipes], page 117, for more information.

number

This attribute governs if the arrow, box, etc. is numbered or not. See Section 8.4 [Numbering], page 80, for details.

compress

If this attribute is set to `yes`, the element is drawn as close to the ones above it as possible without touching those. It is useful to save space, see Section 9.10.3 [Compression and Vertical Spacing], page 136, for a detailed description.

vspacing

Can be set to a number interpreted in pixels or to the string `compress`. Governs how much vertical space is added before the element (can be negative). This attribute is another form (superset) of the `compress` attribute; `compress=yes` is equivalent to `vspacing=compress`, whereas `compress=no` is equivalent to `vspacing=0`.

collapsed

This attribute can be used for group entities and boxes to collapse them.

indicator

If this is set to yes on a collapsed group entity or box, indicators will show hidden entities and other chart elements.

The attributes below can be specified for most elements, but cannot be made part of a style

label

This gives the label of the element (for elements having one). It can be abbreviated with the colon notation, see Section 8.1 [Labels], page 75.

url

This assigns a link target to the label, such as an URL or a Doxygen target. Note that `box` tags cannot be turned into a link using this attribute, use the `\L()` escape instead. See Section 8.3 [Links], page 79, for more info.

refname

Use this attribute to name the element for later reference. Used primarily to refer to elements via their numbers using the '`\r(name)`' escape in labels.

draw_time

Use this attribute to draw elements earlier or later and thereby control how they overlap. See more in Section 9.13.2 [Symbols], page 144.

parallel This can take a `yes` or a `no` and is equivalent to prepending the element with the `parallel` keyword, see Section 9.9 [Parallel Blocks], page 127.

9.10.2 Word Wrapping and Long Labels

Before Msc-generator 3.6 the user was required to manually specify line breaks in labels. Using the `text.wrap` attribute you can instruct Msc-generator to break lines automatically depending on how much horizontal space is available. For labels with this attribute set the line breaks of the source file inside the label are ignored. However, the line breaks inserted into the label via the `\n` escape sequence are still honoured. You can set the `text.wrap` attribute of labels globally via the `text.wrap` chart option, but you can also override this setting individually for each label.

Word wrapping is not available for graphviz graphs.

9.10.2.1 Word Wrapping for Signalling Charts

You cannot set this attribute for signalling chart entities. Their label is always typeset with `text.wrap=no` exactly as you specify in the source file.

Note that this feature is most useful if you do not use automatic horizontal scaling `hscale=auto`, since in that case the distance between entities is determined from the size of the labels - and with `text.wrap=yes` there is no inherent size for most labels. For notes, which float and whose width is not determined by the spacing of entities, a new `width` attribute is inserted, which can specify the width of the note making word wrapping meaningful.

You have effective 3 easy way to typeset long labels.

- Word wrapping: Use `text.wrap=yes` (and perhaps a fixed `hscale`), in this case the long labels wrap into multiple lines.
- Automatic scaling: Use `hscale=auto` and no word wrapping, in this case entities are spaced apart, so that there is enough space for all labels.
- None: No word wrapping or automatic scaling (default): long labels expand beyond their available space, which may be sometimes ugly.


You can some combinations, as well.

- Even with `hscale=auto` you can make some long labels word wrap by applying `text.wrap=yes` only to the specific arrow, box, divider or comment. Specifying a long label with word wrapping will not cause entities to be spaced apart to make room for it, but instead the label is typeset into the space available (determined by other labels). Adding horizontal spacing with the `hspase` command can be applied to manually push entities somewhat apart (but perhaps not to the full length of the long label, which will be wrapped into the space available).
- Even with a fixed `hscale` You can push entities further by using the `hspase` command and thereby make enough room for a long label. You can create exactly as much as needed by using the label text as the argument for `hspase`, see below

```

hscale=1;
a,b,c,d;
a->b: A short label;
c->d: A long label needing space;
hspace c-d: A long label needing space

```



9.10.2.2 Word Wrapping for Block Diagrams

For block diagrams, this setting only takes effect if you specify a concrete size (as in pixels) for the block and not when its size is made equal to another block. Since all blocks have a default size, this is practically always the case, unless you set the size (width or height for left or right oriented labels) to the size of another block (or blocks).


Since Msc-generator always increases the size you specify to accommodate the label (and the content of the block), if you specify a too small size, word wrapping will be performed to the width necessary for the longest word in the label.

Note that if the label's orientation (set via `label_orient`) is vertical, you need to set `height` instead of `width`.

```


text.wrap=yes;
box: 1. Long label, wrapped to its
    \ulooooongest\u word {
    use top=prev, middle=;
    box: 2. Long label, \unwrapped\u
        to its longest word. ;
    box: 3. Long label, wrapped to size
        [width=150];
    below box: 4. Long label, \unwrapped\u
        to its longest word.
    [label_orient=left];
}
box: Long label, wrapped to size
    [height=150, label_orient=right];
above text: Note:
    longest words are \underlined.
    [width=40, text.size.normal=8];

```



9.10.3 Compression and Vertical Spacing

In this section we explain how Msc-generator sets the vertical space between elements. It can apply a set amount of vertical space or can use the `compression` mechanism. The latter aims to reduce the height of chart graphics by vertically pushing chart elements closer to each other. See the two examples below copied from the end of Section 4.1 [Defining Arrows], page 15. They differ only in that the second begins with `compress=yes`.



Each element (except entities) has a `compress` and a `vspacing` attribute. When the former is set to `yes`, the element is first placed fully under the element before it, then it is shifted upwards until it bumps into some already drawn element. The same effect can be achieved by using `vspacing=compress`. If compression is not used the element is placed below the previous element by the value of the `vspacing` attribute (understood in pixels). E.g., using `vspacing=10` adds 10 pixels between the element and the element before it. The `vspacing` attribute is the superset of the `compress` attribute, setting `compress` to `yes` or to `no` is equivalent to `vspacing=compress` and `vspacing=0`, respectively.

Compression and vertical spacing can be set individually for each element, but to save typing by setting the `compress` or `vspacing` chart option, you can effectively set the `compress` or `vspacing` attribute of all elements after. This is similar, how the `numbering` chart option effects the `number` attribute. If you then want to exempt specific elements from compression or add more space individually (so that they are somewhat further from the element above), just specify the `compress` or `vspacing` attribute for the element in question.

Styles can also influence compression and vertical spacing the same way as numbering, that is you can set the `compress` or `vspacing` options for a style, which will effect compression and vertical spacing of elements you assign the style to.

Note that to insert extra vertical spacing you can also use the `vspace` command, see Section 9.13.1 [Spacing], page 143.

9.10.4 Default and Refinement Styles for Signalling Charts

For signalling charts styles can contain any of the attributes listed in Section 9.10 [Signalling Chart Attributes and Styles], page 133. Styles are explained more in Section 8.8 [Defining Styles], page 90.

There is a built-in default style for each signalling chart element: `arrow`, `box`, `emptybox`, `divider`, `blockarrow`, `pipe` `entity`, `entitygroup`, `symbol`, `indicator`¹³, `title`, `subtitle`,

¹³ The style `indicator` determines the appearance of the small symbols that indicate elements hidden due to a collapsed box or entity group.

`note`, `comment`, `endnote`, `vertical`¹⁴, `vertical_brace`, `vertical_bracket`, `vertical_range`, `vertical_pointer`, `symbol`¹⁵ and `text`¹⁶.

There are also predefined styles for grouped entities and boxes for when they are collapsed: `entitygroup_collapsed`, `box_collapsed` and `box_collapsed_arrow`, the latter is used when a box is collapsed to a bidirectional arrow. Attributes of large entitygroups default to the `entitygroup_large` style.

The following refinement styles are defined further governing the appearance of elements

- for arrows: ‘`arrow_self`’, ‘`->`’, ‘`=>`’, ‘`>`’, ‘`>>`’, ‘`==>`’ and ‘`=>>`’.¹⁷
- for block arrows: ‘`block->`’, ‘`block=>`’, ‘`block>`’ and ‘`block>>`’.
- for boxes: ‘`--`’, ‘`==`’, ‘`++`’ and ‘`..`’
- for pipes: ‘`pipe--`’, ‘`pipe==`’, ‘`pipe++`’ and ‘`pipe..`’
- for dividers: ‘`---`’ and ‘`...`’
- for verticals: `vertical->`, `vertical>`, `vertical>>`, `vertical=>`, `vertical--`, `vertical++`, `vertical..` and `vertical==`.

Redefining refinement styles enables you to quickly define, e.g., various arrow styles and use the various symbols as shorthand for these. Usually style names containing non-letter characters have to be quoted, but for the above styles the parser is expected to recognize them without quotation. So both below are valid.

```
defstyle "->" [arrow.size=tiny];
defstyle -> [arrow.size=tiny];
```

Finally there are two more pre-defined styles for signalling charts: `strong` and `weak`. By adding these to any element you will get a more and less emphasized look, respectively. The benefit of these compared to making elements stronger or weaker by yourself is that they are defined in all chart designs in a visually appropriate manner. Thus you do not need to change anything when changing chart design just keep using them unaltered.

As a related comment we note that chart designs modify all the above styles and the default value for the `hscale`, `compress`, `vspacing`, `numbering`, `indicator`, `angle` and `text` chart options, too.

Thus, in summary the actual attributes of an element are set using the following logic.

1. If you specify an attribute directly at the element (perhaps via applying a style), the specified value is used¹⁸.
2. Otherwise, if the attribute is set in the refinement style (at the point and in the scope of where the element is defined), the value there is used. For arrows starting and ending at the same entity, two refinement styles are applied, first ‘`arrow_self`’, then the one based on the arrow symbol. (In the latter only line styles are set, so this is why you can use e.g., `=>` to make an arrow double-lined.)

¹⁴ referring to vertical boxes and block arrows

¹⁵ referring to all symbols

¹⁶ referring to `text at` commands

¹⁷ ‘`arrow_self`’ is applied to arrows starting and ending at the same entity after and in addition to style ‘`arrow`’. The others are applied (potentially after ‘`arrow_style`’ to line segments, also for bi-directional arrows. Thus there is no separate ‘`<->`’ style, for example.

¹⁸ If you specify the attribute several times, the last one is used.

3. Otherwise, if the attribute is set in the default style of the element, the value there is used.
4. Otherwise, the value of the applicable chart option is used, such as `text.*`, `compress`, `vspace`, `indicator`, `numbering`, `auto_heading` and `angle`. In order for these chart options to be effective default styles usually have no value specified for these attributes. You can set these attributes in styles, e.g., to set font type for empty boxes, which will take precedence over chart options.

9.11 Chart Options

Chart options are global settings that impact overall chart appearance or set defaults for chart elements. Chart options can be specified at any place in the input file, but typically they are specified before anything else. The syntax is as below.

```
option = value, ... ;
```

The following chart options are defined.

msc This option takes a chart design name as parameter and sets, how the chart will be drawn. It is usually specified as the first thing in the file before any other chart option. However, it can be specified multiple times, in which case its effect takes place downward from the chart option. If not specified then the ‘plain’ design is used. Note that this option can be overridden from the command line and also from the Windows GUI. Also note that only full designs can be applied with the ‘=’ symbol, partial designs shall use ‘+=’. See Section 8.9 [Chart Designs], page 91, for more on chart designs.

hscale This option takes a number or `auto`, and specifies the default horizontal distance between entities. The default is 1, so to space entities wider apart, use a larger value. When specifying `auto` entity positions will be automatically set according to the spacing needs of elements. In this case the `pos` attribute of entities will be ignored except when influencing the order of the entities. See the end of Section 4.2 [Defining Entities], page 19, for examples. Similar to `msc`, if you specify this attribute multiple times, the last one takes precedence.

numbering

This option takes `yes` or `no` value, the default is `no`. Any element you define will take the default value of its `number` attribute from this option. See more on numbering in Section 8.4 [Numbering], page 80.

compress This option takes a boolean value, and defaults to `off`. Any element you define will take the default value of its `compress` attribute from this option. See more on numbering in Section 9.10.3 [Compression and Vertical Spacing], page 136.

vspacing Can be set to a number interpreted in pixels or to the string `compress`. Governs how much vertical space is added before each element (can be negative). This option is another form (superset) of the `compress` option; `compress=yes` is equivalent to `vspacing=compress`, whereas `compress=no` is equivalent to `vspacing=0`.

angle Specifies the default value for arrow slanting. Its value is measured in degrees, can take values from 0 to 45 degrees and its default value is zero.

indicator

Similar to the **compress** option above this chart option can be used to influence the default value of the **indicator** attribute for grouped entities and boxes. The simplest way to turn all indicators on or off is to specify this chart option at the beginning of the file.

auto_heading

Sets the default value for the ‘**auto_heading**’ attribute of ‘**newpage**’ commands. Setting to yes will cause all ‘**newpage;**’ commands to create an entity heading on the subsequent page making additional ‘**heading;**’ commands unnecessary. The default is no.

classic_parallel_layout

If set to yes, parallel blocks are laid out with an old algorithm, which allows and ignores overlaps between the elements in the different parallel blocks. Defaults to no, and is kept only for backwards compatibility.

pedantic This option takes a boolean value. It defaults to no, but can also be set by the command line or using **Edit|Preferences...** on Windows. When turned on, then all entities must be defined before being used. If an entity name is not recognized in an arrow or box definition an error is generated. However, the implicit definition is accepted. Setting pedantic affects only the definitions after it and you can set it multiple times on and off. However it makes little sense.

text.ident**text.format****text.color****text.wrap**

This chart option can be used to set the default text format. It will be the default for all labels. Any styles or attributes specified will overwrite the formatting specified here. Its syntax is the same as that of the **text.*** attributes.

numbering.pre**numbering.post**

These options specify what shall be prepended and appended to label numbers. Their default value is the empty string and a semicolon followed by a space, respectively. The value of these options are ignored when a label number is inserted due to the ‘\N’ escape sequence. See Section 8.4 [Numbering], page 80, for more.

numbering.format

Specifies the format of automatic numbering for labels. Can be an arbitrary string (usually quoted) and may also contain formatting escapes. Any occurrence of ‘123’, ‘arabic’, ‘iii’, ‘roman’, ‘abc’, ‘letters’ (or uppercase versions) will be replaced to the actual number in the specified format. The string can contain multiple of the strings above, that will be interpreted as a multi-level numbering format. It is an error to describe more levels than the chart has at the location of the option. In this case an error is printed and the option is not changed. Describin fewer levels will result in Msc-generator omitting the

top level numbers from labels. For example, if the numbering is at 2.4.1 and one specifies ‘123.123’ for number format, Msc-generator will display only 4.1. Such truncation, however, will not change the number of levels, merely how the number is displayed.

`numbering.append`

This option can be used to append a new level to numbering. Its syntax is the same as for `numbering.format`. E.g., opening a second level of arabic numbers separated by a colon from the first level can be done by specifying ‘.123’ (use quotation marks). It is possible to add more than levels at once. All added levels start from the value of 1 (or ‘i’ or ‘a’, for roman numbers or letters, respectively).

`numbering.increment`

Sets the amount added to the number for every new numbered element. Specifying negative values will make counting go backwards.

`background.color`

`background.gradient`

These are similar to `fill.*` attributes and specify the background color of the chart. By default the background is set to white in the *plain* design of every language. You can change the background color multiple times, each change taking effect at the place where you issue the background chart option. This is useful to split your chart to multiple sections visually. By setting `background.color=none` the background will be transparent for the rest of the chart.

`file.info`


It takes a (quoted) string of human-readable text as value. It is useful to describe what is this file and what it contains. It is used so far only to annotate design libraries, so that if you open an OLE object with a shape not present in your system you can get some info on what file it is from. You can specify this option multiple times their values get concatenated.

`file.url` It takes a quoted URL as value providing a potential place to download this file from.

```

compress=yes;
C: Client;
S: Server;
B: Backend;
background.color="blue+90";
->C: Hit;
C=>S: Request;
S>>B: Query;
S<<B: Response;
C=<S: Reply;
background.color=none;
...:\i Some time elapses;
background.color="green+90";
->C: Hit [compress=no];
C=>S: Request;
S>>B: Query;
S<<B: Error [color=red];
C=<S: Error [color=red];

```



```

comment.line./*
comment.fill./*

```

If you have comments on the chart these govern the background of the comments and the attributes of the line separating the comments from the chart. As with background changing them applies downwards from the point of the chart option. See Section 9.8 [Notes and Comments], page 125, for more information on comments.

9.12 Multiple Pages

Msc-generator supports multi-page charts. These may be useful when you want to print a long chart. Also, when you only want to show some parts of a chart in a compound document, but want to keep the rest of the text, too. In the latter case just put the parts to show on a different page and show only that page in the compound document.

By default the whole chart is a single page. The chart can be manually broken into multiple pages by inserting ‘`newpage;`’ commands. The chart then can be viewed either as a whole or page by page. You can have as many pages in a document as you want. Adding the ‘`[auto_heading=yes]`’ option to the command will result in displaying an automatic entity heading at the top the page after the page break - but only when the chart is viewed page-by-page. If you want this for all such manually inserted, simply set the ‘`auto_heading`’ chart option to yes.

You can also make Msc-generator to paginate the chart for a given page size. On the command line this is available via the ‘`-p -a`’ options, on Windows, there is a checkbox on the ribbon. You can ask Msc-generator to insert headings to the top of the new pages by specifying ‘`-ah`’ or ticking the ‘Auto Headings’ checkbox.

The command-line version of Msc-generator creates as many output files as many pages there are. If there is more than one page, it appends the page number to the filename you specify. Specifying the ‘`-p`’ option for PDF output allows you to have a single, multi-page output file. In the Windows GUI if you export from Print Preview to PDF, a sinlge multi-page file is created using the page size, orientation, margins and alignment selected in Print Preview.

9.13 Free Drawing

Sometimes one wants to add simple drawing elements to a chart, such as circle an arrowhead or comment, dots or other shapes. Msc-generator supports naturally only limited drawing capabilities, but here they are.

9.13.1 Spacing

Arbitrary vertical space can be added using the `vspace` command.

```
vspace number [attributes];
vspace: label [attributes];
```

In the first form the vertical space is specified as a number in points. In the secod form, the height of the given label will be used. This command also has a specific attribute, called `compressable`, which specifies if the space should be ignored if compress is on. It defaults to `no`.

Horizontal spacing between the entities can be controlled either via the `pos` and `relative` entity attributes or can be made fully automatic by specifying `hscale=auto;`, see Section 9.2.1 [Entity Positioning], page 99, and Section 9.11 [Chart Options], page 139.


The `hspace` command is useful in the latter case to force a certain horizontal distance between two (not necessarily neighbouring) entity. The space can be larger than the one specified with `hspace` if the layout requires so, but never smaller.

```
hspace entity-entity number [attributes];
hspace entity-entity: label [attributes];
hspace left comment number [attributes];
hspace right comment number [attributes];
```

The syntax is similar to that of the `vspace` command, both a number or a label can be used to specify the horizontal distance. Before the distance, the two entities need to be specified. Any one can be omitted, in this case the distance is proscribed between the edge of the chart and the entity¹⁹. Two special versions of the `hspace` command exist to specify the spacing for the comments on the right and left sides.

The `hspace` command can be specified anywhere in the file with the same effect.

```
hscale = auto;
a, b, c, d;
a->b:
    This far apart;
hspace a-b:
    This faaaaaaaaaar apart;
a--d;
vspace 40;
a--d;
hspace c-d 40;
```



¹⁹ Note that the edge will not be the physical edge, merely the invisible line from which arrows connect to when only one entity is specified, such as `a->`; or `->a;`.

9.13.2 Symbols

Currently Msc-generator can draw circles (ellipses), ellipses (three dots), cross symbols (a big ‘X’), arbitrary shapes, rectangles (optionally with text) or just plain text. We call these *symbols*.

```
symbol symbol_type rel_pos [refname] [attributes];
symbol symbol_type at entity [attributes];
symbol symbol_type marker-marker hpos1 hpos2 [attributes];
```

By specifying either `arc`, `rectangle`, ..., `text`, `cross` or `shape` after the `symbol` keyword one instructs Msc-generator to draw one circle/ellipsis, rectangle, ellipses, plain text²⁰, a cross symbol or a shape, respectively. The position of the symbol can be specified three ways as above.

The first (most natural) method places the symbol relative to another element. This element can be specified using its `refname` attribute. If the *reference name* is omitted, then the last element before the symbol is used. (This is similar to how Verticals can be laid out, see Section 9.6 [Verticals], page 119.) The `rel_pos` can be one of `leftof`, `rightof`, `above`, `below` or `around` specifying what relation the symbol shall be to the referenced element. The `around` option will also size the symbol to encompass all of the referenced element, useful for `cross`, `rectangle` or `arc`. Using this placement method makes the symbol a *by ref* symbol. Such symbols (and verticals) are positioned in relation to another element and are laid out together with that element. If there are multiple such *by ref* symbols (or verticals) attached for an element, they are laid out in the order they appear in the source file. For such symbols the following additional attributes are available.

<code>align</code>	This specifies where along the side of the referenced element the symbol shall be placed. It can be set to <code>left</code> , <code>center</code> or <code>right</code> for <code>above/below</code> and <code>top</code> , <code>center</code> or <code>bottom</code> for <code>leftof/rightof</code> . The <code>above_centerline</code> , <code>centerline</code> and <code>below_centerline</code> values are also available for <code>leftof/rightof</code> . Not applicable for <code>around</code> . The default is <code>center</code> .
<code>offset</code>	This shall be a number (defaults to 0). The symbol is shifted by this number of pixels left (negative) or right (positive) for <code>above/below</code> and up (negative) or down (positive) for <code>leftof/rightof</code> . Not applicable for <code>around</code> .
<code>distance</code>	This shall be a number (defaults to 0). The symbol is placed by this number of pixels further away from (positive) or closer to (negative) the referenced element. For <code>around</code> it effectively expands or shrinks the symbol (applying a distance on all 4 sides).
<code>adjust</code>	This attribute can be used to adjust the left/right/top/bottom edge of the symbol after placement. You can specify one to four numbers. If one is specified, it applies to all edges, if two they apply to left/right, respectively, if three, they apply to left/right/top, and if four they apply to left/right/top/bottom. The numbers are interpreted as pixels.

²⁰ We have to note that `text` is just syntactic sugar for a `rectangle` with no line or fill. Rectangles can also contain text.


```

compress=yes;
Backend, _Server, Client;
Client<-> Hit;
Client=>Server;
symbol shape above [shape=def.actor, size=0.2, color=blue];
symbol ... below;
Server>>Backend: Query [refname=query];
Server<<Backend: Response;
symbol rectangle around [line.radius=2, fill.color=yellow,5
distance=5];
Client<=Server: Reply;
symbol cross rightof [align=centerline, color=red];

symbol arc leftof query [align=bottom, size=big, distance=-18,
offset=3, fill.color=teal,50];
note: this [ fill.color=teal,50];
symbol rectangle around query [fill.color=violet,50,
line.radius=2,
adjust=-30,2,2,-2];

symbol rectangle around [fill.color=aqua,50, line.radius=2,
adjust=-30,2,2,2];

```




The second method (using `at`) places the symbol at the horizontal position of an entity. You can also use `--`, `-`, `+` or `++` to place the symbol left or right of the entity line (or maybe a bit further out) and you can also use a number after the entity to shift the symbol to the right (positive) or to the left (negative). The symbol is drawn at the vertical position where they are specified in the file, just like regular elements and the space occupied by the symbol is considered at layout.

```

hscale=auto;
a,_b,_c;
parallel a->b;
parallel symbol shape at b++
[shape=def.trapezoid, xsize=20,
fill.color=olive, fill.gradient=up];
symbol ... at a-b [fill.color=black];
a<-b;
symbol arc at a
[xsize=10, ysize=10, fill.color=blue,
compress=yes];
symbol cross at c [line.color=red];
hide c;


```



With the third method, the symbol will be placed after the layout of the chart. The layout engine will ignore the symbol when placing other elements and the symbol may end up drawn overlapping other elements (this may be your intention). You have complete freedom in selecting both its horizontal and vertical position. Vertical position is specified via markers, similar as for verticals, see Section 9.6 [Verticals], page 119. The vertical size of the object can be specified two ways. Either you specify two markers (as above), in which case the symbol will vertically span from one to the other; or you omit one of the markers,

in which case the `ysize` attribute specifies the height (in points)²¹. If the dash is in front of the marker, the bottom of the symbol will be aligned with the marker. If the dash is after the marker, then the marker designates the top of the symbol. In the example below we see three rectangles. One stretches between two markers, the second is bottom aligned, while the third is top aligned.

```
a, b, c, d;
a->b: Message 1
    goes here;
mark m1;
b->c: Message 2
    goes here;
mark m2;
c->d: Messgae 3
    goes here;
defstyle symbol [fill.color="yellow,200"];
symbol rectangle m1-m2 left at a+ [xsize=30];
symbol rectangle -m2 left at b+ [xsize=30, ysize=30];
symbol rectangle m1-    left at c+ [xsize=30, ysize=30];
```



The horizontal position of the symbol is specified via one or two *horizontal position specifiers*. They specify the horizontal position of either the left or right edge of the symbol or of its center. This is governed by the first keyword

```
left|center|right at entity-entity [number]
left|center|right at entity-
left|center|right at entity-
left|center|right at entity [number]
left|center|right at entity+
left|center|right at entity++
```


Then, after the `at` keyword one specifies either one entity with additional modifiers or two entities. In the former case the horizontal position will be at the middle of the entity's line or somewhat left or right of it depending on the modifiers. In the latter the horizontal position will be between the two entities. Two of the forms can also take a number, which is interpreted as pixels and will shift the position to the right for positive values and to the left for negative values.

If you specify two such horizontal position specifiers one after the other, they describe both the placement of the symbol and its width. If you specify one, the width of the symbol

²¹ In case of rectangles and text, you can use the natural size of the label you specify as the height of the symbol by omitting the `ysize` attribute. If you specify a shape, it is enough to specify one of the sizes, the other will be calculated to keep the original aspect ratio of the shape.

can be specified using the `xsize` attribute²². This may sound a bit complicated, so here is an example with 5 in-line symbols.

```
a,_b,_c,_d;
hspace -a 100; #make room on left side
symbol text left at b+: Two\nlines;
symbol rectangle right at b-: Two longer\nlines;
defstyle symbol [fill.color="yellow,200"];
symbol rectangle right at a- [xsize=30, ysize=30];
symbol rectangle left at a +30 [xsize=30, ysize=30];
symbol rectangle center at c-- [xsize=30, ysize=30];
symbol rectangle left at a-- right at d++ [ysize=30];
symbol rectangle center at b left at a [ysize=30];
```



Whether the symbol is drawn behind or in front of other elements can be controlled by the ‘`draw_time`’ attribute. It can take the following values.

`before_background`

Elements with this property will be drawn before the background. This has effect only if the background is transparent (to some degree or part).

`before_entity_lines`

Elements with this property will be drawn before the entity lines are laid out in the order as they are specified in the chart description.

`after_entity_lines`

Elements with this property will be drawn just after the entity lines are laid out, but before regular elements are drawn.

`default`

This is the default, elements with no `draw_time` will be drawn this time in the order as specified in the chart description.

`after_default`

Elements with this property will be drawn last, after all the above elements in the order as they are specified in the chart description.

²² Similar to height, in case of rectangles and text, you can use the natural size of the label you specify as the width of the symbol by omitting the `xsize` attribute.

Note that from v3.3.4 any element can specify the `draw_time` attribute. It will not impact the layout only the drawing order (what is called the *z-order*).

```

a, b, c, d;
vspace 0;
activate a, b, c;
a->b: Message 1
    goes here;
mark m1;
b->c: Message 2
    goes here;
mark m2;
c->d: Messgae 3
    goes here;
deactivate a, b, c;
---: end;
defstyle symbol [fill.color="yellow,200"];
symbol rectangle m1-m2 center at a- [xsize=30, draw_time=before_entity_lines];
symbol rectangle -m2 center at b [xsize=30, ysize=30, draw_time=after_entity_lines];
symbol rectangle m1- center at c- [xsize=30, ysize=30, draw_time=after_default];

```

As you can see the first (leftmost) rectangle was drawn below the entity lines, the second (middle) one between the entity lines and the arrows, while the last (rightmost) one was drawn on top of the arrows.


Finally we show a few examples of how symbols may be used.

```

mark top;
a, b, c, d;
symbol rectangle top-bottom left at a-b +10 right at c-d -10
[fill.color=lgray, line.type=dashed, draw_time=before_entity_lines];

a->b: Message 1;
b->c: Message 2;
b->c: Message 3;
b->c: Message 4;
symbol ... center at b-c;
b->c: Message \in;
mark circletop [offset=-5];
box c--c: OK: enough;
mark circlebottom [offset=+5];
symbol arc circletop-circlebottom center at c
[fill.color=none, line.width=3, line.color=red, xsize=120];
c->d: Message \in+1;
mark bottom;

```




9.13.3 Inlining other charts

You can use the `symbol` command to include an entire chart using any language that Msc-generator supports. The syntax is the same as for other symbols (Section 9.13.2 [Symbols], page 144), but the `symbol_type` shall be the name of the language `signalling`, `gv` or `block`, see Section 7.1 [Multiple Chart Types], page 56. After the `symbol_type` you can specify the location where to insert and the text of the other chart enclosed between `<<<` and `>>>` symbols.

```

Client1, Client2, Server, Backend1, Backend2;

[layout=one_by_one_merge] {
    defstyle arrow [color=green-50];
    Client1->Server; Server->Backend1;
    Server<-Backend1; Client1<-Server;
} {
    defstyle arrow [color=blue];
    Client2->Server; Server->Backend2;
    Server<-Backend2; Client2<-Server;
} {
    symbol block left at Client1+ right at Client2- <<<
        background.color=lgray+40;
        col { box Client1, Client2; }
        box [line.type=dotted, fill.color=lgray] {
            box Server;
            col { box Backend1, Backend2; }
        }
        arrows use color=green-50;
        Client1->Server->Backend1;
        arrows use color=blue;
        Client2->Server->Backend2;
    >>>;
}
  
```



Any feature of the included language can be used, including defining procedures, designs or including files. No reference, entity, block or node name is shared between the parent chart and the included one. At layout the inlined chart is considered as a simple rectangular block. Note that a corresponding feature is available for Block Diagrams, see Section 11.2 [Block Types and Definition], page 158.

9.13.4 Inline text

Sometimes one just wants to add some text to the diagram and in this case the `symbol text` syntax may be a bit heavy and difficult to do. As an easier way to do that Msc-generator offers the `text at` command.


```
text at pos [attributes]: label;
```

This draws just text (you must specify a label) at the vertical position the command is written. You can use simple horizontal position specifiers, like below to place the text centered in-between two entities; left of an entity, centered around an entity or right of an entity. You can optionally specify a number, which will be interpreted as a pixel offset to the right (negative value to the left.)

```
entity-entity [number]
entity-- [number]
entity- [number]
entity [number]
entity+ [number]
entity++ [number]
```

You can influence the default appearance via the `text` built-in style.

```
hscale=auto;
A, _B, _C;
text at A--: left side\nright ident;
text at A-: left side\nright ident\ncloser;
text at B: centered;
text at B-C: centered again;
text at A+: right side\nleft ident;
```



9.14 Commands

Besides entity definitions, arrows, dividers, boxes, parallel block definitions and options, msc-generator also has a few commands.

nudge	This command inserts a small vertical space useful to misaligning two arrows in parallel blocks, see Section 9.9 [Parallel Blocks], page 127.
hspace	This command forces horizontal distance between two (not necessarily neighbouring) entity. See Section 9.13.1 [Spacing], page 143.
vspace	This command inserts an arbitrary size vertical space, see Section 9.13.1 [Spacing], page 143.
newpage	This command starts a new page, see Section 9.12 [Multiple Pages], page 142.
heading	This command displays all entity headings that are currently turned on. It is useful especially after a newpage command. Note that if there are any immediately preceding or following entity definition commands before or after heading, only one copy of the entity headings is drawn.
show	
hide	Prepending these in front of an entity definition (or later mention) will set the ‘show’ attribute of those entities (there can be a comma separated list) to yes or no, respectively. Using them alone with no following entity names, will show or hide all entities of the chart defined before this point. (Just a shorthand to save typing a lot of entity names.) Specifically show will show all hidden entities. If you want to display a heading for all entities, follow it by a heading command.
activate	
deactivate	Prepending these in front of an entity definition (or later mention) will set the ‘active’ attribute of those entities (there can be a comma separated list) to yes or no, respectively. In addition, when these commands are used to activate or deactivate certain entities immediately after an arrow, the activation or deactivation will take place at the tip of the arrow and not after it. This is to indicate that the activation or deactivation happened as a result of the arrow. This effect is not applied if an entity is activated or deactivated by setting its active attribute. Using them alone with no following entity names, will activate or deactivate all entities of the chart that currently show. (Just a shorthand to save typing a lot of entity names.) The activation status of hidden entities is left unchanged.
mark	This command creates a marker by storing the vertical position of this command. Symbols, verticals and notes can then refer to this location. See Section 9.6 [Verticals], page 119, for more information.
note	
comment	
endnote	These commands are useful to annotate the chart, see Section 9.8 [Notes and Comments], page 125.
symbol	
text	These commands can be used to draw arbitrary graphics to the chart, see Section 9.13 [Free Drawing], page 143.
defcolor	This command is used to define or re-define color names, see Section 8.5 [Specifying Colors], page 83.

defstyle This command is used to define or re-define styles, see Section 8.8 [Defining Styles], page 90.

defdesign

This command is used to define new designs, see Section 8.9 [Chart Designs], page 91.

defshape This command is used to define new shapes, see Section 8.10 [Defining Shapes], page 92.

9.15 Mscgen Backwards Compatibility

Msc-generator was forked (and completely re-written) from mscgen version 0.08 and is a superset of the language of that version even today. During the years the original tool has developed quite a lot and now contains features that are not present or not compatible with that of Msc-generator. From version 4.5 Msc-generator aims to support all the features and syntax of mscgen to provide a true superset of features. Since there are conflicts in the syntax, full support is possible only in a special *mscgen compatibility mode* of Msc-generator.

By default, Msc-generator enters this mode, if the chart text starts with `msc {`. This is because mscgen requires you to start your chart like this, but Msc-generator does not²³. You can force or prevent mscgen compatibility mode using appropriate command-line switches (`--force-mscgen` and `--prevent-mscgen`, respectively) or preference settings in the MFC GUI. Also, on the MFC GUI, the `mscgen-compat` indicator of the status bar turns to red in compatibility mode. (The CLI GUI always operates in autodetect mode.)

Msc-generator attempts to be liberal in what it accepts in both modes. That is, Msc-generator attributes are available also in mscgen mode (So you can, for example, use vertical constructs, hide or show entities, let entities be defined later or use the style mechanism.) Similar, mscgen attribute names and chart options are understood even when not in mscgen compatibility mode. This includes the hashmark syntax for color specifications, like `text.color="#c0ffff"`. For attributes or other syntax I intend to deprecate a warning is given.

A major tool for backwards compatibility with mscgen is the `mscgen` partial design, which contains new styles for arrowhead symbols. E.g., the symbol `->` represents a filled triangle arrowhead in Msc-generator, while it represents a half line arrowhead in mscgen. Using the `mscgen` design via the `msc+=mscgen;` command sets (when not in mscgen compatibility mode) the styles as used by mscgen. In mscgen compatibility mode, such styles are automatically applied.

Some mscgen attributes are interpreted by Msc-generator not exactly the same way as in the mscgen tool due to the different internal workings. The below hold for both the compatibility and regular modes of Msc-generator.

- The `textbgcolor` attribute of mscgen in empty boxes is interpreted as `fill.color`, as in mscgen. It is not possible to have a different fill color and text background for an empty box.
- The `wordwraparcs` chart option is interpreted as `text.wrap`, resulting in word wrapping for all text. Since in mscgen boxes are word wrapped by default, the `mscgen`

²³ If you specify a design, like `msc=qsd {` Msc-generator will NOT enter compatibility mode.

design sets word wrappin in the `emptybox` style. This command sets it everywhere (as intended in mscgen).

- The `arcskip` attribute of mscgen makes an arrow slant down to the vertical position of the next arrow. This is approximated with a fixed `slant_depth`.
- Line breaks in quoted strings are accepted by mscgen (and by Msc-generator in mscgen mode). This is not the case when not in mscgen mode. In general I think it is a bad idea. Use the `\n` text formatting escape or use colon labels.

There are a few conflicting syntax though that must be interpreted differently in the two modes.

- The arrows symbols `::` and `:>` exist only in mscgen mode. These conflict with colon label syntax. You can use `==>` instead of the latter one in both mscgen mode and outside (to represent a double line arrow with a filled triangle arrowhead).
- The `-x` arrow symbol is used to indicate loss in mscgen mode. This arrowsymbol is not understood outside the mscgen mode. Replace it with `->*` instead.
- The `->*` broadcast arrow construct is interpreted in mscgen mode only, since in Msc-generator it means a lost arrow. Even in mscgen mode it is interpreted as a broadcast arrow only if there is no target entity after the asterisk. If there is one (which is not valid mscgen syntax), it is interpreted as a lost arrow according to Msc-generator syntax.
- The `hscale` attribute in mscgen mode is interpreted as setting the width of the chart to 600 pixels times the specified value. In Msc-generator it is interpreted as setting the inter-entity distance to 130 pixels times the specified value (and can of course be set to `auto` to automatically select the spacing between the entities). Due to user feedback the `hscale` attribute is interpreted the Msc-generator way even in compatibility mode²⁴.
- Since `--`, `..` and `==` are arrow symbols (representing headless arrows, essentially just lines), they cannot be used to define empty boxes in mscgen mode using syntax like this: `a==b [label="label"];`, because in mscgen this is a line, not an empty box. So you need to prepend the `box` keyword to indicate that this is supposed to be a box.

²⁴ After all the point of the compatibility mode is to enable migration towards Msc-generator features.

10 Graph Language Reference

Msc-generator supports an extended version of the DOT language, which is fully backwards compatible with the original saving two exceptions.

- New keywords cannot be used as node names. These are `defstyle`, `defproc`, `replay`, `defdesign`, `usedesign`, `include`, `cluster`, `if`, `then` and `else`. If you want these as nodes put them in-between quotation marks.
- Msc-generator does not support graphviz HTML labels. They get parsed, but interpreted only as text.

The language of the graphs (the DOT language) is documented well on the graphviz pages at <http://www.graphviz.org/> (<http://www.graphviz.org/>). Below we just list the additions by Msc-generator.

10.1 Graph Attributes

Msc-generator allows you to specify node, edge and subgraph labels via a colon syntax similar to signalling charts (see Section 8.1 [Labels], page 75). However, since colons are already used in the DOT language to express node ports, two colons shall be used to start labels, like `node::label;`. Such colon labels are terminated by semicolons and opening square brackets or braces - even if in DOT, nodes can be separated via commas or even spaces. When a label is specified via the colon notation all text formatting escapes used by signalling charts can be used (see Section 8.2 [Text Formatting], page 76). Note that in this case it is not possible to use the escape sequences of graphviz (such as `\N` to insert the name of the node).

The format of labels (sepecified either via the `label` attribute or the colon notation above) can be influenced via the `label_format` attribute. Assigning a series of text formatting escapes to this attribute (see Section 8.2 [Text Formatting], page 76) will apply their effect to the label. Note that the formatting is only applied to the label of the node, edge or cluster subgraph, not to any of headlabel, taillabel or xlabel. Labels can contain links via the `\L()` escape, the formatting of which can be set with the `link_format` attribute. On the other hand, the `URL`, `href`, `target`, `tooltip`, `labelURL`, `labeltarget`, `labeltooltip`, `edgeURL`, `edgehref`, `edgetarget`, `edgetooltip`, `headURL`, `headtarget`, `headtooltip`, `tailURL`, `tailhref`, `tailtarget`, `tailtooltip` attributes are ignored.

Msc-generator supports shadows for nodes and cluster subgraphs, via the `shadow_offset`, `shadow_blur` and `shadow_color` attributes. Read more on these in Section 8.6 [Common Attributes], page 85.

Since the `style` attribute of graphviz can take multiple values separated by commas, like `style="dotted,filled"` it is hard to work with it using the Msc-generator style system. For example, if in an Msc-generator style you define `defstyle node [style=filled];`, and later you create a node as `a [style=dotted];`, the attribute would normally completely overwrite the style of the node removing the fill. So in Msc-generator, the '`style`' attribute is applied specially. Any assignment to this attibute will *add* a new value in a comma separated way, so `a [style=dotted, style=filled];` is equivalent to `a [code="dotted,filled"];`. If you want to *remove* a previously added component, use the minus sign, such as `style="-filled"`. If you want to drop previous values to `style` and have the current value be the only one, use an exclamation mark like `style="!bold"`.

Finally, there are a few rules on how attributes can be specified, which are somewhat against the logic of other Msc-generator languages.

- You can use a comma separated list of node names to define nodes, but be aware that any attribute you specify will apply to all nodes in the list. So in `a, b, c::Label`; all three nodes will have the same label. (This is from the original DOT language.)
- Attributes after an edge will apply to the edge and not to the last node. So in `a->b [color=red]`; the edge will be red and not node `b`. (This is from the original DOT language.)
- You can assign attributes to subgraphs in the subgraph heading via square brackets. For example: `subgraph cluster_one [style=rounded] {...};`. (This is an Msc-generator extension, in the original language you can only specify subgraph attributes inside the subgraph.)

10.2 Default and Refinement Styles for Graphviz Graphs

You can define and redefine *styles*. As with signalling charts a style is a package of attributes with a name (see Section 8.8 [Defining Styles], page 90). Any attribute can be assigned to styles (even `label`).

To apply a style to a node, edge or subgraph, use its name alone enclosed in square brackets, like `a->b::Label [style];`. You can have another square bracket after or before with additional attributes, like `a->b::Label [color=red] [style] [arrowhead=onormal];`.

Graphviz has a default style for edges, nodes, cluster subgraphs and graphs named `edge`, `node`, `cluster` and `graph`, respectively. The default style for graph, node and edge can be set using the `graph`, `node` and `edge` commands as well, such as

```
node [shape=record];
```

Msc-generator allows you to collapse cluster subgraphs. Such collapsed subgraphs show up as nodes, but will have `cluster_collapsed` as the default style (boxes with double lines in the plain design). Similar edges that were pointed to nodes inside the collapsed cluster, will have the style `edge_collapsed`.

If you want to adjust an attribute for all versions of an arrow symbol (forward, backward, bidir and arrowless) list all of them in the `defstyle` command, such as in

```
defstyle ->, <->, <<->>, --- [color=red];
```

Note that no quotation marks are needed for these special style names. (Otherwise non-alphanumeric names need to be enclosed in quotation marks.) Specifically, Msc-generator has the following refinement styles for arrows: '`->`', '`=>`', '`>`', '`>>`', '`->>`' and '`==>`', to represent single, double, dotted, dashed line; double arrowhead and double line, respectively. The last arrow type is provided for you to re-define, by default it is the same as '`=>`'. You can also use them backwards (such as '`<-`') or bidirectionally ('`<->`'). There is also a directionless variant for both, without arrowheads: '`--`', '`==`', '`..`', '`++`', '`---`' and '`====`'.

Note that when using the `edge` command to set the default style for edges, not only the `edge` style is modified, but all refinement styles, too. This is to produce the same output as graphviz for the same input files. If you want to modify only the `edge` style, use the `defstyle edge` command.

Msc-generator supports design definitions, for graphs as well. Use the `defdesign` command to define a named design, like below.

```

defdesign omeapple {
    defstyle node [color="#d23f7a", style=filled,
                  fillcolor="#c20a50:#f3cedc", gradientangle=270,
                  fontcolor=white, penwidth=2, label_format="\b"];
    defstyle cluster [color="#c20a50", style=dashed,
                      fontcolor="#c20a50"];
    defstyle edge [penwidth=2, fontcolor="#270210"];
}

```

You can invoke such designs via the `usedesign` command like `usedesign opeapple`. You can also select among the designs on the GUI.

Thus, in summary the actual attributes of an element are set using the following logic. (There are minor variances for each language.)

1. If you specify an attribute directly at the element (perhaps via applying a style), the specified value is used¹.
2. Otherwise, if the attribute is set in the refinement style (at the point and in the scope of where the element is defined), the value there is used.
3. Otherwise, if the attribute is set in the default style of the element in any of the designs applied, the value in the last design that has this attribute is used.

10.3 Clusters

When using the `dot` layout algorithm, graphviz makes it possible to make subgraphs visible by drawing a rectangle around the nodes inside the subgraph. By convention this feature is activated if the name of the subgraph starts with ‘cluster’. In Msc-generator, you can use the `cluster` keyword instead of `subgraph`, to create a visible subgraph. You do not need to append the ‘`cluster_`’ prefix to its name. Also, whatever name you specify will also become the default label. Omitting the name will result in an empty label and an auto-generated name, like ‘`cluster_xxx`’, where ‘`xxx`’ is an auto-incremented number.

Msc-generator also supports cluster subgraph collapsing/expansion. When in the GUI, you hoover over a cluster subgraph, a collapse icon (red minus sign) appears at the top-right corner of the cluster. (Similar to how it happens with boxes and group entities in signalling charts.) Clicking that (or double clicking anywhere in the cluster) collapses the subgraph into a single, double-lined node of shape box. In the GUI there are buttons to collapse or expand all subgraphs with one click.

When collapsing a subgraph multiple edges from within the subgraph to the same node outside will be kept if the graph is not `strict`. Use the `collapse.strict=yes`; chart option to remove duplicate edges on collapsing a subgraph. In this case `edge_collapsed` style is applied to the remaining edge. Similar, nodes that result after the collapse of a subgraph are applied the `cluster_collapsed` style instead of `cluster`.

¹ If you specify the attribute several times, the last one is used.

11 Block Diagram Language Reference

A block diagram consists of *blocks* and *arrows*. Blocks can contain other blocks - these has to be specified immediately after the containing blocks enclosed in braces. Blocks that contain other blocks are called *container blocks* or *containers*. Putting one block into another has two consequences. First, the container is resized to be able to cover all its content. If you want to place a block outside or overlapping with the contour of a container, add it outside. Second, the name of the container will become the name prefix of all its content blocks, see Section 11.1 [Block Name Resolution], page 157. The most common block is a rectangle (or box for short), but other shapes can be used (and defined), as well.


Arrows can be specified anywhere in the chart and their location has significance only in when they are drawn (in Z order) or if they contain coordinates referring to their parent. In the latter case the identity of the parent matters. Note also that blocks not defined at the time of defining an arrow may get auto-created if `pedantic=no`; chart option is specified (the default). So defining an arrow before or after the blocks it uses is different.

11.1 Block Name Resolution

Blocks may have a name, which is useful if you want to later refer to them (to align other blocks or to specify arrows). Several blocks may have the same name, but if you refer to them the reference must be unambiguous. Using a (full) block name twice is not good practice.

You can refer to a block using its *full name*, which is the dot-separated list of its parents and its name. So if block C is inside block B, which is inside block A, then the full name of C is A.B.C. (See the red block in the example below.)

```
box_A: {
    box_B {
        box_C [color=red];
        box_D;
        C->D;
        C->A.D [color=green];
    }
    box_C [middle=B.D];
    box_D;
    C->D;
    D->B.D [color=blue];
}
```



You can also omit the common prefix of the place of reference and the block. For example, the black arrows in the above example are inside the blocks A.B and A and they refer to blocks C/D with the same prefixes, that is to blocks A.B.C/A.B.D and A.C/A.D, respectively.

Similar, the blue arrow, refers to blocks B and B.D, which have the full name A.B and A.B.D, respectively. Using the full name is not needed because we refer to them from within block A. In contrast, the green arrow must use the full name of block A.D, in order to avoid confusion with block A.B.D, which is within the same block as the green arrow.

In short, name resolution is local and I think fairly intuitive. If you do not give a block a name, all its content will have the same name prefix as their parent. In other words, one level of naming is skipped.

Note that the special block names `prev`, `next`, `first` and `last` cannot be given to a block. (See more on them later.) They always refer to a block within the same parent. Also names including dots cannot be assigned to blocks for now, to avoid introducing an extra level in the name hierarchy.

11.2 Block Types and Definition

There are several types of blocks.

- Boxes* are the most common type. They are rectangular in shape, may have rounded or bevelled corners, may contain a label and/or other blocks. They can be defined via the `box <name>` or `boxcol <name>` commands. The former will have its content oriented left-to-right (as a row), while the second top-to-bottom (as a column), see more on block alignment below. The name is case sensitive and optional.

```
box_a: Plain;
box_b: Filled [fill.color=lgray, text.color=white];
box_c: Round\ncorners [line.corner=round, fill.gradient=down,
                      fill.color2=lgray];
below a box_d: Note\nstyle [fill.color=yellow,
                           line.corner=note, line.radius=5];
bottom box_e: Subtle shading [line.type=none, fill.color=lgray,
                           text.color=gray,
                           shadow.offset=4, shadow.blur=4];
```


- Shapes* behave like boxes, but may have arbitrary shape, defined by the `defshape` command. They can be defined via the `shape <shape> <name>` command or using an asterisk as a shorthand `*<shape> <name>`. The name is case sensitive and optional. `text` is a special kind of shape, that has no outline or fill, just a label. You can invoke it via the `text <name>` command, usually adding the label.

```
*trapezoid [color=lgray] {
  use col;
  *oval_comment: This is Joe;
  *def.actor_joe [color=gray, size=30];
  comment->joe;
}
```

- Any box or shape can be made into a *multiblock* series by specifying the `multi` keyword and an optional number before them. This makes the original block appear in several overlapped versions. The content (if any) is kept only once: for the front block in the series. The name of content blocks remain the same as without the multi clause. That is, in case `multi box A { box B; }` block B will have the full name A.B. The name A will refer to the entire multiblock series, while `A.front` and `A.back` will only refer to the front or back block of the series. These may be important when routing arrows to elements of a block series. You can use the optional number after the multi keyword


to specify the number of instances (defaults to 3) and the `multi_offset.x` and `multi_offset.y` attributes to specify in which direction and how much shall the subsequent instances appear.

```
multi 2 *trapezoid [color=lgray,
                    multi_offset.x=10] {
    use col;
    *oval comment: These are
           Joes;
    multi *def.actor_joe [color=gray, size=30
                          comment->joe;
}
```



4. *Invisible containers* are rectangular blocks that contain other blocks, but have no visual appearance: no line, fill or shadow and no label either. They have zero margin, as well. They are only used to group other boxes and to govern their default layout. They can be defined using the `row <name>`, `col <name>` or `cell <name>` commands and their name is also case sensitive and optional. Using `row <name>` or `col <name>` will result in a default row or column arrangement of blocks inside, respectively. In contrast, blocks in a `cell` are continued to be laid out by default as in the parent block. `cell` is useful when you want to temporarily break a row or column, add a small comment, for example, and then continue the row or column, as in the example below.

```
box: Big\nText\nBlock;
cell {
    use margin=0;
    box: Another\nBig\nText
          Block;
    top text: \-small
              comment;
}
box: Whatever;
```




5. Other charts (of any language supported by Msc-generator) can be included as a block. This is done via the `chart <language> <name>` command. The name is case sensitive and optional. After potential attributes, you can specify the chart to include enclosed between `<<<` and `>>>` tokens. The natural size of the block will be the size of the included chart.

```

box_Client [line.radius=3, line.corner=bevel, line.width=2];
space;
box_Processing [fill.color=aqua+80, line.radius=5] {
    chart signalling <<<
        background.color=None;
        hscale=auto;
        _Server, _LB [show=no], _Backend1, _Backend2;
        ->Server;
        show LB;
        Server->LB;
        LB->Backend2;
        hide LB;
        Backend2 ->Backend2:Proc;
        join <->:external\Inreq;
        Server<- Backend2;
        <-Server;
    >>>;
};
space;
box_zone: zone nw [line.type=dotted] {
    *def.db_db [size=40, line.width=2, middle=Client];
    *def.db [size=40, line.width=2];
}
below +30 *def.actor_op [size=20, color=orange];
below -10 text: \-Operator;
Client<->Processing<->db;
op=>zone: config [label.pos=middle];

```




6. There are two special kinds of blocks, that do not appear, but control layout only. Spaces are merely adding extra distance between two blocks in an easy way without any visual appearance. You can define them via the `space <number>` command, where the number is a distance in pixels added and is optional.¹ Breaks do not even add space, merely separate blocks in a sequence, see Section 11.4.1 [Default alignment], page 169. You can add them, via the `break;` command.
7. You can encircle other blocks via the `around` command. List the blocks to include separated by plus signs and follow by a box or shape specification. Msc-generator will prepare a block as big as to just include the listed blocks into its content area. You can specify a label, too. You can also specify content, which will be included into the content area, but will by default be in no alignment relation to the encircled blocks. I am not sure what this feature is good for. By default around blocks have no fill.

```

use label="";
col_A {a, b, c, d;}
col_B {a, b, c, d}
col_C {a, b, c, d}
use line.color=red;
around A.a+A.b *oval;
around B.c+C.d box [line.corner=bevel,
                     content_margin=10];

```




8. You can define blocks that are *joins* of other blocks. Use the `join` keyword followed by the blocks to join separated by plus signs. This is useful to quickly create new shapes from existing ones, such as an L-shape from two boxes.

¹ If the number is omitted space will try to add "just the right amount" of space. Currently this is the not so clever, fixed 10 pixels, but it may change later.

```


boxcol_a: {
    use color=green;
    box; box; box;
}
boxcol_b: {
    use color=blue;
    box; box;
}
box_c: [left=a, top=a@bottom] {
    use color = red;
    *oval; *oval;
    *def.actor; *oval;
}
join a+c;

```



The unison of the surface area of those other blocks is taken to be the surface of the join block. No name, label or content can be specified. Join blocks can have only line, fill and shadow attributes (not even alignment as their position is fully determined by the joined blocks). If they have any of the above three attribute set, the corresponding attribute in the joined blocks is removed. Thus, for example, if the join block has line attributes (the default), the joined blocks do not draw a line.

On the example below, we created three blocks (original), then joined them with leaving only the line attribute of the join set (Line); or unsetting line and setting fill attributes (Fill); or unsetting line and setting shadow attributes (Shadow). Note that the second is useful to apply a smooth gradient fill to the entire joined area and that in this case the joined area must be drawn before the blocks joined (so that the fill is behind them). In the last example the join only draws the shadow.




Blocks (except space and break) can be followed by attributes, see Section 8.6 [Common Attributes], page 85, (including a colon-label, see Section 8.1 [Labels], page 75) and contained blocks enclosed in braces.

The definition of blocks and arrows must be ended with a semicolon (';'). However, if a block contains other blocks, its closing brace ('}') does not have to be followed by a semicolon (to adhere to C syntax better). In fact, no closing brace in Block Diagrams needs a semicolon after (including design definitions, etc.).

Block and shape definition commands can be used to define multiple blocks in one command. Simply list multiple names separated by commas instead of a single name. You must give a name to all blocks defined like that. Any attributes and content you specify will apply to all blocks defined that way.


```
box A, B [color=lgray] {X->Y; }
box C, D [color=aqua] {F->G->B; }
```



If the `pedantic` chart option is set to `no` then you can define blocks without the `box` keyword or without specifying a shape, just by typing the block name. (Unnamed blocks cannot be defined this way.) The block created will take its attributes from its default style and the running style (defaults to a box).

You can specify more than one block name separated by commas. If you specify any attribute (including a potential colon label), it will apply to all the blocks defined. However, if there already exists a block with the name you list no new block will be created, only its attributes will be updated. This is true if the block of a given name exists outside this scope. Finally, block names used in arrow specifications will also get auto-created. Let us see a few examples.

```
a, b [color=red];      //a, b does not yet exist, create them
box a;                  //always defines a block. Now we have 2 blocks named 'a'
b, c [color=aqua];     //creates 'c', but not b. Only makes 'b' blue
box BIG {
    use shape=oval;   //make the default shape the oval
    c->d;            //'c' already exist (outside BIG), but 'BIG.d' is created oval
}
BIG.d [color=green]; //update the color of 'BIG.d'
BIG.e [color=red];     //error, we cannot create a block with a dot in name
```



Note that the above makes it impossible to forward-reference a block from an arrow definition (as it will get auto-created). To allow that you need `pedantic` set to `yes`. In that case, blocks can only be created by the `box` or `shape` commands or via the asterisk abbreviation. You can still change their attributes later, however.

11.3 Block Attributes

Blocks support a number of attributes governing their visual appearance. Any of the below attributes can be made part of styles (even `label`).

label Blocks may have a label, that can be specified using the colon syntax as described in Section 8.1 [Labels], page 75. They can contain formatting escapes, see Section 8.2 [Text Formatting], page 76. Specific to Block Diagrams is the `*` escape, which is replaced to the name of the block (if any).

text.* Text formatting attributes are described in Section 8.6.3 [Text Formatting Attributes], page 87. Word wrapping is also available via `text.wrap`, see Section 9.10.2.2 [Word Wrapping for Block Diagrams], page 136.

number Numbering is used similar to Signalling Charts, see Section 8.4 [Numbering], page 80.

line.*

fill.*

shadow.* These attributes apply fully to blocks. Boxes, in particular also honour **line.corner** and **line.radius**. See Section 8.6 [Common Attributes], page 85, for more.

margin.top

margin.bottom

margin.left

margin.right

These can be used to set the margin around the block. Any nonnegative number will do. Using the **margin** attribute sets all four of the above to the same value. Alternatively, you can specify four comma separated values, to set all four margins, like **margin=10,10,5,5**. The order is top, bottom, left and right.

imargin.top

imargin.bottom

imargin.left

imargin.right

These can be used to set the margin inside the block. This value is used only when the block contains other blocks or a label. Any nonnegative number will do. Using the **imargin** attribute sets all four of the above (plus also **imargin.label**, see below) to the same value. Alternatively, you can specify four comma separated values, to set all four margins, like **imargin=10,10,5,5**. The order is top, bottom, left and right. The default is 10 pixels for boxes (via the default style ‘**container**’) and 0 pixels for shapes (via the default style ‘**container_shape**’). For boxes and shapes with no content (but perhaps a label) all values default to 0 (via the default style ‘**block**’).

imargin.label

Use this attribute to set the distance between the label and the content of the block in pixels. Any nonnegative number will do. Only applicable if the block has both a label and content. Note that setting the **imargin** attribute to a single number will also set this attribute. See above. Also, the default value for this attribute is the same as for the other **imargin.*** attributes above.

content_margin

If the block has content (or is **around** some other blocks), setting this to false (which is the default) will ignore the margins of the contained blocks, when determining the size of this (the containing) block. The margins of the contained blocks will still be applied, when determining their position in relation to one another. For example, setting this to false and all **imargin** values to zero results in the contained blocks actually touch the container. As a shorthand, you can also set this attribute to a number, which is equivalent to setting it **false** and applying the number to **imargin**. Similar to **imargin** you can set all four internal margins using four comma separated values.

`width`
`height`
`size`

These set the size of the block. `size` is a shorthand for setting both `width` and `height` to the same value. Alternatively, you can specify two numbers separated by a comma to set both height and width, such as `box [size=10,20];`. If the block contains other blocks, this is a minimum size - the block can be larger if needed. Otherwise Msc-generator attempts to make the block exactly this big. If the block has a label larger than this size, it may be increased to make room for the label depending on the value of the `label.mode` attribute (defaults to `enlarge`). You can also specify the name of another block. In this case the two blocks will have the same height and/or width. You can even specify a list of other blocks separated by plus signs (+). In this case the size of the block having this attribute will be equal to the space occupied by the blocks listed (and not the sum of their size). Thus if blocks A and B are far from each other, setting the width of a block C to `width=A+B` will include the distance between A and B.

```
box_A;  

space 50;  

box_B;  

box_C [width=A+B];
```



`aspect_ratio`

Forces a certain X/Y ratio for the block. Cannot be used if both the width and height are set. The value is a floating point number, such as `aspect_ratio=1.5` or a ratio (width divided by height), such as `aspect_ratio="3:2".`)

`label.pos`

Used to set the position of the label inside the block. Five values are valid: `above` (above the content and top aligned), `below` (below the content and bottom aligned), `left` and `right` (left or right of the content and left/right aligned). For blocks without content it can also be `center` in which case the label is centered both horizontally and vertically.

`label.align`

If `label.pos` is `above` or `below`, this can be `left`, `right` or `center`. If `label.pos` is `left` or `right`, this can be `top`, `bottom` or `middle`. It is used to specify the label's alignment in the direction not set by `label.pos`. You can also use a percentage number for finer control. Value 0 means `left` or `top`; value 100 means `right` or `bottom`; while value 50 means center or middle. You can use any value between 0 and 100, but not outside.

`label.orient`

Governs the orientation of the label. It can take four values, `normal`, `upside_down`, `left` or `right`. The default is the first one, which is normal, left-to-right reading.

label.mode

It dictates what to do when both a label and the size of a block is set. The default value `enlarge` will result in scaling the block to accommodate the label. Using `scale` will scale down the text to fit the label. Especially useful for shapes. Finally, `scale_2d` will also scale the text, but differently in the two dimensions.

allow_arrows

If set to yes, arrows do not go around this block, but can cross it. Note that for individual arrows, you can allow crossing of a block via the `cross` and `cross_all` arrow attributes, see Section 11.5 [Arrows in Block Diagrams], page 175.

multi_offset**multi_offset.x****multi_offset.y**

These govern how many pixels subsequent instances in a multiblock series are shifted compared to the previous one. If both are set to a positive number the next element will come left and up by the number of pixels specified. Specifying a negative number will reverse the given direction. Setting `multi_offset` is equivalent to setting the `x` and `y` direction to the same value.

draw_before**draw_after**

You can specify one or more blocks (separated by plus signs or via repeating the attribute) the block shall be drawn before or after. You can use only one of the two attributes for any diagram elements. Setting one will automatically clear the other.

collapsed**indicator**

Setting this attribute on a visible block (`box`, `boxcol` or `shape`) that contains other blocks will skip drawing the visible block. Instead a small indicator is drawn to show that there are hidden elements. You can disable the indicator using the `use indicator=no` command or via the `indicator` block attribute.

Note that if you disable the indicator for a collapsed block, its base style will become `block` and not `container` or `container_shape` (see below), as is the case for blocks with content.


Blocks also rely on a few built-in styles. The attributes of each block default to the style `block` if they are empty and to the style `container` or `container_shape` (for boxes and shapes, respectively) if they contain other blocks.² Then a style for each shape is applied (if found) that has the name `shape_<shapename>`, that is `shape_box` for boxes or `shape_oval` for the shape `oval`. These apply to the specific shape irrespective of whether that is a container or a block. They can be used to set shape specific margin, for example. The shapes `around`, `join`, `indicator` and `text` serve as basic styles for the respective block

² Thus now we have 3 basic default style for blocks. `block` is for boxes and shapes that have no content, just maybe a label. `container` is for boxes having content, while `container_shape` is for shapes having content. The reason for all these is to allow different internal margins by default. Of these three only `container` has a nonzero default `imargin`, specifically 10. The other two has `imargin=0`.

types, while `invis` is used for invisible blocks such as `row`, `col`, `cell`, `break` or `space`. Finally, the special styles `col` and `row` are used to set default alignment for blocks in a column or row, respectively. See Section 11.4.1 [Default alignment], page 169, below.

There are two more special styles, called the *running styles*. They are applied to every block and arrow (respectively) created and can be modified via the `use`, `arrows use` and `blocks use` commands, with the first one having an effect on the running styles for both arrows and blocks. These commands enable quickly setting the attributes of all following blocks and/or arrows. Note that similar to all other style modifications, any change made to the running styles is in effect only till the end of the scope, see Section 8.7 [Scoping], page 89. Thus on the example below, the red line color only applies to block B and the arrow between A and B, but the margin is also applied to block C.

```
box: 0;
box {
    box_A;
    use line.color=red;
    box_B;
    A->B;
};
box_C;
C->A;
```



11.4 Block Layout


The layout of blocks is central to the block diagram language. It is governed primarily by the *alignment* attributes of each block. We have quite a number of them, but they are quite easy to remember. Four are for horizontal alignment: `left`, `center`, `right` and `xpos` and four are for vertical alignment: `top`, `middle`, `bottom` and `ypos`. Since the words `center` and `middle` are easy to mix, `xmiddle` and `xcenter` can be used for `center` and `ymiddle` and `ycenter` for `middle`.

These attributes govern, where the sides or the centerline of the block is placed in the horizontal or vertical direction. As value you can name any side or centerline of any other block (in the same direction only). If you only specify a block name the same side/centerline is used, that is `left=A` means `left=A@left`.

When you set one of the four side attributes (`left`, `right`, `top` or `bottom`) it refers to the visible side of the block. If you want to align to the side of the block *including margins* prepend these four values with the letter `m` both as attribute names and attribute values, e.g., `mleft=a@right` or `top=b@mbottom`. Note that setting either of `left` or `mleft` both set the position of the left side, so setting one will overwrite the other - they are effectively the same attribute with the margin as difference. On the below example you can see that block B is aligned to block A without maring in the X direction and with margins in the Y direction.


You can not select a side for `xpos` or `ypos`, just a block. It is equivalent to setting both `left/right` or `top/bottom` to the same block effectively resulting in the same width or height as the referenced block.

```
box_A: A\nA;
box_B [mbottom=A@mtop, left=A@right];
box_C [center=A, mbottom=B@mtop];
box_D: DDDD [middle=B, mleft=B@mright];
box_E [xpos=D, ypos=A];
```



You can specify multiple blocks separated by the plus sign (+) as value to any of the alignment attributes. In this case Msc-generator takes the bounding box of the listed blocks and aligns to that. If you specify the `m` letter after the @ symbol, the bounding box is calculated from the listed blocks with their margin.


```
box [imargin=0, margin=0,
      line.type=none, fill.color=lgray] {
    box_A;
    below rightof box_B;
}
box_C [middle=A+B, mleft=A+B@mright];
box_D [middle=A+B, mright=A+B@mleft];
box_E [center=A+B, mtop=A+B@mbottom];
box_F [center=A+B, mbottom=A+B@mtop];
```



In addition to the tokens representing box sides or centers, you can align to any part of a block by appending a percentage number after the @ symbol. The value 0% is equivalent to `top` (or `left`), while 100% represents `bottom` (or `right`). Naturally 50% means `middle` (or `center`). However, you can also specify any other number, even smaller than zero or larger than 100%. (You can omit the percent sign.) You can use the `m` letter even with percentage numbers, such as `left=A@m33%`, which means the left thirdpoint of the span between the left side of A minus its left margin and right side of A plus its right margin.

```
use col;
box_A;
box: Aligned to left of A [left=A@0%];
box: A Bit right [left=A@10%];
box: Right by a quarter of A [left=A@25%];
box: Right by a half of A [left=A@50%];
box: Aligned to right of A [left=A@100%];
box: Beyond the right of A [left=A@150%];

use draw_before=A;
(A@left,A) .. (,A+300);
(A@right,A) .. (,A+300);
```



Finally, you can also add a pixel offset at the end of alignment attribute values. Positive values mean right or down. Thus the full syntax for alignment attributes is as below.

`[block[+block]...] [+|-offset] [@[m] (token|percentage) [+|-offset]]`

Where

- `token` can be one side or `middle/center`;
- `percentage` is a number with an optional percentage sign;
- if no blocks are specified, `prev` is assumed;
- if no `@` symbol (and parts after) is specified the token value defaults to the attribute we set; and
- only one of the two `offset` values can be present.

A second way to specify block alignment is to align it to its parent. In this case the `@` symbol shall be omitted and the syntax becomes as below.


```
[m]percentage%[+-offset]]  
[m] [offset]
```

If you specify a percentage, it refers to the position between the inner margins of the parent. Thus, `top=10%` will place the top of the block close to the upper internal margin of the parent, while `0%` and `100%` refers to alignment to the top or bottom internal margin of the parent, respectively. Specifying an offset is added in pixels to the position designated by the percentage. Specifying only an offset (no percentage sign) will be a pixel position added to the top or left internal margin. Prepending a `m` character will use space covered by the content part of the parent block ignoring any inner margin.³⁴

Of course, all contained blocks must end up between the inner margins. Thus when you specify the above, do not use percentages outside the `100%` region. Also, bear this in mind when ignoring the inner margin, e.g., do not specify `top=m0%` unless the inner margin is zero, since this by definition dictates a block to be outside. When the attribute specified violates this rule, it will be ignored with an error, see Section 11.4.4 [Block Layout Conflicts], page 174.

After applying all alignment attributes, exactly one or two of each direction triplet has to be specified. (`xpos` and `ypos` immediately specifies two of them.) If only one is specified, the block's size will determine the other side. That is, in case of specifying `top=A`, but neither of `bottom` or `middle` the bottom side will be determined by adding the height of the block in question. If two of the triplet are specified, they also act as a size specifier, so `left=A`, `right=A` makes the block as wide as `A`.

```
box_A: narrow A;  
box_B: narrow B [mleft=A@mright, middle=A];  
box_C: C: combined width  
[left=A, right=B, mtop=A@mbottom];
```



You can use `prev`, `next`, `first` and `last` instead of a block name. These refer to the block specified before the current one, the one specified after, the first or last block within

³ For shapes this is the area specified by the `T` line in a `defshape` command, see Section 8.10 [Defining Shapes], page 92. For boxes this is the inner edge of the box line. If the corners are bevelled or rounded, this is the largest rectangle that fits into the box.


⁴ Note that when referring to the parent the presence of the '`m`' character will *ignore* the inner margin, whereas when referring to another block, the presence it will *include* the outer margin. This is so that the most common case would not need the '`m`' character.

this parent block.⁵ This allows easy default values. For example, `mleft=prev@mright`, `ycenter=prev` will align just right of the previous block - in a row including outer margins. Note that you must avoid using `prev` and `first` on the first element inside a block and `last` and `next` on the last one. These generate an error.

11.4.1 Default alignment

In order to avoid specifying each attribute by hand, a mechanism of automatically setting these attributes is in place. If you set none of a triplet for a block, and it is not a single child, it inherits the alignment attributes from the *running style* for that direction. Two styles `row` and `col` are defined containing the default alignment attributes for row and column-oriented arrangement, respectively. The `box`, `shape` and `row` blocks automatically add the style `row` to the running style, while the `boxcol` and `col` blocks add the style `col`. You can also manually add these styles by adding `use col;` or `use row;` anywhere. Subsequent boxes will be arranged such.

```
box: A;
box: B;    #follows in a row
use col;
box: C;    #follows in a col
box: D;    #follows in a col
use mleft=prev@mright;
box: E;    #follows diagonally
box: F;    #follows diagonally
```



The `col` style can be defined as

```
defstyle col [mtop=prev@mbottom, center=prev];
```

making each block "stick" to the previous one from below and to center horizontally to it. Applying it to a block or to the running style will make these the default value for that and any following block.

Note that, since applying this to the first block would case an error (there is no `prev` of a first block and it also cannot align to itself), alignment attributes referring to invalid `prev`, `next`, `first` or `last` block names are silently ignored if the attribute is coming from a style. If, on the other hand, you specify `top=prev@bottom` on a first block explicitly, you will get an error.

Unfortunately, the above definition is not sufficient, since if we have a series of blocks and assign a custom alignment for one in the middle, then it will stop "sticking" to the previous one. Like on the example below (where we specified all alignment attributes manually, to prevent any of the defaults to apply. This makes blocks `A` and `B` (as a group) unrelated to any other block vertically. (Clearly block `B` "sticks" to block `A`, but none of them have any other vertical alignments.) Since Msc-generator strives to minimize the size of each block, it will overlap blocks `A` and `B` with blocks `C` and `D` to minimize the height of the unnamed column they are in.

⁵ Specifically, this is not the first and last element within the *scope*, since you can open new scopes any time. This is the first or last block, who has the same parent block.

```

box_ref: ref [height=100];
col {
    box_A: AAAA;
    box_B: BBB [mtop=prev@mbottom, left=prev];
    box_C: CC [mtop=ref@mbottom, left=prev];
    box_D: D [mtop=prev@mbottom, left=prev];
}

```

To prevent this and always provide a suitable default the style `col` is actually defined such that each block "sticks" to both its previous and next blocks.

```
defstyle col [mtop=prev@mbottom, mbottom=next@mtop, center=prev];
```

This can lead to unwanted effects, if you specify alignment for two blocks in a row, with specifying no special alignment for the block in between them. On the figure below we tied block `M` and `O` to the left and right side of the row above and kept the default (sticky) behaviour for block `N` in-between. We gave on where to position block `N`, just that it shall align to the blocks left and right to it. Msc-generator has chosen to break the stickiness between blocks `N` and `O` quite at random.⁶

```

col {
    row {
        box_A: Alpha;
        box_B: Beta;
        box_C: Gamma;
    }
    row {
        box_M [left=A];
        box_N;
        box_O [right=C];
    }
}

```

To fix this situation, clear one of the alignment attributes of `N`, like below.

```

col {
    row {
        box_A: Alpha;
        box_B: Beta;
        box_C: Gamma;
    }
    row {
        box_M [left=A];
        box_N [left=];
        box_O [right=C];
    }
}

```

⁶ Specifically, alignment requirements of blocks deeper in a nester block hierarchy or earlier in the source file take precedence, if the two alignment has the same priority. See Section 11.4.4 [Block Layout Conflicts], page 174.

This prevented inheriting the default alignment attribute `left` but still allowed all other alignment attributes (specifically `right`) to apply to block N.


To summarize the rules for inheriting alignment attributes:

- If a block has none of its alignment attributes set in a direction (where a direction can be *horizontal*⁷ or *vertical*), it inherits the alignment attributes of the running style in that direction. The two directions are independent.
- If a block has none of its alignment attributes set in a direction, but some are *cleared* (set to empty), then it inherits the alignment attributes of the running style in that direction, with the exception of the cleared ones.
- If a block has any of its alignment attributes set in a direction, it inherits none of the alignment attributes in that direction.

Finally, sometimes one wants to break this mutual alignment reference between two neighbouring boxes, perhaps, because alignment will be fully specified for the subsequent one and we don't want the former one to keep "sticking" to it.

```
box_A;
box_B [bottom=A@40%];
box_C [bottom=A@20%];
box_D [bottom=A@0%];
rightof A box_E [top=A@60%];
box_F [top=A@80%];
box_G [top=A@100%];


arrows use draw_before=A;
(A, A@top)..(D,);
(A, A@bottom)..(D,);
```



While this can be done by clearing one attribute of D (`right=`), it is not so intuitive. So as syntactic sugar we can use the `break;` command. Any alignment attribute referencing this zero-sized block is silently dropped (even values, you explicitly assign).

```
box_A;
box_B [mbottom=A@m40%];
box_C [mbottom=A@m20%];
box_D [mbottom=A@m0%];
rightof A box_E [mtop=A@m60%];
box_F [mtop=A@m80%];
box_G [mtop=A@m100%];

arrows use draw_before=A;
(A, A@mtop)..(D,);
(A, A@mbottom)..(D,);
```




11.4.2 Alignment modifiers

Sometimes one wants to depart from the row or column-like arrangement enabled by the above mechanisms. For this you can use the *alignment modifiers* that can be written before

⁷ ...with the horizontal alignment attributes being: `left`, `center`, `right`, `mleft`, `xcenter`, `xcenter`, `mrigh` and `xpos`

any block definition. They can be of two kind: *major* and *minor*. Major alignment modifiers are `leftof`, `rightof`, `above` or `below`. Writing them before any block places that block in the given relation to the previous block (and center/middle aligned in the other direction). You can also use minor alignment modifiers `top`, `bottom`, `left` or `right`, which only set alignment in one direction making the new block align at the specified side. You can use two major alignment modifier of them like `below leftof box A;`, which will place A diagonally south-west from the previous block. Note that the next block will continue to use the default row or column like layout, that is be left or below and middle or center aligned to the *previous* (so far last) block, such as block H is aligned to G on the figure below.


```
use margin=10;
box_A,_B;
above box_C;
leftof box_D;
above box_E;
rightof box_F;
rightof below box_G;
box_H;
A->B->C->D->E->F->G->H;
```



You can also specify one major and one minor alignment specifier, like `below left` to override the center alignment in the horizontal direction and make it left aligned.

You can specify one or more blocks after the alignment modifiers (use the plus sign (+) to separate them) if you want to relate to a block other than the previous one. When using several blocks, Msc-generator will align to their bounding box, similar to how attribute values are interpreted.

```
box_A;
box_B;
below A+B box_C;
A,B->C;
rightof A+B+C box_D;
B,C->D;
above B+C+D box_E;
A,B,C,D->E;
```



You can also specify an offset in addition to (or instead of) the block(s) after alignment modifiers. This number will be added appropriately.

```

box_A: Very large block
  [label.align=left];
rightof A+10 top box: +10 and top;
rightof A+20 box: +20 and center;
leftof A-30 bottom box: -30 and bottc
                                -30 and bottom

```

The diagram illustrates the placement of blocks relative to a large block A. A vertical line labeled 'Very large block' represents block A. To its right, three smaller blocks are placed: one at '+10 and top', one at '+20 and center', and one at '-30 and bottom'. Arrows point from the text labels to their corresponding block positions.

Note that the blocks/number and the second alignment modifier may come in any order.


You can also specify a coordinate after alignment modifiers. This is less useful, so I do not give an example here.

If you specify alignment modifiers before a definition of multiple blocks, it will only apply to the first one. For example, in the example below only block C is below block A, whereas block D simply follows block C following the default row alignment. The same way, you can apply alignment modifiers after the definition of the block, but if you do it for a series of blocks, it only applies to the first one (C.F in the example below).

```

box_A, B [color=lgray] {_X->Y; }
below A
box_C, D [color=aqua] {_F->G->H, B; }
rightof C.H C.F, C.G [line.width=3];

```



11.4.3 Content placement inside another block


There are two more attributes governing alignment. By default a block is just as big as needed to contain its contained blocks. However, we can specify an explicit size larger than this; it may have a very large label that makes it bigger than its content; or its sides may be aligned so that it ends up being larger. In this case the `content.x` and `content.y` attributes govern, how the contained blocks are placed inside. They can take the following values: `left/top`, `center/middle`, `right/bottom` or `justify`.

If you specify `justify` the blocks that are laid out in relation to their previous or next block will not necessarily "stick" to that previous or next block. Additional space may be added to have a justified layout inside the parent. It may sound complicated, but in most practical cases it just works.

```

boxcol
  {box; box; box;}
boxcol [content.y=bottom]
  {box; box; box;}
boxcol [content.y=justify]
  {box; box; box;}
use line.color=blue;
space;
boxcol [height=140]
  {box; box; box;}
boxcol [content.y=bottom, height=140]
  {box; box; box;}
boxcol [content.y=justify, height=140]
  {box; box; box;}

```




Inside a shape the content is placed into the area specified by the T command (used for labels in signalling charts). The whole shape will be sized to be able to contain its contained blocks. On the figure below, we define a simple triangular shape with fill (the part after S 0;) on the area, where content can go (the box designated at the last line)

```

defshape tri {
  S 2;
  M 100 0; L 200 100; L 0 100; E;
  S 0;
  M 50 50; L 150 50; L 150 100; L 50 100; E;
  T 50 50 150 100;
}

*tri [fill.color = lgray] {
  box_A: A;
  below *oval_B: B;
  *tri_C: C [middle=A];
}

```



Note that apart from sizing containers appropriately Msc-generator makes no attempt to avoid overlaps. Blocks go where you specify them and if they overlap they are simply drawn like that.

11.4.4 Block Layout Conflicts

As you may have guessed it is easy to create conflicts between alignment attributes. E.g., top=A@bottom, bottom=A@top is an obvious conflict, that cannot be satisfied. More complicated cases may also arise. For example, below all three boxes are left of another one.

```

box A [left=C@right, size=10];
box B [left=A@right, size=10];
box C [left=B@right, size=10];

```

Or in this example the two attributes of c cannot be met at the same time.

```

box A [top=B, size=10];
box B [bottom=A, size=10];
box C [top=A@bottom, middle=B];

```

Note that when specifying or processing alignment attributes the order of their specification is indifferent. Msc-generator does not work by placing one block and align any blocks referring to it afterwards. Instead, it solves the constraint requirements as a linear problem using the simplex algorithm. Thus specifying `top=B@bottom` in block A is equivalent to specifying `bottom=A@top` in block B. Hence, in case of conflicts it is hard to pinpoint which attribute actually caused the conflict with which other attribute.

Msc-generator has a set of priority levels for resolving conflicts. Attributes explicitly specified for the block or part of a style explicitly assigned to the block take highest priority. Any attribute coming from the running style comes second (e.g., specified via the `use col`; command), then come the default content arrangement of blocks, such as the fact that blocks in `boxcol` are arranged in columns. Lastly any attribute coming from the default style takes fourth priority and default values are the lowest. If an attribute explicitly specified cannot be fulfilled (because other such conflicting attributes), Msc-generator creates an error naming the attribute it ignores to resolve the conflict. If other, lower priority attributes need to be removed (because they conflict something of the same or higher priority than themselves) Msc-generator generates a warning. You can suppress these warnings via the `conflict_report` chart option. It can be set to any of `off` (no conflict warnings at all), `style` (warnings if an attribute from the running style is ignored), `default` (warnings even if attributes from content or default styles are ignored) or `full` (warning for all ignored attributes). The default is `off`.

When two alignment attributes with the same priority conflict, the one belonging to a block deeper in a nested block hierarchy takes precedence. In case the two block are contained in the same number of blocks (such as two block inside the same container) the one specified earlier in the source file takes precedence.⁸

In addition to emitting an error or warning when an attribute conflicts with attributes of higher priority, it also attempts to figure out which higher priority attribute may be the one the conflict is with and indicates it to you.⁹

11.5 Arrows in Block Diagrams

11.5.1 Defining Arrows and Lines

Arrows (and lines) can be defined between any two blocks using *arrow symbols*. When you specify a block as start or end of an arrow, its center is targeted. Alternatively, you can specify `port` for the block (if some other place than its center is needed) or a `coordinate` instead of a block as start and/or endpoint.

If the arrow (or line) starts (or ends) at a block (or at a coordinate having the same, single block in the specification of both X and Y coordinates) and the arrow starts from the inside of the block, it is clipped to go only to the perimeter of the block. This can be turned off by setting the `routing.clip_block` attribute to `no`.

⁸ When an object is copied, its content may be reordered (see more on copying later) and content can be added/removed, too. In this case the precedence between attributes of the content blocks is according to the newly created order, even if an element later defined is inserted before an element earlier defined.

⁹ Specifically, for every layout conflict error or warning emitted, it cycles through all the higher priority attributes and checks if the layout can be solved by removing them. This algorithm does not find any solution if a conflict is caused by a combination of multiple higher priority attributes.

Currently the usual four arrow symbols are supported: `->` for solid line, `>` for dotted line, `>>` for dashed line and `=>` for double line. They can point backwards so `A->B;` is equivalent to `B<-A;.` Bidirectional arrows (`<->`, `<>`, `<>>` or `<=>`) or simple lines (`--`, `..`, `++` or `==`) are also valid.


You can specify multiple blocks and coordinates on both sides of the arrow symbol separated by commas. In this case an arrow is added from each block on the left side to each block on the right side (using the same arrow symbol). This is an easy way to quickly add a lot of arrows.

You can chain arrow definitions, so `A->B->C;` is equivalent to `A->B;` `B->C;.` You can use different arrowheads in a chain and can have more than one block, so `A->B,C<=D;` is equivalent to `A->B;` `A->C;` `B<=D;` `C<=D;.`

You can specify attributes for arrows in square brackets after the arrow specification. Attributes apply to all arrows if you have several blocks, but only for the latest arrow symbol in a chain. For example, `A->B,C [attr1] <=D [attr2];` makes the first set of attributes to apply to the first two arrows (staring from A) and the second set of attributes to the second two arrows (starting from D).

```
use margin=15;
box _A,_B;
below box _C;
rightof box _D;

A->B, C [color=red] => D [color=blue];
```



Note that if the `pedantic` chart option is set to `no` (the default) then any block used to define an arrow will automatically get created if it has not been defined before.

11.5.2 Ports and Directions

When an arrow or line starts or ends in a block, you can add a *port* after the block name and the '`@`' symbol ('`@`').¹⁰ This will make the arrow or line start or end at the part of the block denoted by the port. Shapes can define where ports lie, see Section 8.10 [Defining Shapes], page 92. Boxes have eight ports defined (`topleft`, `top`, `topright`, `bottomleft`, `bottom`, `bottomright`), but you can also use `n`, `s`, `e` and `w` (for north, south, east and west, respectively) followed by a percentage to specify a port along the perimeter of the block. The direction specifies the side of the block, whereas the percentage specifies the position along that side.


In addition to the above ports, you can apply a *compass point* to any block as a port. These can be `n`, `ne`, `e`, `se`, `s`, `sw` or `w` and will make the arrow start/end at the countour of the block seen in the specified compass point from its center. You can also specify a number in degrees, which is interpreted as clockwise from the north (top) direction.

¹⁰ Ports and directions work similar to graphviz, with the exception of using the '`@`' symbol ('`@`') instead of the colon ('`:`') to separate the block, the port and the direction. In graphviz the direction is called a *compass point*, because one can only use a compass point to specify a direction.

```

use col;
arrows use text.size.normal=10;
a, b, c [width=100, margin=20];
b [width=20];
use label.pos=right;
b@ne->a: ne;
b@e->a@ne: e;
b@170->c: 170 degrees;
use label.pos=left;
b@-20->a: -20\ndegrees [label.align=30];

```




In addition to ports you can also specify a *direction*, which govern what direction the arrow or line leaves the block. This can be any of the above compass points (or degree numbers), or *perp*, which is a direction perpendicular to the contour of the block.

```

*oval_a: oval a;
*oval_b: oval b;
below a left
  box_c: box c;
*oval_d: oval d [xmiddle=b];

a@n->b@n;
b@sw@perp->a;
b@nw->a;
c@topleft->a@w;
c@topright@90->d;
c@bottomright@180->d;

```



Ports usually have a default direction. This equals the direction of the compass point or degree number if that is used as a port. If you omit the direction after specifying a port then the default direction of the port is used. If the port has no default direction or there is no port specified (like, for example, with the simple `a->b;`), the direction will point towards the other end of the arrow or line.

When you specify a direction (via a port default or explicitly), Msc-generator starts the arrow in the given direction, before gradually turning to where the arrow must go. This is done by setting a waypoint 10 pixels from the contour of the block in the direction you specify. If you want a larger or smaller distance (leading to a wider or tighter arc) you can use the `distance` attribute for the starting or ending blocks. That is the line below will result in a tight arc at the block a, but a wider arc at block b.

```
a@ne -> b@nw [distance=a@5, distance=b@20];
```

Setting the `distance` for all blocks also impacts the waypoint distance, but only if set to a value larger than 10. If you want to reduce the waypoint below 10, use the `distance` attribute specifically for the starting or ending block as in the above example.

11.5.3 Fine-tuning Arrow Ends

Msc-generator allows you to fine-tune the exact location of the arrow endpoints. There are two kinds of modification after you have selected an arrow end via blocks and ports.

- Shift the endpoint by a fixed amount of pixels along the X and/or Y axis. For this add `+x<pixels>` or `-x<pixels>` after the block and/or port/dir specification. You can add either or both an X and an Y coordinate offset.
- Shift in relation to the block of the endpoint. In this case add a percentage, like `+x<percent>%` or `-x<percent>%`. Note that this is a relative shift added to the endpoint calculated from block and/or port/dir specification. Thus for example `a+x10%->b` will start the arrow from coordinate (`a@60%`, `a@50%`) (a bit to the right from the center of `a`), since in the absence of a port the arrow starts from the middle of the given block.
- Make the arrow longer or shorter. This can be achieved by simply adding or subtracting a pixel amount to the arrow end specification by appending `+<pixel>` (no x or y) or `-<pixel>`. Negative values make the arrow shorter, positive values longer. Lengthening of a curvy arrow is made by continuing it in a straight line. You can also add a percentage, which is calculated from the original length of the arrow.

You can mix and match any number of the above modifiers after the block/port/dir specification - they will be cumulated and their order is irrelevant. The shift operations will be applied first, then the length modifier ones. They can be combined with both `via` attributes and any routing algorithm.

```
arrows use routing.arrow_distance=0; //turn automatic overlap avoidance off
a -> b; //From the center of 'a' to the center of 'b'
a+x5 -> b; //From five pixels right of the center of 'a'
a+x5+y5 -> b; //From five pixels right and down from the center of 'a'
a -> b-5; //From the center of 'a' to the center of 'b' - but stop 5 pixels before 'b'
```

11.5.4 Automatic De-Overlapping at Arrow Ends

If you define multiple arrows between two blocks normally they would overlap. If you specify multiple arrows from a block using the same port and direction, they will start overlapping. Normally neither of these overlaps are pleasant, so Msc-generator aims to *de-overlap* arrow ends by shifting them apart a bit.¹¹

For each arrow you can specify the `routing.arrow_distance` attribute to specify how many pixels to set them apart. Msc-generator takes the width of the arrows into account, but not the width of their arrowhead. If you specify zero, then the given arrow will not be modified.

¹¹ Arrow ends specified via coordinates, where both the X and Y coordinate uses the same, single block and the arrow endpoint is inside that block, will also get de-overlapped. This is a good way to specify a port for a block, whose shape does not have the port position, you need.


```

cell_a {
    box_A;
    col {B, C};
    A@e->B,C [routing.arrow_distance = 0];
}

cell_b {
    box_A;
    col {B, C};
    A@e->B,C;
}

cell_c {
    box_A;
    col {B, C};
    A@e->B,C [routing.arrow_distance = 15];
}

```




After shifting the endpoints, arrows may still overlap, for example if they go around the same block they will follow the contour of that block. By setting `routing.block_others` to `yes` (the default) on the arrow drawn earlier will prevent any later arrows that had overlapping ends with it from crossing it or getting closer to it than `routing.arrow_distance`. You can specify a number via `routing.order` to govern in which the arrows in overlapping by their end are re-laid out after their endpoints have been shifted. Note that arrows never block other arrows that have no overlapping ending with them.

```

cell_a {
    use bottom=first;
    box_A;
    box_B[size=40];
    box_C,D;
    A->C;
    A->D [color=red];
}

below cell_b {
    use bottom=first;
    box_A;
    box_B[size=40];
    box_C,_D;
    use routing.block_others=no;
    A->C;
    A->D [color=red];
}

```




Below we show a more complicated example, where we want all three arrows to go between blocks `y` and `v`. This is complicated, because if we draw the either the black or the arrow first, it will go very close both to `y` and `v` and will leave no room for the other arrow (red or black, respectively). So we engineer to draw the red before, but proscribe a larger distance between it and `y` to leave room for the black arrow. The black arrow is then assigned a `routing.order` of 1. (Which is larger than the default of 0 and will cause the black arrow to be drawn after the red.) Adding the same distance to the blue arrow will ensure its overlap with the red one (and hence the proper de-overlapping and no crossing).

```

cell_x {
    box_x,_y,_v,_z;
    y,z [ymiddle=x@top];
    top x v;
    #x->z [routing.order=1];
    x->z [color=red, distance=y@10];
    #x->v [color=blue, distance=y@10];
}
below cell_y {
    box_x,_y,_v,_z;
    y,z [ymiddle=x@top];
    top x v;
    #x->z [routing.order=1];
    x->z [color=red, distance=y@10];
    x->v [color=blue, distance=y@10];
}
below cell_z {
    box_x,_y,_v,_z;
    y,z [ymiddle=x@top];
    top x v;
    x->z [routing.order=1];
    x->z [color=red, distance=y@10];
    x->v [color=blue, distance=y@10];
}

```




You can specify `routing.try_harder` to `yes` to make Msc-generator try all permutations in the order of how the arrow ends are shifted. Doing a complete search is quite slow, so this is not recommended for a large set of overlapped arrow ends. If not set, we stop the search as soon as a layout is found, where no arrow overlaps with another in the group of overlapping arrow ends.

Note that if two arrows go on the exact same path they will be routed together. To split them to different paths, set the `routing.allow_joint_layout` attribute to `no` for one of them. (This attribute defaults to `yes`.) In the below example we have the same situation for two pairs of arrows (one pair from block A to C and another pair from block C to E). The first pair is routed together, but the second pair is not.

```

box_A,_B,_C,_D,_E;
A->C;
A->C;
C->E [routing.allow_joint_layout=no];
C->E;

```



11.5.5 Defining Coordinates

Coordinates can be specified using values similar to the value of alignment attributes (see Section 11.4 [Block Layout], page 166). You need two of those separated by a comma and enclosed in parenthesis. All the possibilities of how to form these values apply, with the difference that if you omit the `at` symbol and the parts after (just specify a block or blocks with or without an offset), the center of the block(s) is taken (with the offset potentially). A few examples

`(A,B)` The center of A and the middle of B.

(A+10, B@top)

10 pixels right of the centerline of A and the top of B.

(A+B, A+B)

The midpoint of the bounding box of A and B

(A@33%, B@m33%)

The left thirdpoint of A (ignoring its margins) and the top thirdpoint of B (including its margins).

(10,20) A coordinate relative to the parent's inner margin's (top-left corner).


(50%+10, 0)

Ten pixels right of the vertical centerline of the parent at the top internal margin of it.

(m0%+15, 100%)

Fifteen pixels right of the inner line of the parent and at the bottom inner margin of it. For a detailed explanation on how the 'm' character modifies number-only coordinates referring to the parent, see Section 11.4 [Block Layout], page 166.

```
box_A;
below rightof +30 box_B;
(A,B) -> (A+10, B@top) [color=red];
(A+B, A+B) -> (B@m33%, A@33%);
(A@left-5, A@top) ++ (,B@bottom);
(A@left-7, A@top) ++ (,B@bottom);
```



If you use a coordinate at least on one end of an arrow, you can omit one of the coordinates. This will result in a horizontal or vertical arrow, taking the omitted coordinate from the other end.

After coordinates, you can also specify a direction (but not a port number) after the @ symbol. You can use both compass points and degrees as for blocks. A second @ symbol with a number may follow the direction specifying the distance of the waypoint (since there is no block to attach a distance to). After that you can further add fine-tuning modifiers to the arrow end as specified in Section 11.5.3 [Fine-tuning Arrow Ends], page 178. Thus this is a valid arrow: (a@10,b@20)@sw@10+x10-5.

If an arrow (or line) starts (or ends) at a coordinate, where the specification of both the X and Y coordinates contain the same single block, that block is used to clip the arrow (if the arrow starts inside the block). This allows coordinates to be used to specify a location for the block not exposed as a port. For such arrow ends, arrow de-overlapping is also executed as explained in Section 11.5.4 [Automatic De-Overlapping at Arrow Ends], page 178. Arrows starting from coordinates which have different or multiple blocks in the X and Y coordinate specification do not get de-overlapped.

11.5.6 Arrow Labels and Markers

You can add labels and markers to arrows/lines. Markers are dots, tick marks, etc. along the line.

To add a single label, simply add a label to the arrow using the `label` attribute or the colon-label syntax as for blocks: `a->b: This is the label.`; The position of the label can be set using the `label.align` attribute, with the values of 0 and 100 representing the start and end of the arrow (with values proportionally in between, but you can also have negative values or above 100%). With `label.pos` you can set how the label is placed in relation to the arrow line. It can take one of `left`, `right`, `above`, `below` or `center`, the latter meaning onto the line.

To add a single marker, use the `label.align` to set its position and optionally the `marker.type` attribute to set its shape, using any of the arrowhead types listed in Section 11.5.7 [Arrow and Line Attributes], page 183. You can apply to any of the arrowhead attributes to `marker.*`, specifically `marker.type`, `marker.gvtype` (to set the arrowhead type in graphviz style), `marker.size` and `marker.color`.


You can add more labels by specifying them in curly braces after the arrow. Their syntax is as follows.

```
[mark] [extend] <position> [<attributes>];
```

`<position>` is specified as a) a percentage number (followed by a percentage sign); b) a percentage number (followed by a percentage sign) followed by an optional sign and an offset number in pixels; or c) an offset number in pixels (no percentage sign). The percentage number is interpreted so that 0% means the start of the arrow, 100% means its entire length. A positive percentage number is measured from the start of the arrow forward, while a negative value means measuring it from the end of the arrow backwards. Thus, 99% is the same as -1%. The offset value is then added/subtracted from this position (interpreted in pixels). If only a pixel value is specified, a positive value is measured from the start of the arrow forward, a negative value backward.

If you specify `extend`, the start of the measurement changes: for positive values we start to measure from the end of the arrow forward (beyond the end of the arrow), for negative values we start measuring from the start of the arrows (before the start). It is actually simpler than it sounds, here are a few examples.

```
box_a; space 200; box_b;
a->b {
    use label.orient=left;
    use text.color=gray;
    mark 10%: 10%;
    mark -30%: -30%;
    mark -30%+20: -30%+10;
    mark -30%-20: -30%-10;
    mark 10%+20: 10%+20;
    mark 80: 80;
    mark extend -20%: -20% (extend);
    mark extend 20%: +20% (extend);
}
```



Adding the `mark` keyword before the position will add a marker at the designated position. In this case there is no need to specify a label, if you just want the marker. The default type for markers is `marker.gvstyle=tick`, which is a thin line perpendicular to the arrow line.


Among attributes you can use `marker.*` to set the marker style; `label`, `text.*`, `label.pos` and `label.orient` to specify the label and its details; and `line.*`, `fill.*` and `shadow.*` to specify a box around the label.

As a shorthand, you can omit the curly braces and simply specify the labels after the arrow specification.

The label orientation can be influenced with `label.orient` the same way as for blocks, but for arrow labels, two more values are available: `tangent` and `tangent_upside_down`. These rotate the label to match the tangent of the arrow line at the label position. Layout makes an effort to lay out the label to be readable left-to-right, but you can also request it to be upside down with `tangent_upside_down`.

```
col_C {
    box_A; space 20; box_B;
    arrows use label.pos=right;
    A@e->B@e {
        mark -12%: aaa [label.orient=tangent];
        mark 6%: bbb [label.orient=tangent_upside_down, label.pos=right];
        mark 50%: ccc [label.orient=tangent];
    }
}
space 50;

row_R {
    box_A; space 20; box_B;
    arrows use label.pos=below;
    A@s->B@s {
        mark -12%: aaa [label.orient=tangent];
        mark 6%: bbb [label.orient=tangent_upside_down];
        mark 50%: ccc [label.orient=tangent];
    }
}
```



11.5.7 Arrow and Line Attributes

The following attributes apply to arrows. All of them can be made part of a style, except `via` and `routing.order`.

line.* These set the line style of the arrow. See Section 8.6 [Common Attributes], page 85, for more. `line.corner` and `line.radius` is used when the `routing` attribute is set to `polygon`.

text.*

label.* These attributes specify one label of the arrow.

marker.* Together with `label.align` these specify one marker on the arrow.

via


Sometimes you want to better influence what path the arrow takes. In this case use the `via` attribute. You can specify any of the four sides of any block, for example `block@top` or any corner, such as `block@bottomleft`. The arrow will go around the named block on that side. You can also specify a coordinate - this will be a mandatory waypoint for the arrow. You can specify any number of `via` attributes mixing the blockside and coordinate types as you wish. The arrow will visit them in the specified order. Note that Msc-generator still tries to pick the shortest route (while staying on the specified side of the specified block(s) and going through the mandatory waypoints). Also note that if the

arrow starts from or ends at a block and due to the `via` attributes it goes through that block several times, it is trimmed to the first crossing.


```

col_c1 {
    row {
        box_A;
        below box_B;
        box_C [top=A];
    }
    B->C [via=A@topleft,
            line.color=red];
    use margin=0;
    text: via A@topleft;
}
below 20 col_c2 {
    row {
        box_A;
        below box_B;
        box_C [top=A];
    }
    B->C [via=A@right,
            line.color=red];
    use margin=0;
    text: via A@right;
}

```



`via A@topleft`



`via A@right`


`cross`
`cross_all`

You can specify these attributes several times, if needed. Their value can be a block name (or a list of block names separated by plus signs). The arrow or line is allowed to cross any block, if listed in the `cross` attribute; or any block and its children if listed in the `cross_all` attribute. This complements the `allow_arrows` attribute of blocks. While that attribute allows any and all arrows to cross the block it is set on, the `cross` and `cross_all` attributes can be specified selectively per arrow.

distance This attribute dictates how large is the distance between the arrow and the blocks it goes around. Setting it to single number (meant in pixels) will impact all blocks it goes around. Setting it in the syntax `<block>@<number>` will make the arrow to use this distance only for `<block>`. You can apply this attribute multiple times, or use the plus sign to list several blocks, like `distance=<block>+<block>+...@<number>`.

routing This attribute governs, how arrows are laid out. If set to `straight` the arrow will go between the two blocks in a straight line, potentially over other blocks. If set to `horizontal` or `vertical`, the straight line will be parallel to the X or Y axis, respectively. `grid` will make it either horizontal - or vertical if horizontal routing is not possible. If set to `polygon`, the arrow will go the shortest path around blocks in the way. `manhattan` will result in using only horizontal and vertical edges attempting to go around blocks in the way. For the latter two

modes, the `line.corner` and `line.radius` attributes can be used to round the vertices of the polygon. If set to `curvy` (default), the arrow will avoid blocks in the way but a smoothing is applied to make the arrow one smooth line. There are limitations to this, but Msc-generator attempts to do a good job.



`routing.factor`

This attribute governs how the arrow layout algorithm trades off between minimizing arrow length and minimizing the number of turns the arrow makes. The value 0 means considering only arrow length (default), while 1 means considering only minimizing turns. Values in between balance. Note that arrow layout starts with a straight line between start and end (or waypoints) and then attempts to go round blocks in-between. Thus, even if you prefer minimizing turns, it will not by itself consider a big detour. Use the `via` attribute for that.

`routing.clip_block`

If set to `yes` (default) arrows/lines starting/ending at a block will be clipped by that block. Also applies if the ending is specified as a coordinate using the same, single block in the specification of both the X and Y coordinate.


`route.arrow_distance`

Overlapping arrow endings will get spaced apart this much, see Section 11.5.4 [Automatic De-Overlapping at Arrow Ends], page 178. Default is 5 pixels.

`routing.extend`

If arrow endings get outside the block due to de-overlapping, we extend them to the perimeter of the block, if this is set to `yes` (default).

```
cell_x {
    box_A;
    below rightof 20 box_B;
    A@se->B@n, B@nw, B@w;
}
cell_y {
    box_A;
    below rightof 20 box_B;
    A@se->B@n, B@nw, B@w
    [routing.extend=no];
}
```



`routing.block_others`

If set, it prevents overlap between the middle of arrows subject to de-overlapping, see Section 11.5.4 [Automatic De-Overlapping at Arrow Ends], page 178.

`routing.try_harder`

If set, at de-overlapping Msc-generator will try all permutations of arrow ordering until the best arrow layout is found (fewer turns and shorter arrows). If not

set, the search is stopped at the first layout where the de-overlapped arrows do not overlap.

`routing.allow_joint_layout`

This attribute defaults to `yes` meaning that at de-overlapping arrows that initially fully overlap Msc-generator will route this arrow side-by-side along the same path with the others it initially overlapped. If set to `no`, this arrow will be laid out independently. It can be useful if you want the arrows to go around an obstacle on different paths.

`routing.order`

If the above is set, the layout order of arrows is important. You can assign a number to this attribute to govern the repeated layout of arrows after their endpoints have been shifted due to overlaps.

`draw_before`

`draw_after`

You can specify one or more blocks (separated by plus signs or via repeating the attribute) the block shall be drawn before or after. You can use only one of the two attributes for any diagram elements. Setting one will automatically clear the other.

Arrows and lines have a single default style `arrow`. Then we have eight refinement styles `--`, `==`, `..`, `++`, `->`, `=>`, `>` and `>>`. The former four sets `routing=straight`, but otherwise they are just setting the line style by default. Of course, you can change them. Arrow labels have the default style `label`. This is also the default style for label-less markers, so there you can set the default marker type, too.

11.6 Replicating parts of the Diagram

To minimize typing, Msc-generator offers tools to replicate parts of the Diagram.

First, you can define a procedure, which you can replay many times, see Section 8.11 [Procedures], page 94. Second, you can copy an already placed block modifying its attributes and also its content (if any). Third, you can create many copies of a set of blocks with one line.

11.6.1 Copying a Block

Using the `copy` command, you can repeat a previously defined block.

```
copy <block_name> [as <new_name>] [<attributes>] [<content_update>];
```

Using the `as` clause you can give a name to the newly created copy (so that you can reference it later). Specifying any attribute will override the attributes of the original block. The type of the original block cannot be changed. The content of the original block is also copied, including arrows.

After the attributes you can also specify a list of updates to the content of the block in curly braces. Here you can use the following commands.

```
add [before <block>] <block_description>
```

Add a new block (of any kind) to the end of the content list, or before `<block>`. The added block can itself have content you can specify. Alternatively, the added block can itself be a copy or even an arrow.

drop <block>

Remove a block from the content list. You can specify a children using nested dot notation, so if the block to copy contains a block `inner`, which in turn contain a block `innermost`, you can drop the latter by saying `drop inner.innermost`. You can also drop several blocks, list them separated by commas.

[recursive] update <block>|blocks|arrows|all <attributes>

Using this commands you can apply changes to the attributes of a block, all blocks, arrows or simply all content of the block under copy. If you add `recursive`, the attribute update also applies to content of the content recursively. Similar to drop, you can specify a block to update that is not directly the children of the copied block. You can also specify a list of blocks, separated by commas - this applies the change to a specific set of contained blocks.

replace <block> <block description>

This command removes `<block>` and adds the block specified afterwards in its place. Again, similar to dropping or updating `<block>` can denote an indirect child.

move [before <block1>] <block2>

This command moves `<block2>` to the end of the content list or before block `<block1>`. You can only reorder the content list of the copied block here not that of any children.

use <attributes>**blocks use <attributes>****arrows use <attributes>**

These commands (as elsewhere) update the running style impacting block and/or arrows you add or replace subsequently.

Currently there is no way to drop or replace an arrow, but you can add arrows and replace blocks to arrows.

11.6.2 Block Templates

In addition to copy blocks that were previously added to the diagram, you can also define templates of blocks. Block Templates are fully defined blocks (of any kind), that are not part of the diagram. They do not appear and cannot be referenced (e.g., for alignment or arrows). They, however, can serve as a basis of other blocks, that do appear.

To define a template, simply add `template` in front of any block. The name of the block becomes the name of the template.¹² The template captures the default styles and running styles at the location of definition and replaying it will not be impacted by a later change to them.

To create an actual box from the template use the `copy <template>` command. Add `as <name>` if you want the block created to have a different name. Note that the `copy` command can also be used to copy existing blocks. If you have a template and a block of

¹² If you have defined the template inside another block (wierd), the full name of the block becomes the name of the template.

the same name, `copy` will use the template. When creating a box from a template, you can change attributes and content of the template, similar to the case of copying an existing block, see Section 11.6.1 [Copying a Block], page 186.

Templates are part of the chart and do not disappear after closing the scope (at the closing brace). Template definitions cannot be made part of design specifications, but can be part of procedures. In the latter case replaying the procedure will result in the definition of the template at the location of the replay. You can also create a modification of a template by copying it. Simply type `template` before a `copy` command and the result will not be a block that appears but a template. This construct is also suitable to create a template (possibly with modifications) from an existing block - simply type `template copy <block_name> as <template_name>`.

11.6.3 Repeating a Block Many Times

This functionality is not yet implemented.

11.7 Chart Options and Commands

You can set the following chart options by simply adding `<option>=<value>` to the chart at any place. All of them can be made part of designs (with the exception of `numbering.append`).

`numbering`

`numbering.*`

These options control automatic numbering and are described in Section 8.4 [Numbering], page 80.

`text.*` These attributes can be used to set the default font in the chart. See Section 8.6.3 [Text Formatting Attributes], page 87.

`background.color`

`background.color2`

`background.gradient`

These can be used to set the background fill of the diagram. You can specify them at any point during the diagram specification - the last setting will be used.

`conflict_report`

This option can be used to set the report level of alignment conflicts, see Section 11.4.4 [Block Layout Conflicts], page 174. You can set it multiple times and it honours scoping - its effect lasts until the next closing brace.

`pedantic` When set to `no` any block mentioned in an attribute update or arrow that has not yet been defined, will be automatically created. This prevents forward referencing blocks in arrow definitions, but is handy to quickly define a lot of blocks. See Section 11.2 [Block Types and Definition], page 158.

Block diagrams support the generic `defcolor`, `defstyle` and `defshape` commands to define colors, styles and shapes, respectively (see Section 8.5 [Specifying Colors], page 83, Section 8.8 [Defining Styles], page 90, Section 8.10 [Defining Shapes], page 92); `usedesign <design_name>`; to apply a design (both full or partial designs); `include` to include another

text file (see Section 8.13 [File Inclusion], page 98); and `defproc`, `replay` and `if` commands to define and use procedures (see Section 8.11 [Procedures], page 94). These are all generic commands available also in signalling charts and graphs (with some minor variations).

In addition to the commands, Block Diagrams support another command: `use`. After `use` you can specify a list of attributes, the same way as for any chart element (but you do not need the square brackets). This has an effect of applying the listed attributes on all subsequent blocks and arrows (until the scope ends, see Section 8.7 [Scoping], page 89). This allows a quick shorthand to save typing. The feature is implemented via a hidden *running style*, which is applied to all newly defined blocks or arrows after the default and enhancement styles, but before the explicit attributes. Thus the values you set this way can override default and enhancement style values, but can be overridden by explicitly specified attributes. Note that you can list styles after `use`, for example, `use col;` will apply the `col` style to any subsequent elements, making blocks align from top to bottom.